

# **Computer Structure 234267**

**Lecturers:  
Lihu Rappoport  
Adi Yoaz**

# General Course Information

## ◆ **Grade**

- ❖ 20% : 4 exercises (all mandatory) תרגול – submission in pairs
- ❖ 80% final exam

## ◆ **Come to the lectures and to the tutorials !**

- ❖ The material for the exam includes all that is taught during the lectures and the tutorials – the foils do not contain everything
- ❖ While in the classroom wear masks properly over the nose, and keep distance

## ◆ **Course web site**

- ❖ <http://webcourse.cs.technion.ac.il/236267/>
- ❖ Foils will be on the web several days before each lecture

# Class Focus

## ◆ CPU

- ❖ Introduction: performance, instruction set (RISC vs. CISC)
- ❖ Pipeline, hazards
- ❖ Branch prediction
- ❖ Out-of-order execution

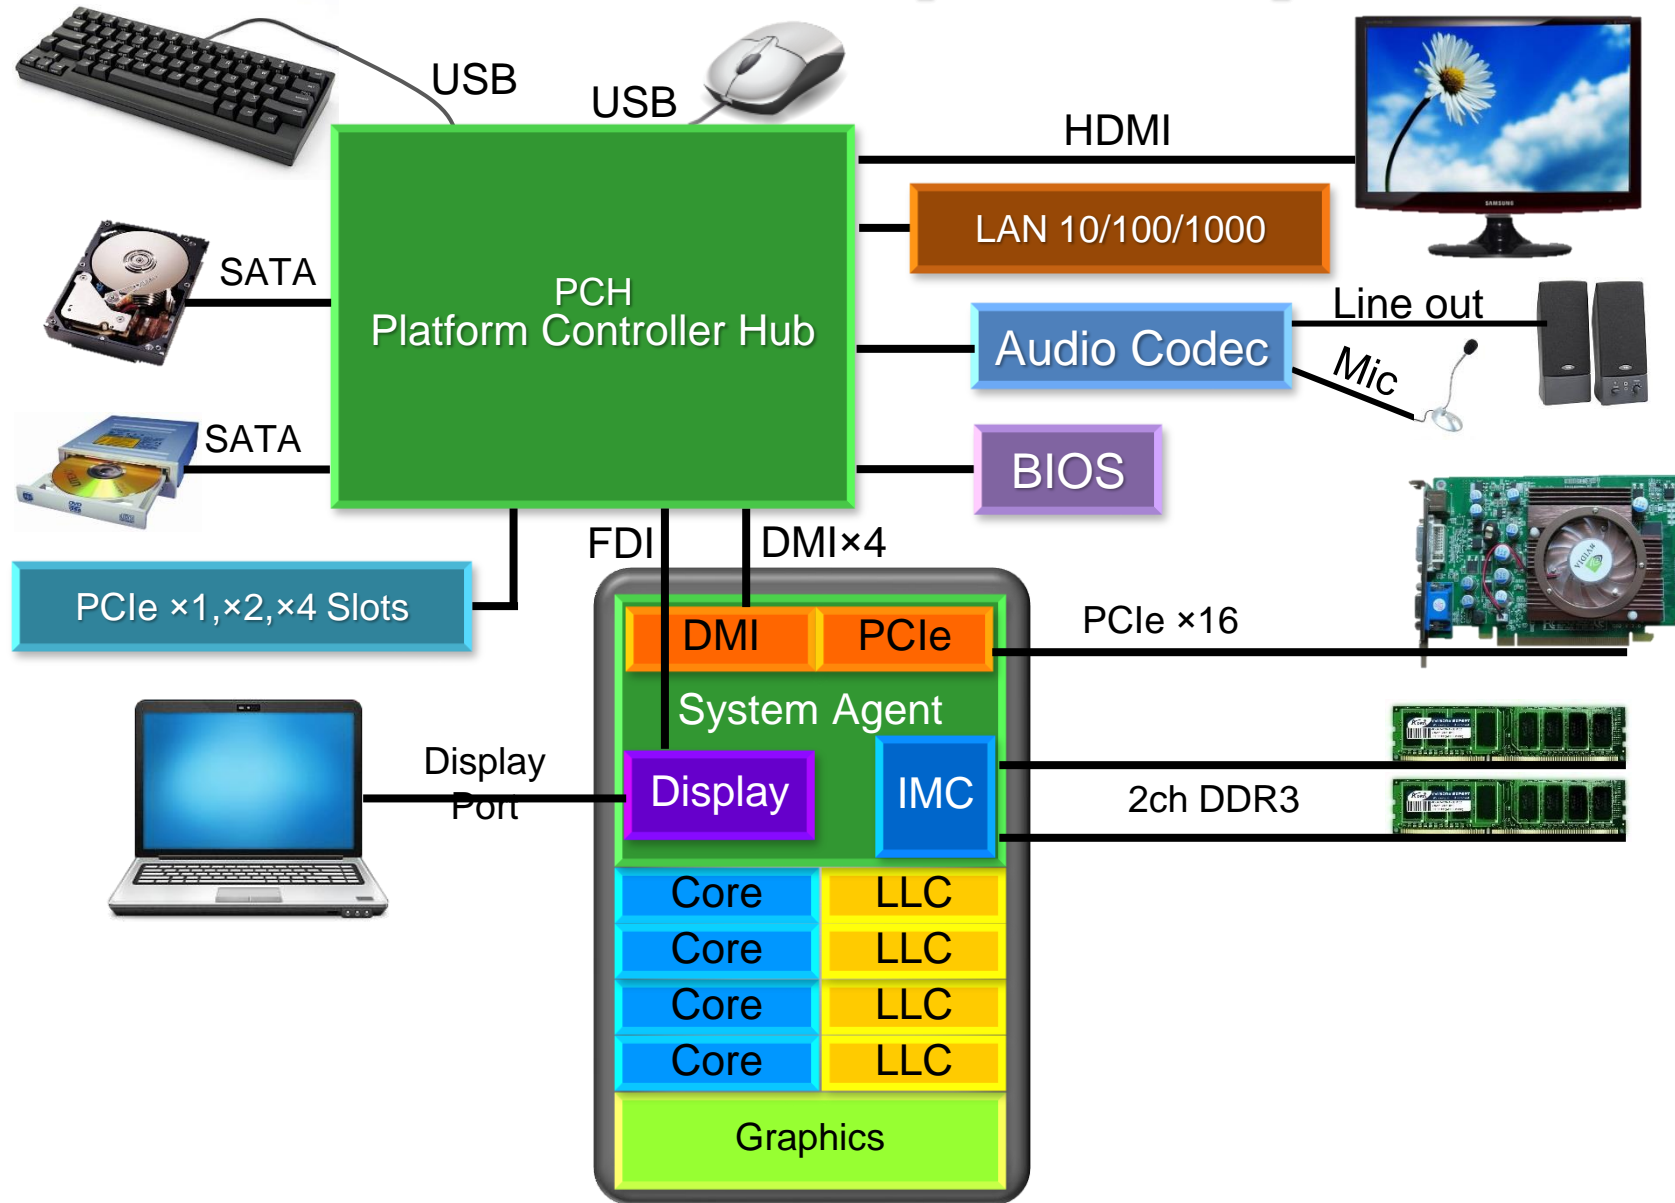
## ◆ Memory Hierarchy

- ❖ Cache and cache coherency
- ❖ Main memory
- ❖ Virtual Memory

## ◆ More topics

- ❖ Multi-threading
- ❖ System
- ❖ Power considerations

# Personal Computer System



# Architecture & Microarchitecture

## ◆ Architecture

**The processor features seen by the “user”**

- ❖ Instruction set, addressing modes, data width, ...

## ◆ Micro-architecture

**The internal implementation of a processor**

- ❖ Caches size and structure, number of execution units, ...

## ◆ Processors with different $\mu$ Arch can support the same architecture


## ◆ Compatibility

- ❖ A new processor can run existing software
- ❖ The processor  $\mu$ Arch is new, but it supports the Architecture of past generations, possibly adding to it

# Moore's Law: The number of transistors on microchips doubles every two years Our World in Data

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

## Transistor count



## MOSFET scaling

- 10 µm – 1971
- 6 µm – 1974
- 3 µm – 1977
- 1.5 µm – 1981
- 1 µm – 1984
- 800 nm – 1987
- 600 nm – 1990
- 350 nm – 1993
- 250 nm – 1996
- 180 nm – 1999
- 130 nm – 2001
- 90 nm – 2003
- 65 nm – 2005
- 45 nm – 2007
- 32 nm – 2009
- 22 nm – 2012
- 14 nm – 2014
- 10 nm – 2016
- 7 nm – 2018
- 5 nm – 2020

## Future

- 3 nm ~ 2022
- 2 nm ~ 2024

Data source: Wikipedia ([wikipedia.org/wiki/Transistor\\_count](https://wikipedia.org/wiki/Transistor_count))

OurWorldinData.org – Research and data to make progress against the world's largest problems.

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser.

# CPU Performance

## ◆ CPUs work according to a clock signal

- ❖ Clock cycle is measured in nsec ( $10^{-9}$  of a second)
- ❖ Clock frequency ( $= 1/\text{clock cycle}$ ) measured in GHz ( $10^9$  cyc/sec)

## ◆ CPI – Cycles Per Instruction

- ❖ Average #cycles per Instruction (in a given program)

$$\text{CPI} = \frac{\text{Total number of cycles required to execute the program}}{\text{Total number of instructions executed in the program}}$$

- ❖ IPC ( $= 1/\text{CPI}$ ) : Instructions per cycles

# Calculating the CPI of a Program

- ◆  $IC_i$ : #times instruction of type  $i$  is executed in the program

- ◆  $IC$ : #instruction executed in the program:  $IC = \sum_{i=1}^n IC_i$

- ◆  $F_i$ : relative frequency of instruction of type  $i$ :  $F_i = IC_i/IC$

- ◆  $CPI_i$  – #cycles to execute instruction of type  $i$

- ◆ e.g.:  $CPI_{add} = 1$ ,  $CPI_{mul} = 3$

- ◆ #cycles required to execute the entire program:

$$\#cyc = \sum_{i=1}^n CPI_i * IC_i = CPI * IC$$

- ◆  $CPI$ :

$$CPI = \frac{\#cyc}{IC} = \frac{\sum_{i=1}^n CPI_i * IC_i}{IC} = \sum_{i=1}^n CPI_i * \frac{IC_i}{IC} = \sum_{i=1}^n CPI_i * F_i$$



# CPU Time

- ◆ **CPU Time - time required to execute a program**


$$\text{CPU Time} = \text{IC} \times \text{CPI} \times \text{clock cycle}$$

- ◆ **Our goal: minimize CPU Time**

- ❖ Minimize clock cycle: more GHz (process, circuit, uArch)
- ❖ Minimize CPI: uArch (e.g.: more execution units)
- ❖ Minimize IC: architecture (e.g.: vector instructions)

# Amdahl's Law

Suppose enhancement E accelerates a fraction  $F$  of the task by a factor  $S$ , and the remainder of the task is unaffected, then:



$$t'_{\text{exe}} = t_{\text{exe}} \times \left[ (1 - F) + \frac{F}{S} \right]$$

$$\text{Speedup}_{\text{overall}} = \frac{t_{\text{exe}} - t'_{\text{exe}}}{t_{\text{exe}}} = 1 - \frac{t'_{\text{exe}}}{t_{\text{exe}}} = 1 - \left[ (1 - F) + \frac{F}{S} \right] = F - \frac{F}{S}$$

# Amdahl's Law: Example

- Floating point instructions improved to run 2× faster
- 10% of the executed instructions are FP instruction



$$t'_{\text{exe}} = t_{\text{exe}} \times (0.9 + 0.1 / 2) = 0.95 \times t_{\text{exe}}$$

$$\text{Speedup}_{\text{overall}} = 1 - 0.95 = 0.05 = 5\%$$

**Corollary:**  
**Make The Common Case Fast**

# Evaluating Performance of future CPUs

- ◆ **Use a performance simulator to evaluate the performance of a new feature / algorithm**
  - ❖ Models the uarch to a great detail
  - ❖ Run 100's of representative applications
- ◆ **Produce the performance s-curve**
  - ❖ Sort the applications according to the IPC increase
  - ❖ Baseline (0) is the processor without the new feature



# Using Benchmarks for Comparing Performance


## ◆ **Benchmark types**

- ❖ Real applications, or representative parts of real apps
- ❖ Synthetic benchmarks which represent typical workloads

## ◆ **Benchmarks are targeted at the specific system usages**

- ❖ SPEC INT – integer applications: data compression, C compiler, Perl interpreter, database system, chess-playing, Text-processing, ...
- ❖ SPEC FP – floating point applications: fluid dynamics, quantum chemistry, image ray-tracing, speech recognition, ...
- ❖ SYSMark – reflects usage patterns of business users
  - office productivity: Acrobat, Excel, PowerPoint, Word
  - media creation: Photoshop, Premiere, Sketchup
  - data/financial analysis: Excel, WinZip
- ❖ TPC Benchmarks - transaction-processing throughput

# Instruction Set Design

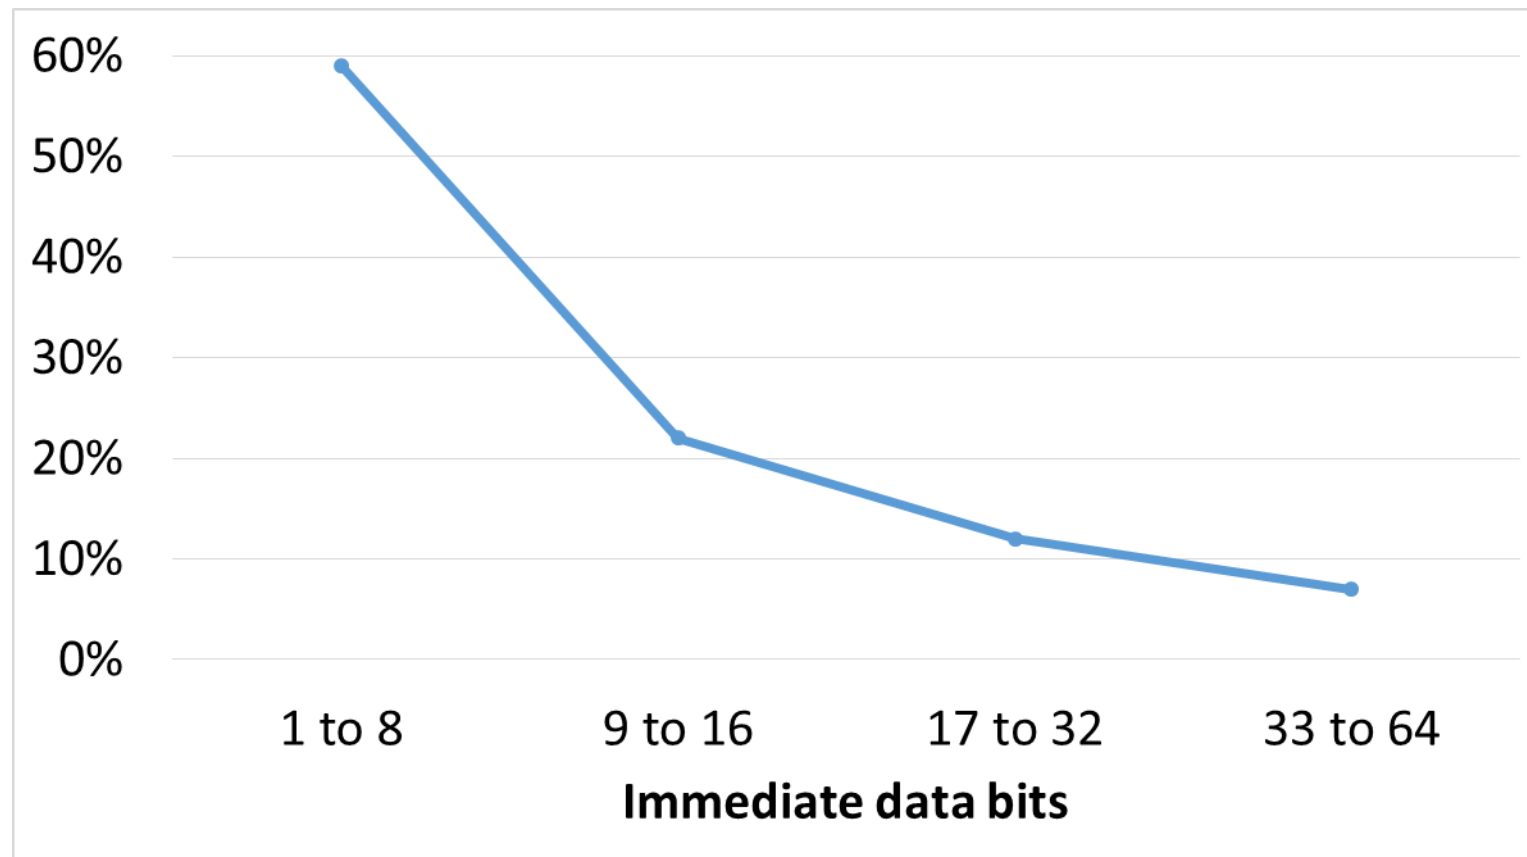


# ISA Considerations

- ◆ **Reduce the IC to reduce execution time**
  - ❖ A single vector instruction performs the work of multiple scalar instructions
- ◆ **Simple instructions  $\Rightarrow$  simpler HW implementation**
  - ❖ Higher frequency, lower power, lower cost
- ◆ **Code size**
  - ❖ Long instructions take more time to fetch and require a larger memory
    - Important for large code footprint workloads

# Architectural Consideration Example

- ◆ **Most instructions use a short immediate value**
  - ❖ Having all instructions support the maximum size needed would increase code size






# ISA Extensions for Reducing IC


## ◆ Example

- ❖ 128-bit packed (vector) / scalar single precision FP (4×32)
- ❖ 128-bit registers (XMM0 – XMM7)

### Packed:



### Scalar:



- ❖ A single “packed add” instruction replaces 4 “add” instructions
- ❖ In order to fully gain the performance, need also new HW
  - execution unit that performs “packed add” in a single cycle

# CISC Processors

- ◆ **CISC – Complex Instruction Set Computer** – a high-level machine language
  - ❖ Example: x86
- ◆ **Complex instructions and complex addressing modes**
  - ⇒ complicates the processor and slows down the simple, common instructions
  - ⇒ contradicts **Make The Common Case Fast**
- ◆ **Small number of registers** ⇒ requires many memory accesses
- ◆ **Non-orthogonal registers:** some operations supported only on specific registers
  - ⇒ Not compiler friendly
- ◆ **ALU operations directly on memory**
  - `add eax, 7680[ecx] ; eax ← MEM[7680+ecx] + eax`
- ◆ **Variable length instructions**
  - ❖ Common instructions get short codes ⇒ save code length
  - ❖ Difficult to decode few instructions in parallel: only after an instruction is decoded, its length (and when the next instruction starts) is known
  - ❖ An instruction may cross a cache line or a page

# Top 10 Instructions

<u>Rank</u>	<u>instruction</u>	<u>% of total executed</u>
1	load	22%
2	conditional branch	20%
3	compare	16%
4	store	12%
5	add	8%
6	and	6%
7	sub	5%
8	move register-register	4%
9	call	1%
10	return	1%
	<b>Total</b>	<hr/> <b>96%</b>

**Simple instructions dominate instruction frequency**

# RISC Processors

## ◆ RISC – Reduced Instruction Set Computer

- ❖ Simple instructions enable fast hardware: make the common case fast
- ❖ Example: ARM

## ◆ Characteristics

- ❖ A small instruction set, with few instruction formats
- ❖ Simple instructions that execute simple tasks, most of them single-cycle (with pipeline)
- ❖ A few addressing methods, with ALU operations on registers only
  - Memory is accessed using Load and Store instructions only
- ❖ Many orthogonal registers – all instructions available with all registers
- ❖ Three address machine:   Add dst, src1, src2
- ❖ Fixed length instructions
- ❖ Compiler friendly – easier to utilize the ISA: orthogonal, simple

# RISC Processors (Cont.)

- ◆ **Simple architecture  $\Rightarrow$  Simple micro-architecture**
  - ❖ Simple, small and fast control logic
  - ❖ Simpler to design and validate
  - ❖ Leave space for large on die caches
  - ❖ Shorter time-to-market
- ◆ **Existing RISC processor are not “pure” RISC**
  - ❖ E.g., support division which takes many cycles
- ◆ **Modern CISC processors use RISC  $\mu$ -arch ideas**
  - ❖ Internally translate the CISC instructions into RISC-like operations
    - the inside core looks much like that of a RISC processor