

# **Computer Structure**

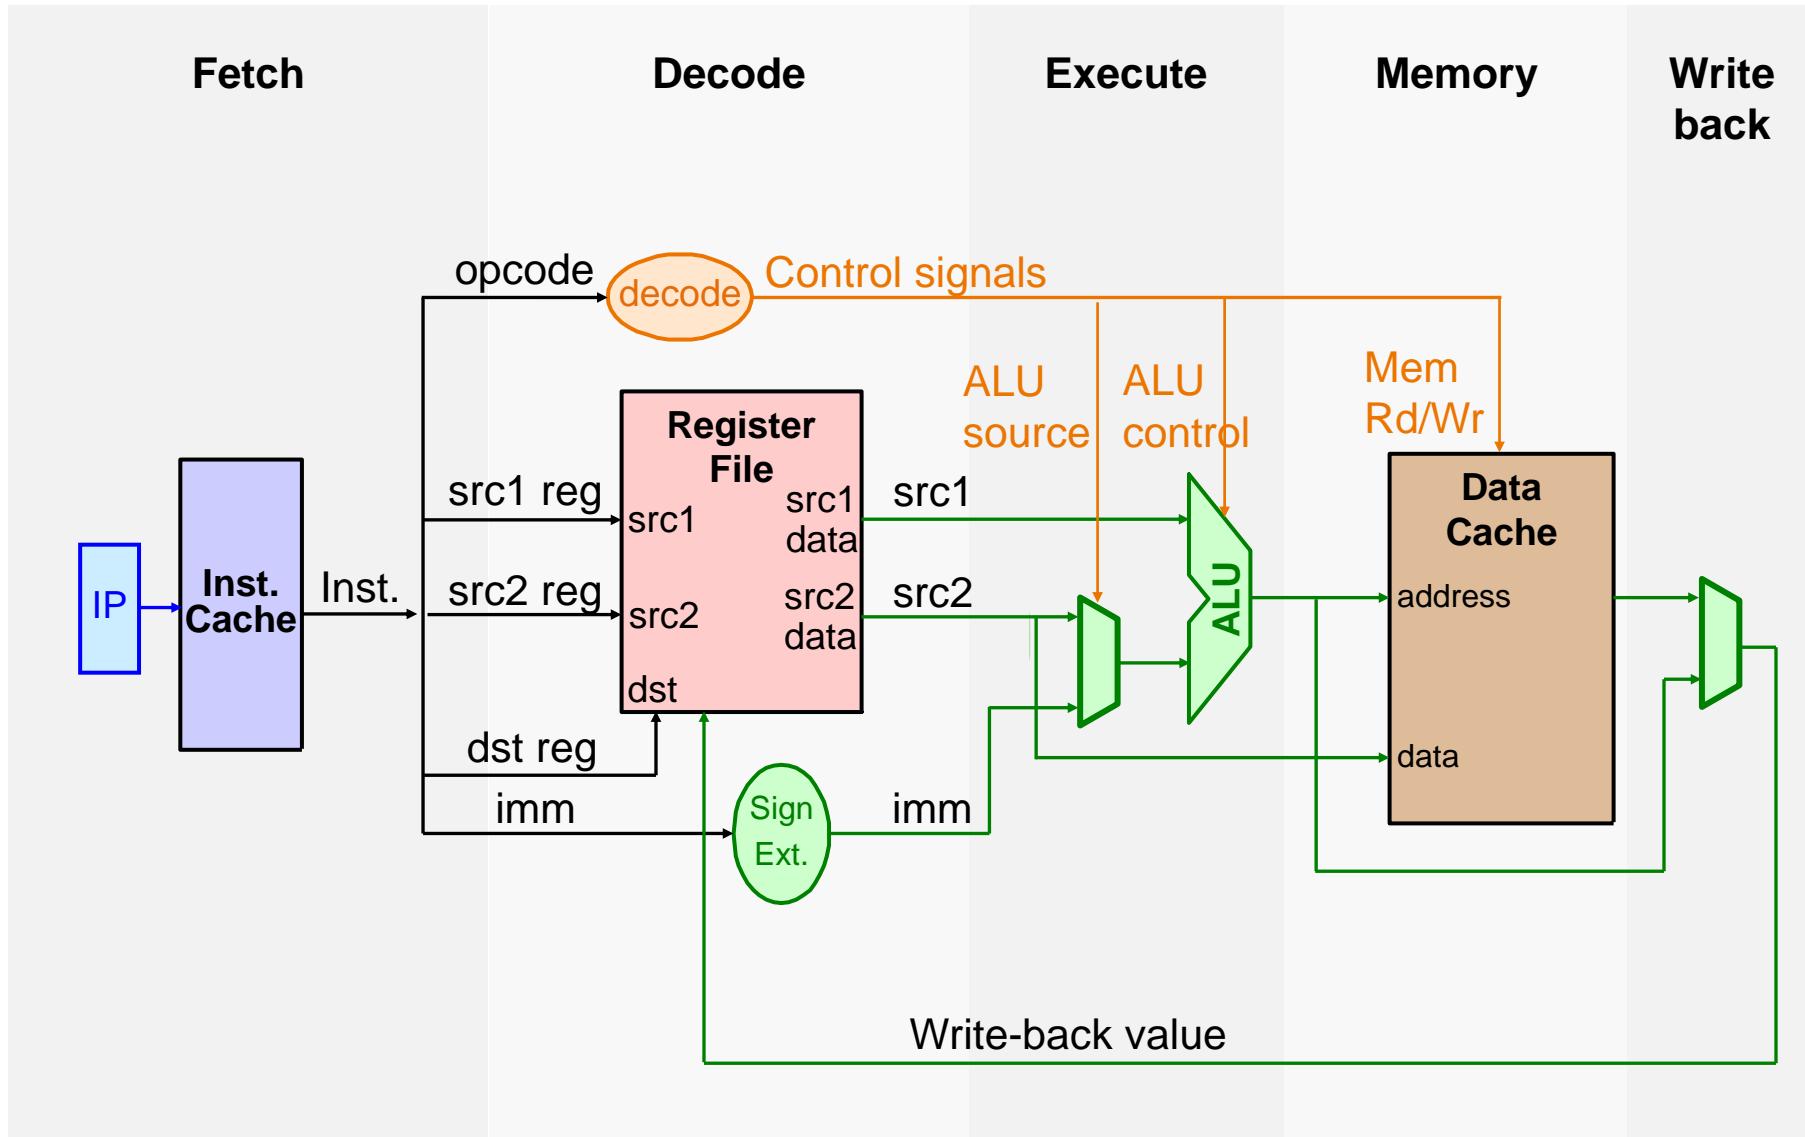
## **Pipeline**

**Lecturers:**

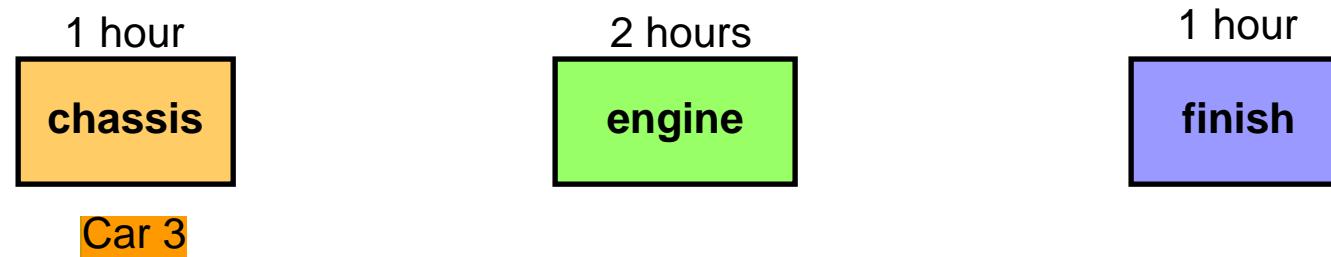
**Lihu Rappoport**

**Adi Yoaz**

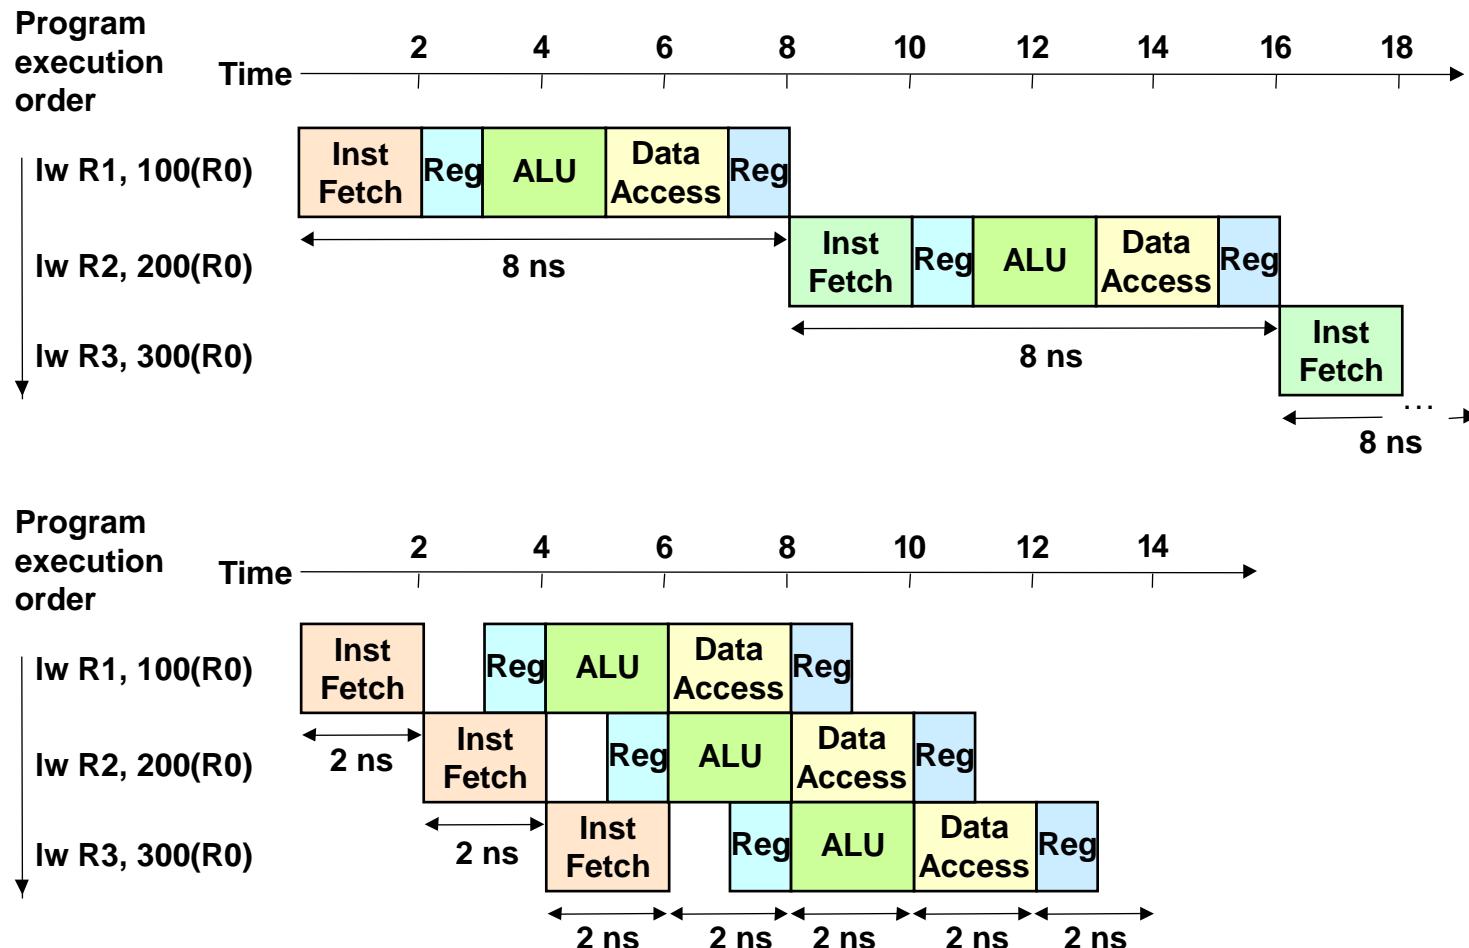
# A Basic Processor



# Pipelined Car Assembly



# Pipelining Instructions

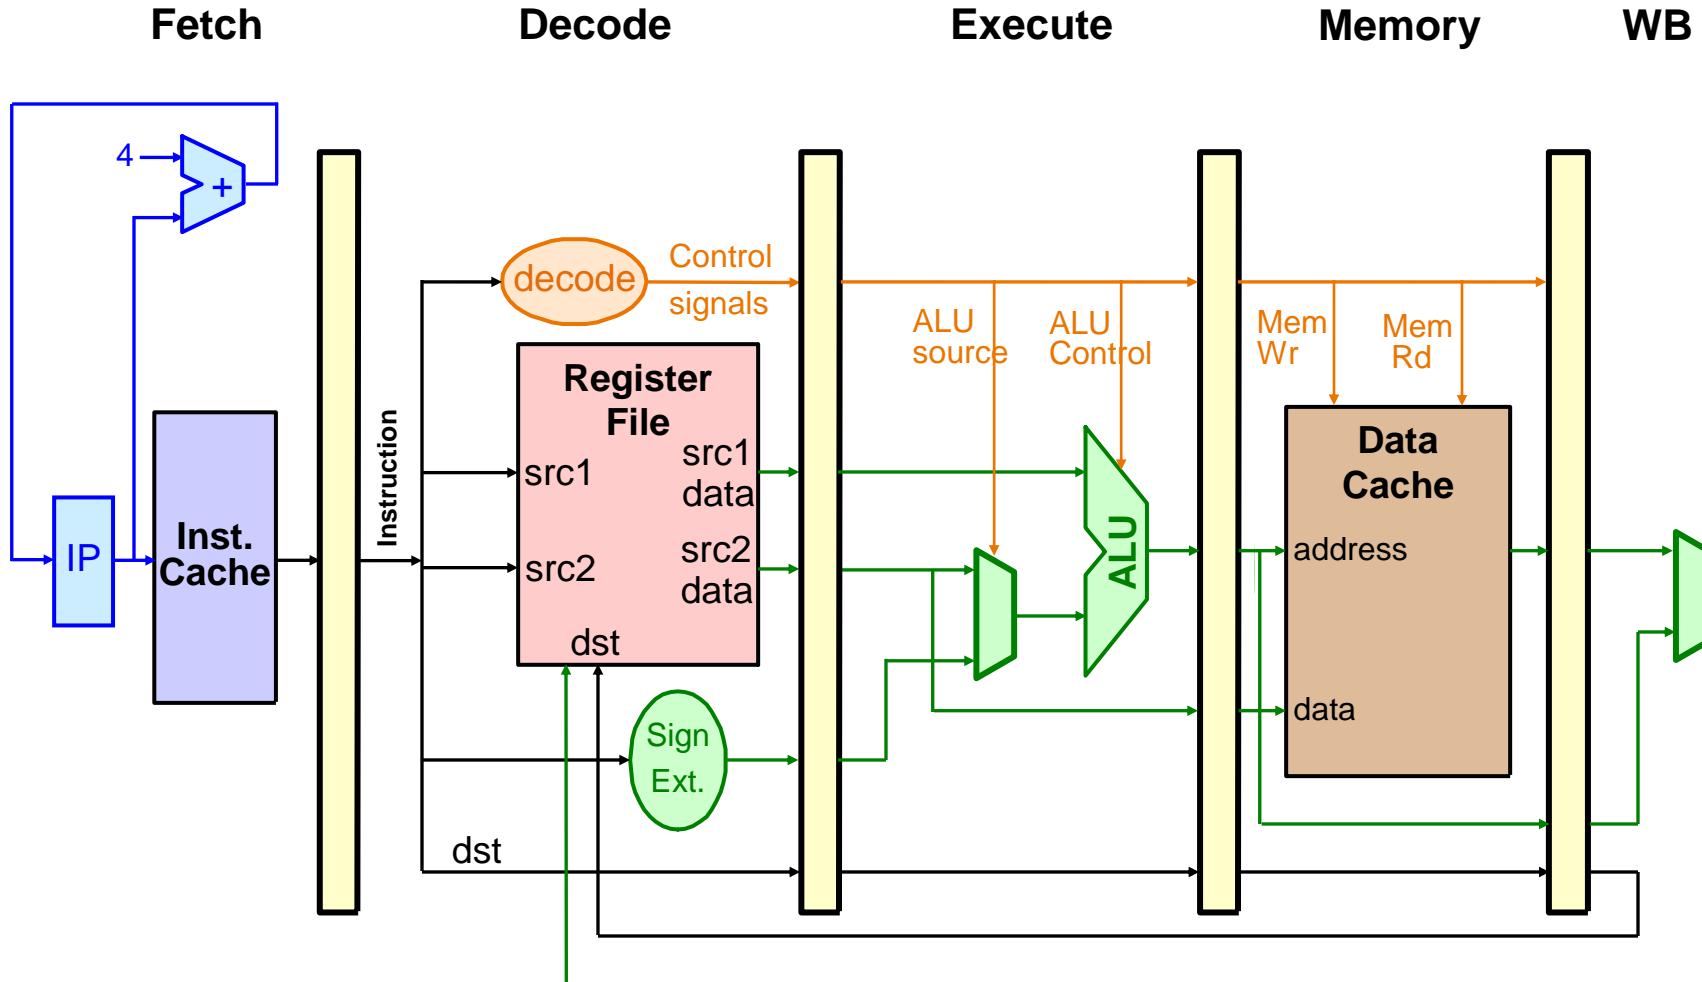


Ideal speedup is number of stages in the pipeline. Do we achieve this?

# Pipelining

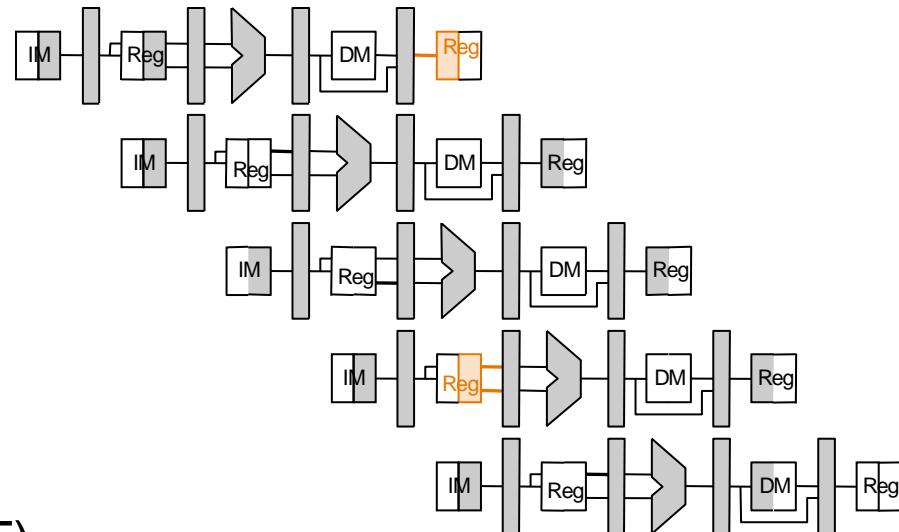
- ◆ Pipelining does not reduce the **latency** of single task, it increases the **throughput** of entire workload
- ◆ Potential speedup = Number of pipe stages
  - ❖ Pipeline rate is limited by the slowest pipeline stage
    - ⇒ Partition the pipe to many pipe stages
    - ⇒ Make the longest pipe stage to be as short as possible
    - ⇒ Balance the work in the pipe stages
- ◆ Pipeline adds overhead (e.g., latches)
  - ❖ Time to “fill” pipeline and time to “drain” it reduces speedup
  - ❖ Stall for dependencies
    - ⇒ Too many pipe-stages start to loose performance
- ◆ IPC of an ideal pipelined machine is 1
  - ❖ Every clock one instruction finishes

# Pipelined CPU

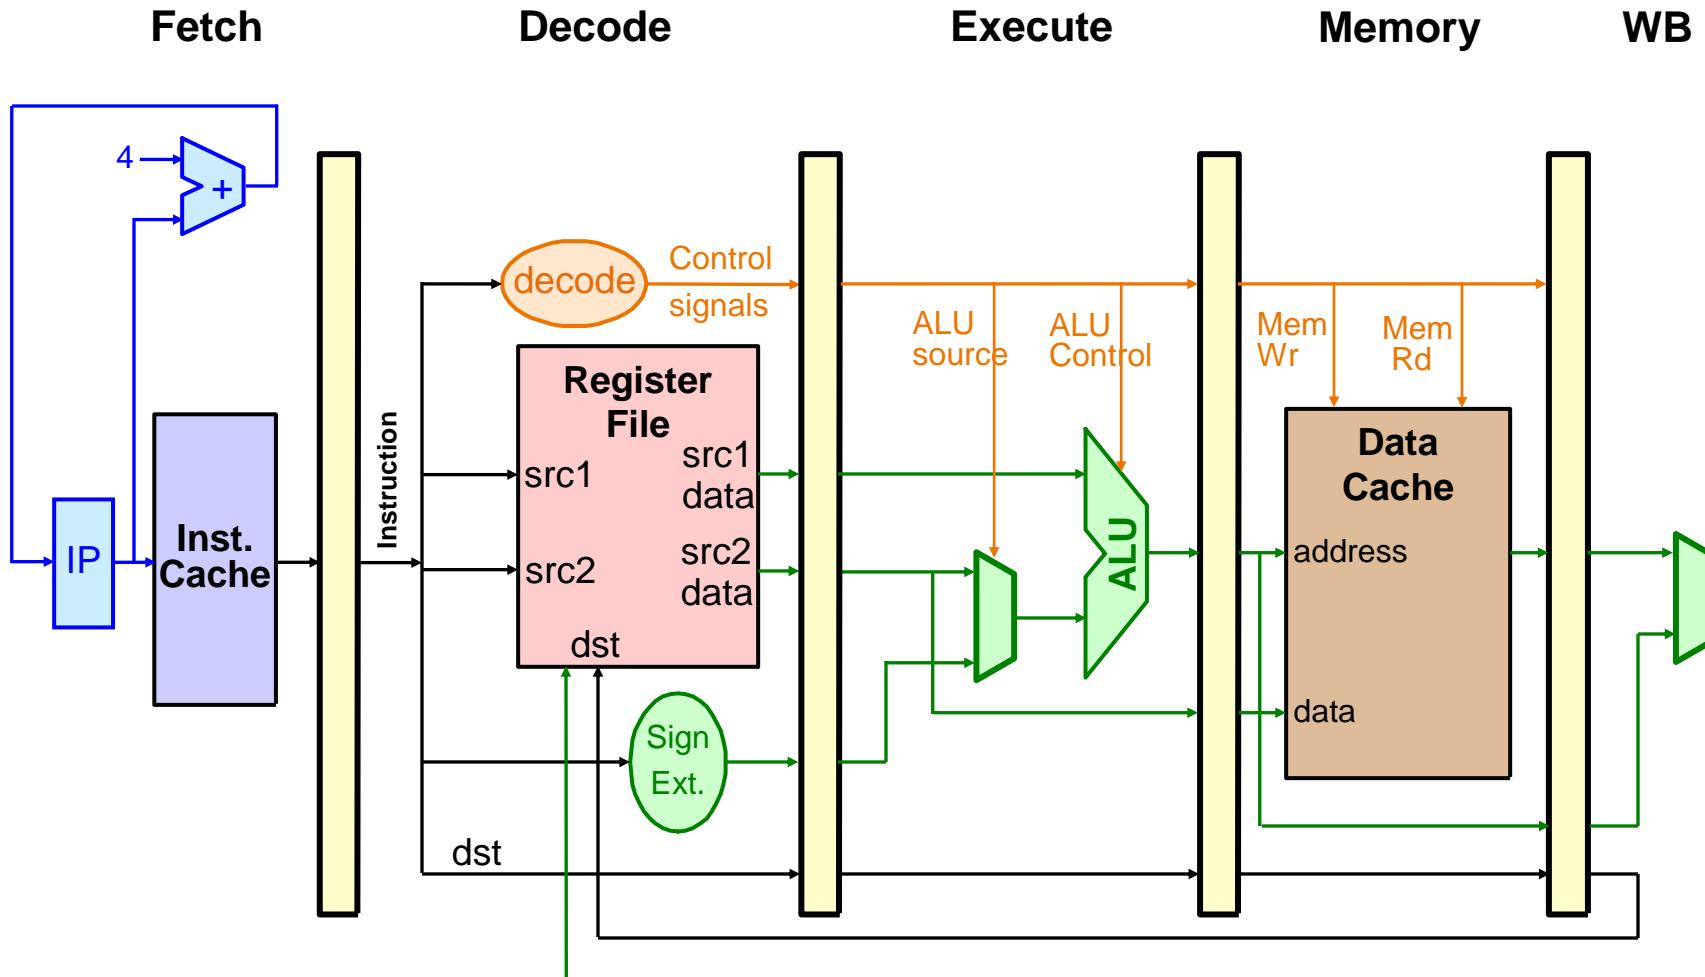


# Structural Hazard

- ◆ Different instructions using the same resource at the same time
- ◆ Register File:
  - ❖ Accessed in 2 stages:
    - Read during stage 2 (ID)
    - Write during stage 5 (WB)
  - ❖ Solution: 2 read ports, 1 write port
- ◆ Memory
  - ❖ Accessed in 2 stages:
    - Instruction Fetch during stage 1 (IF)
    - Data read/write during stage 4 (MEM)
  - ❖ Solution: separate instruction cache and data cache
- ◆ Each functional unit can only be used once per instruction
- ◆ Each functional unit must be used at the same stage for all instructions



# Pipeline Example: cycle 1

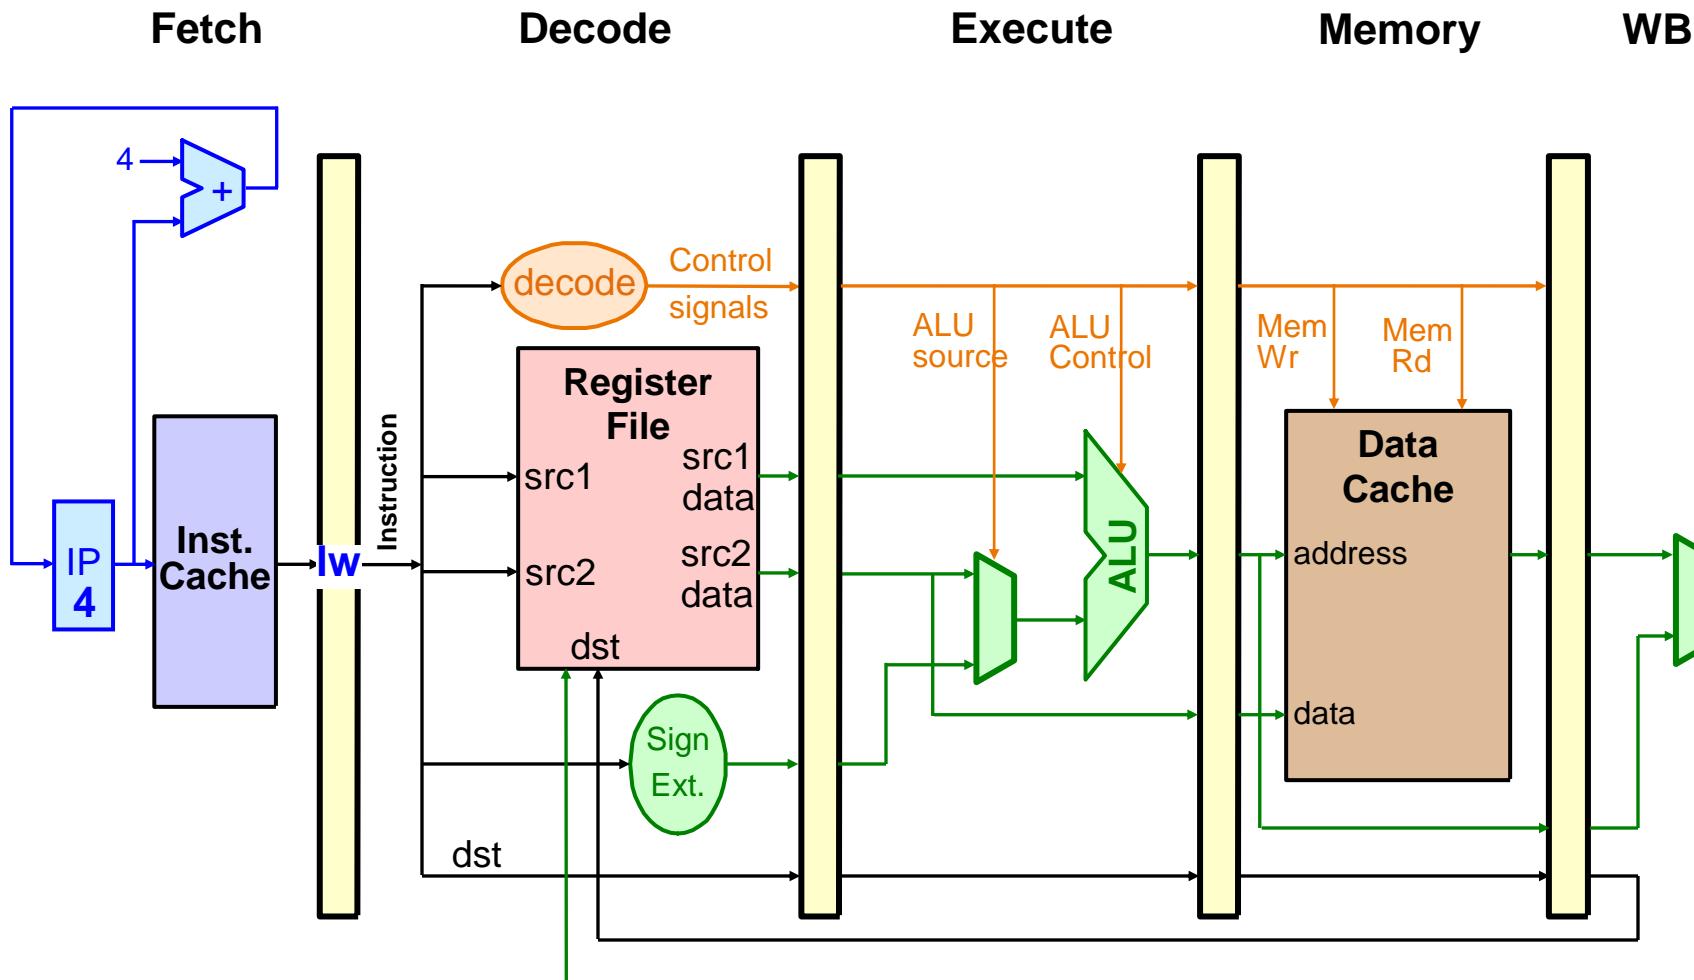


```

0  lw  R10,9(R1)
4  sub R11,R2,R3
8  and R12,R4,R5
12 or  R13,R6,R7

```

# Pipeline Example: cycle 2

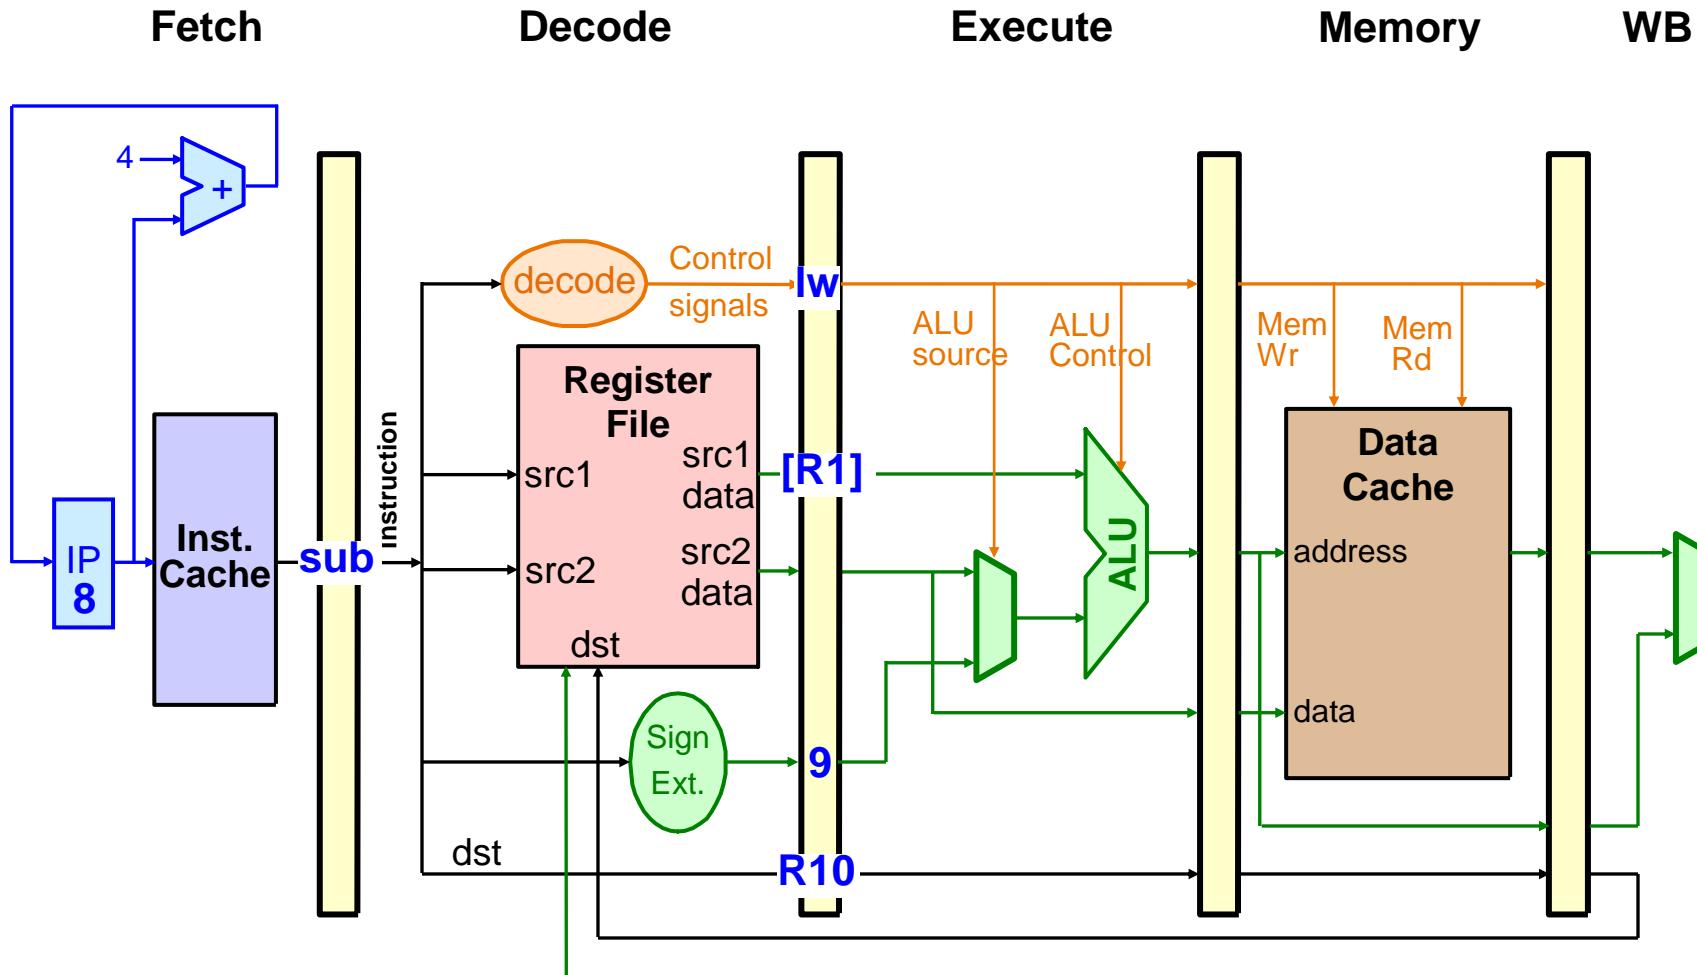


```

0  lw  R10,9(R1)
4  sub R11,R2,R3
8  and R12,R4,R5
12 or  R13,R6,R7

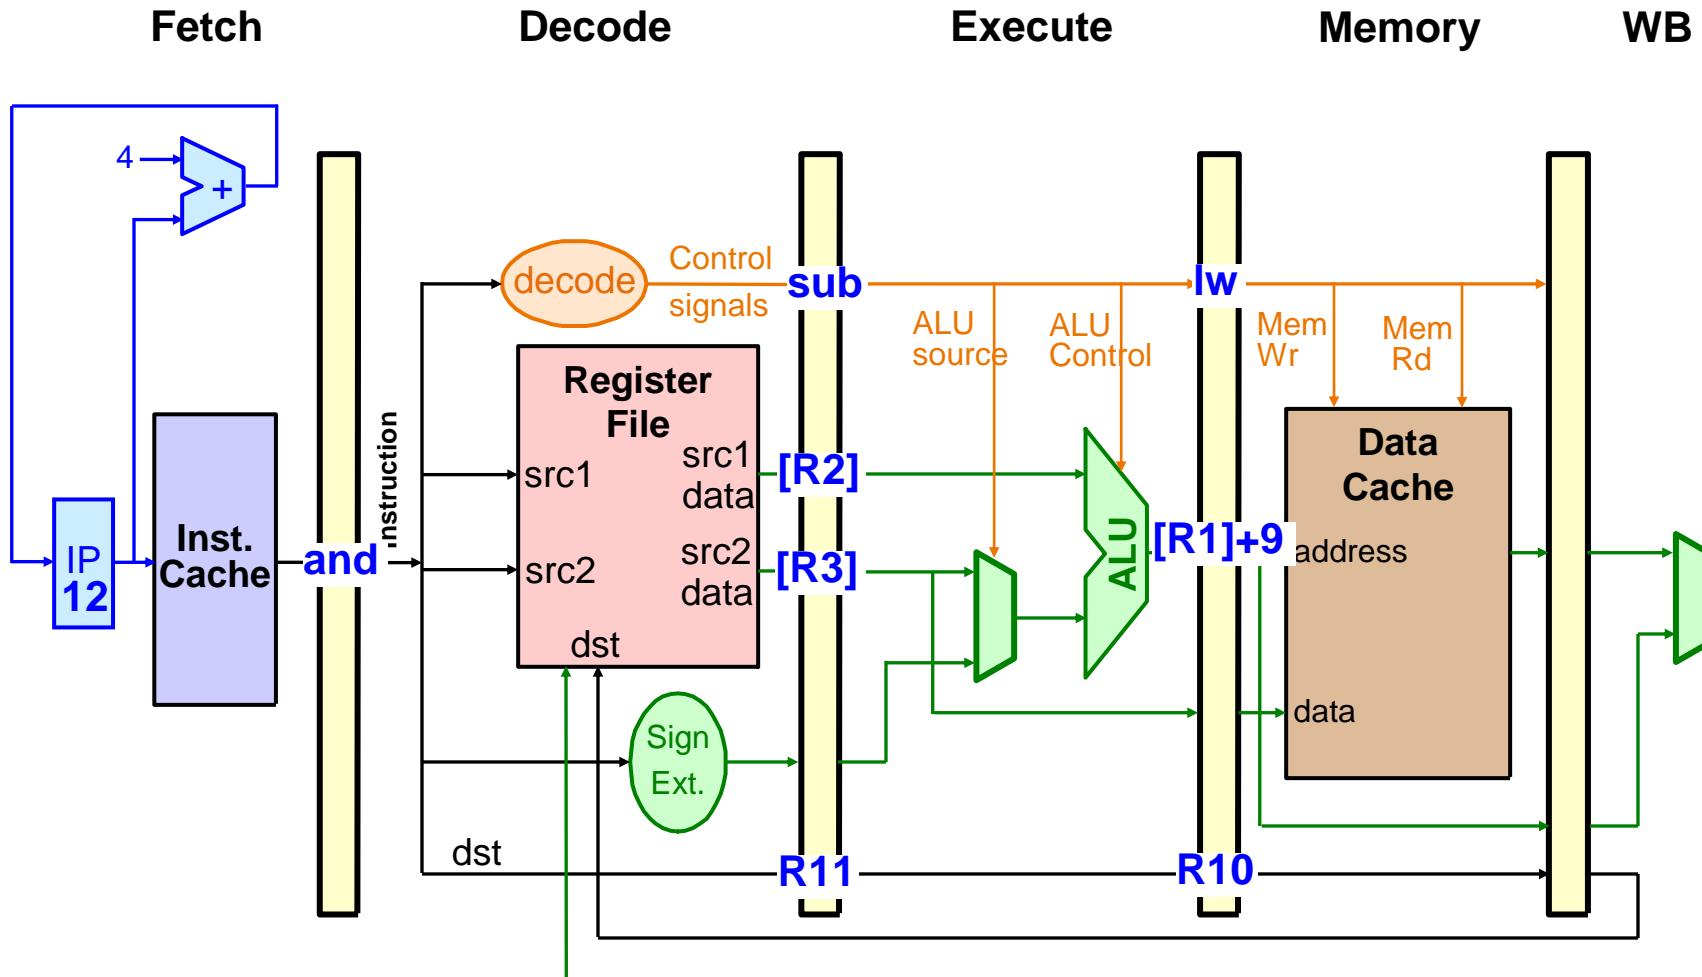
```

# Pipeline Example: cycle 3



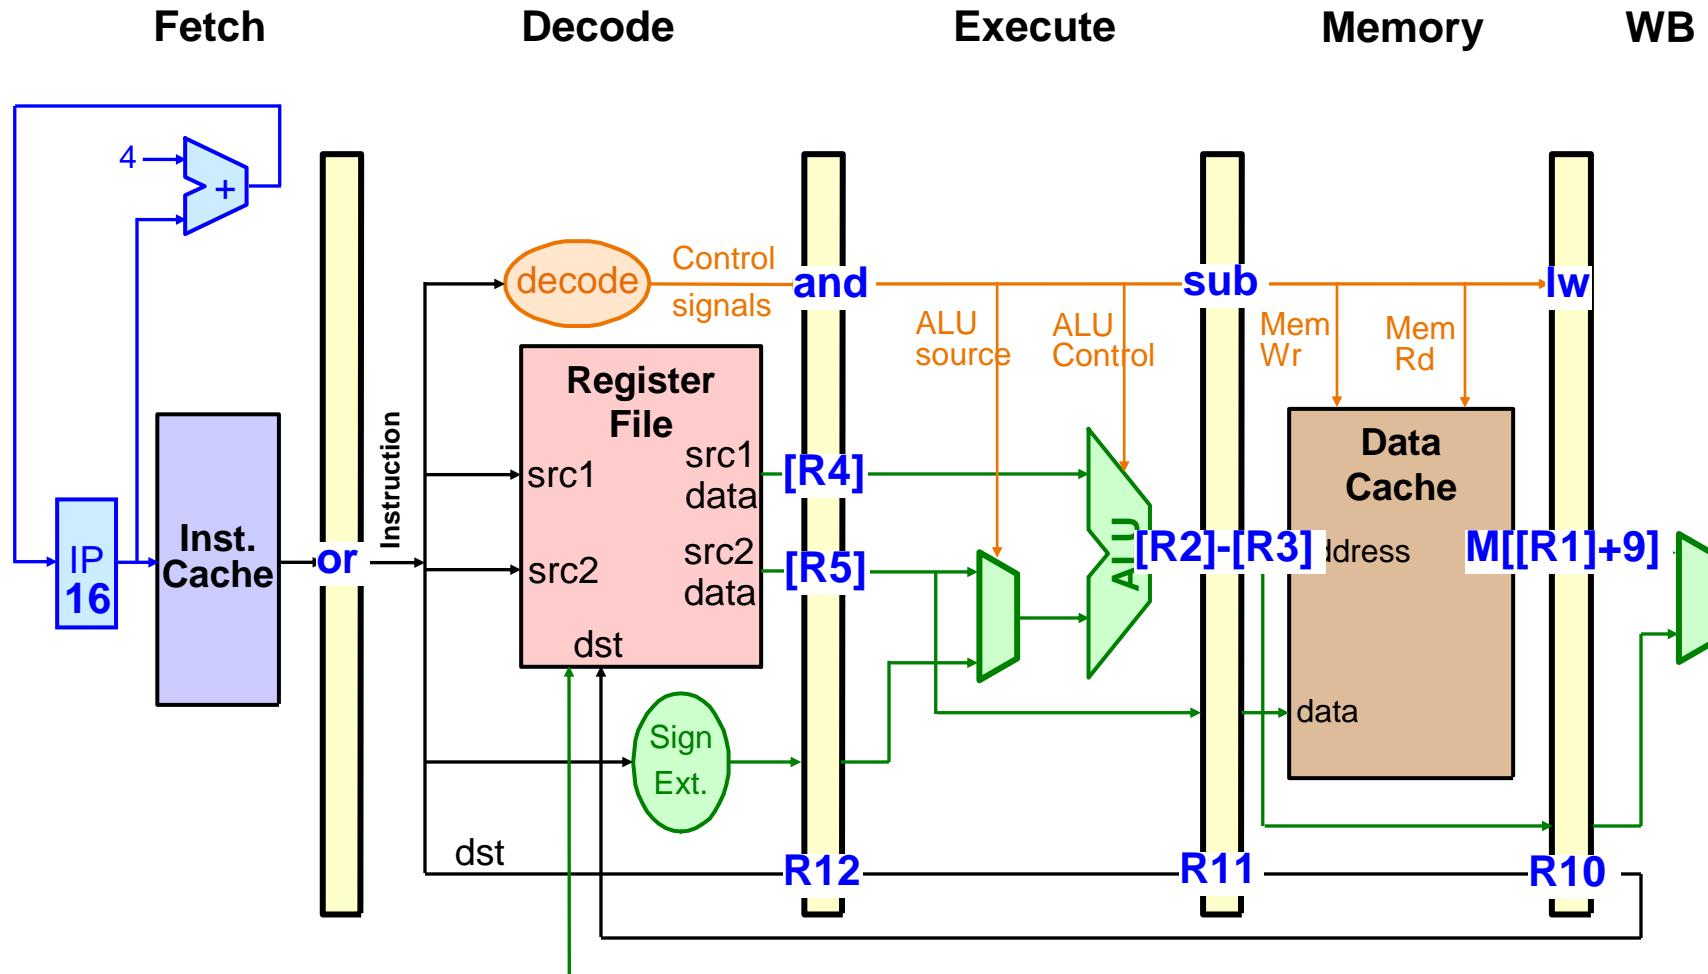
0	<b>lw</b>	R10,9(R1)
4	<b>sub</b>	R11,R2,R3
8	<b>and</b>	R12,R4,R5
12	<b>or</b>	R13,R6,R7

# Pipeline Example: cycle 4



0	lw	R10, 9(R1)
4	sub	R11, R2, R3
8	and	R12, R4, R5
12	or	R13, R6, R7

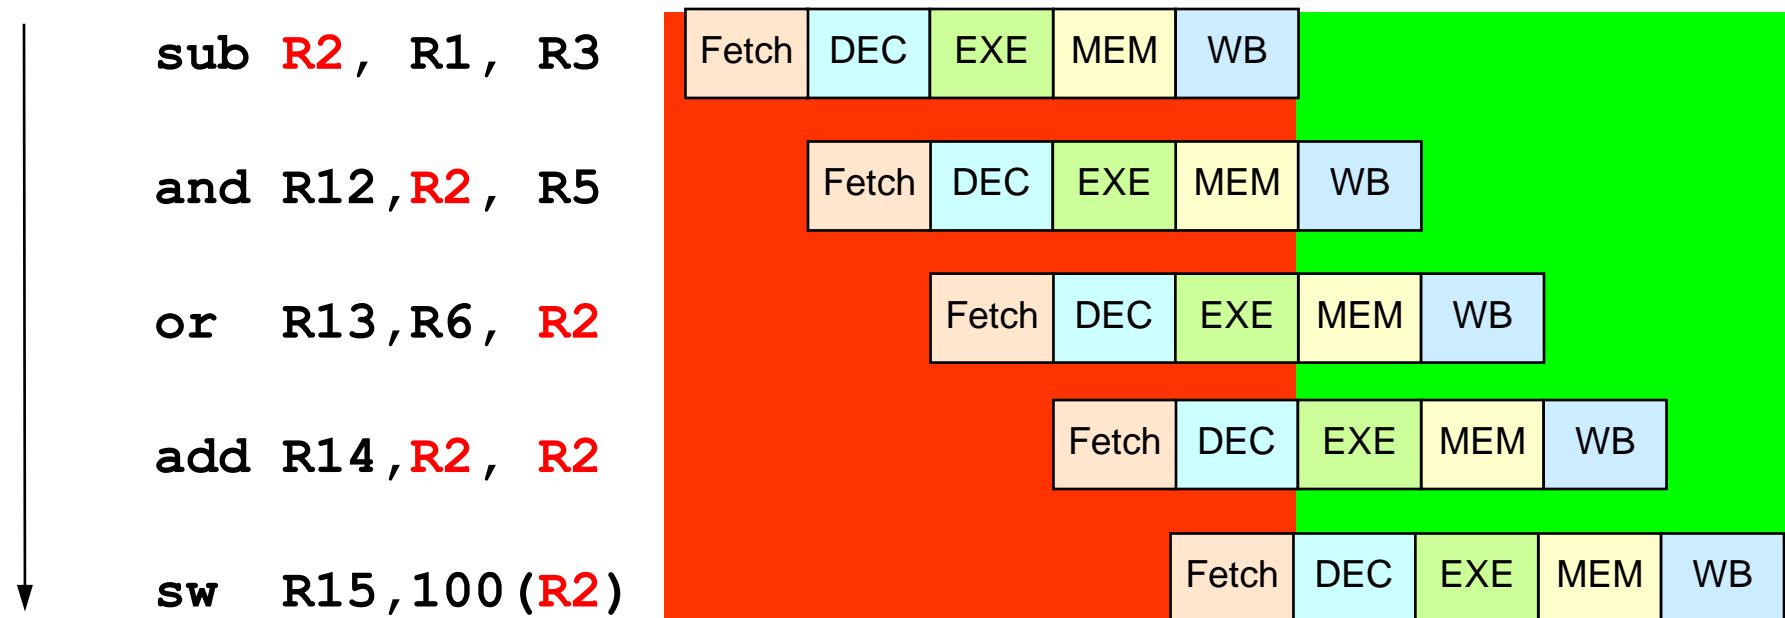
# Pipeline Example: cycle 5



0	lw	R10, 9(R1)
4	sub	R11, R2, R3
8	and	R12, R4, R5
12	or	R13, R6, R7

# RAW Dependency

Program  
execution  
order

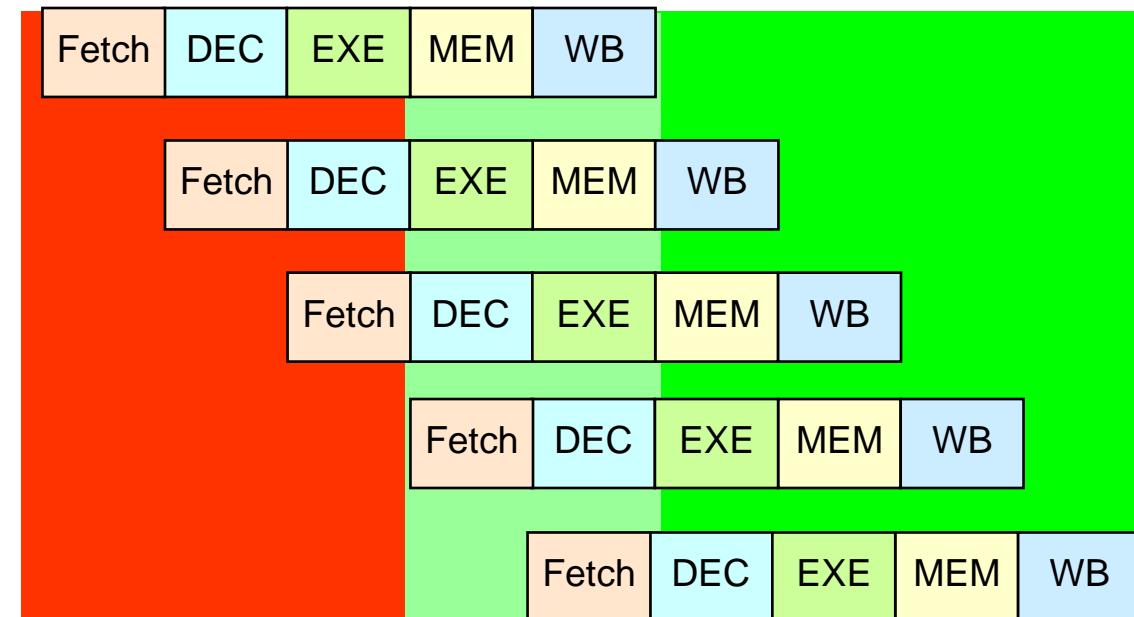


# Using Bypass to Solve RAW Dependency

Program  
execution  
order

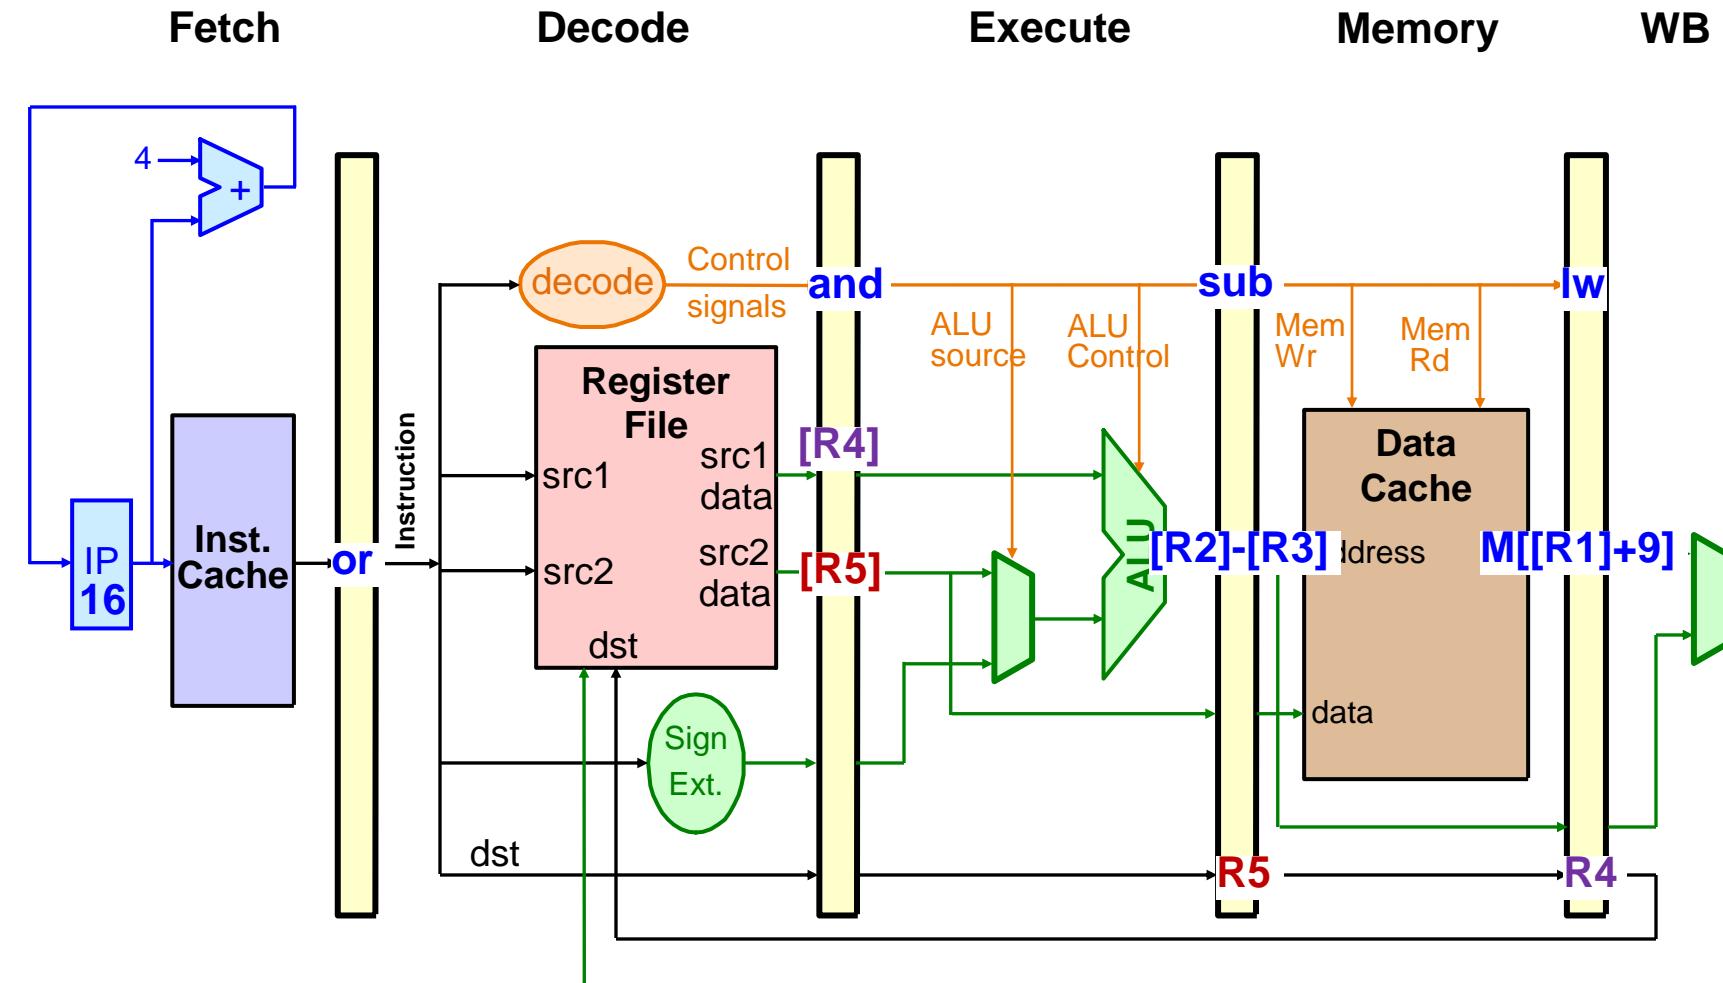
↓

```
sub R2, R1, R3
and R12, R2, R5
or R13, R6, R2
add R14, R2, R2
sw R15, 100(R2)
```



Bypass result directly from EXE output to EXE input

# RAW Dependency

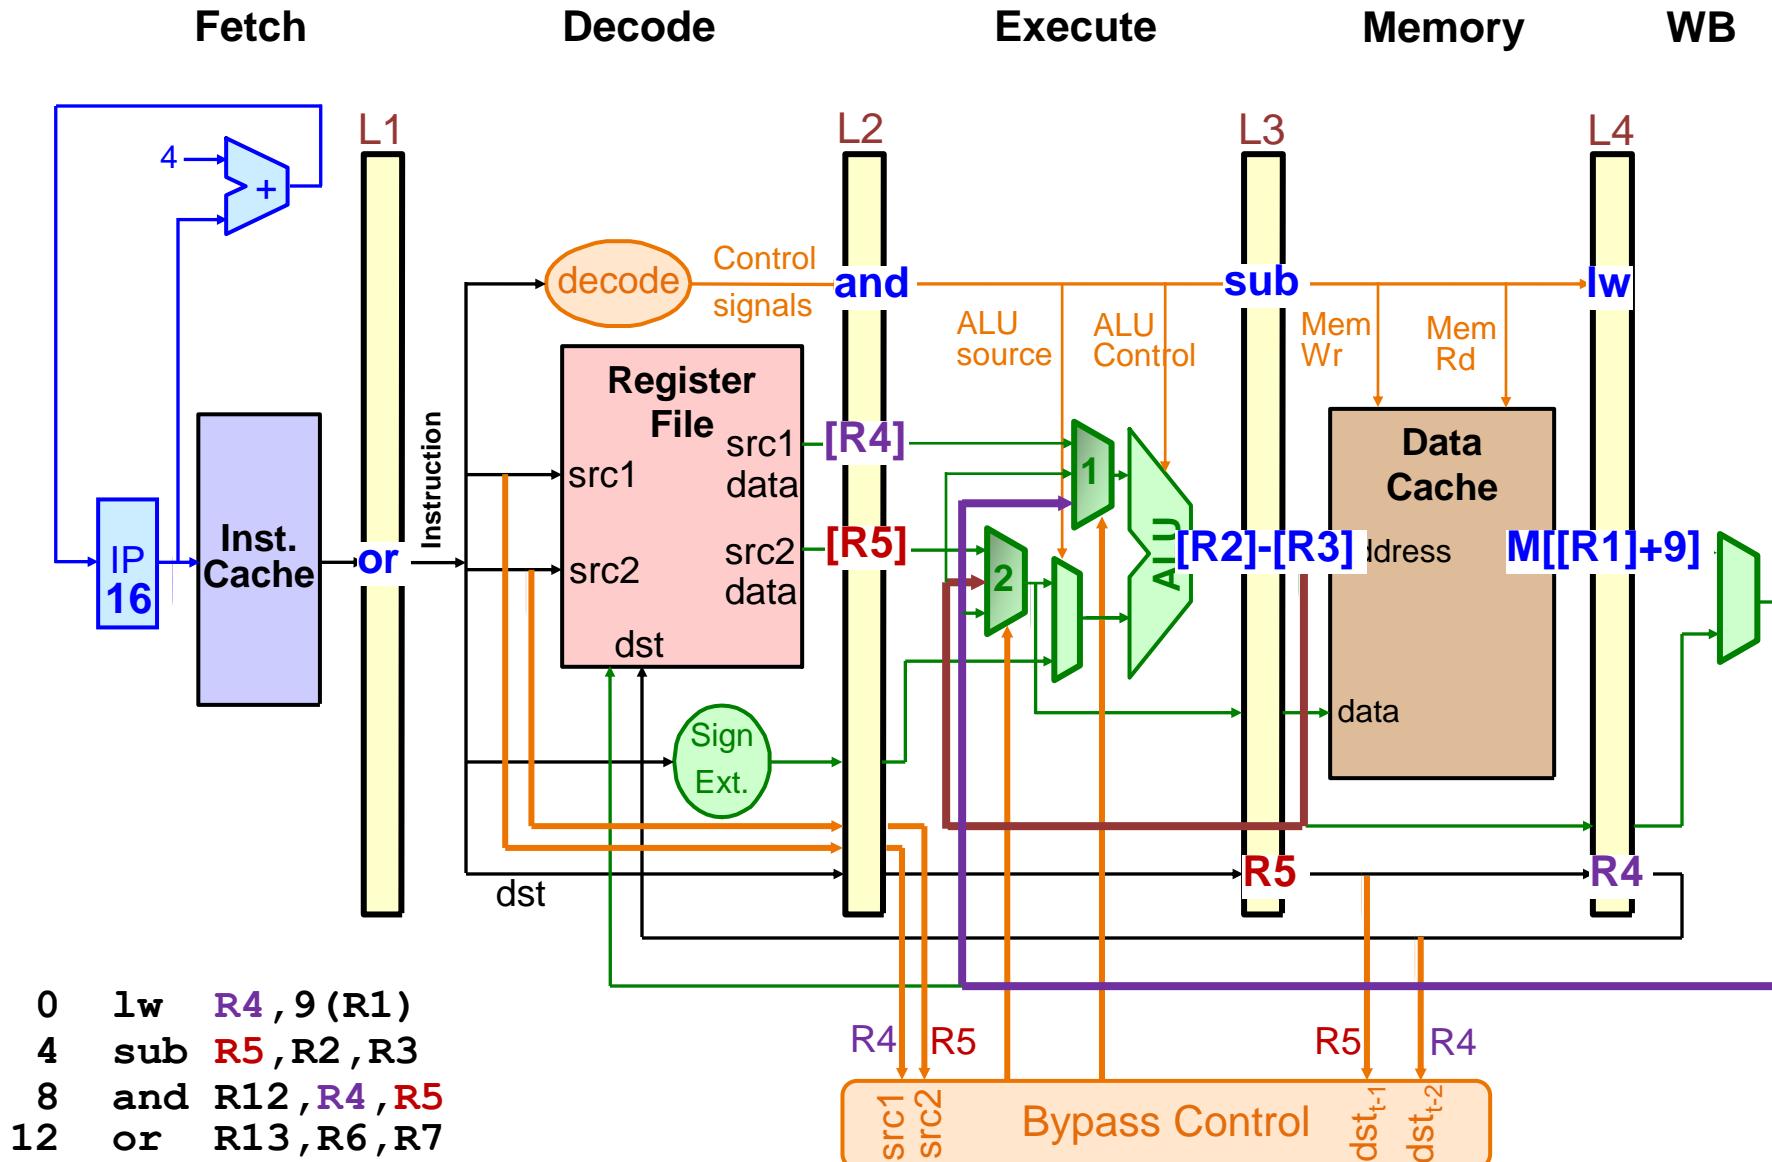


```

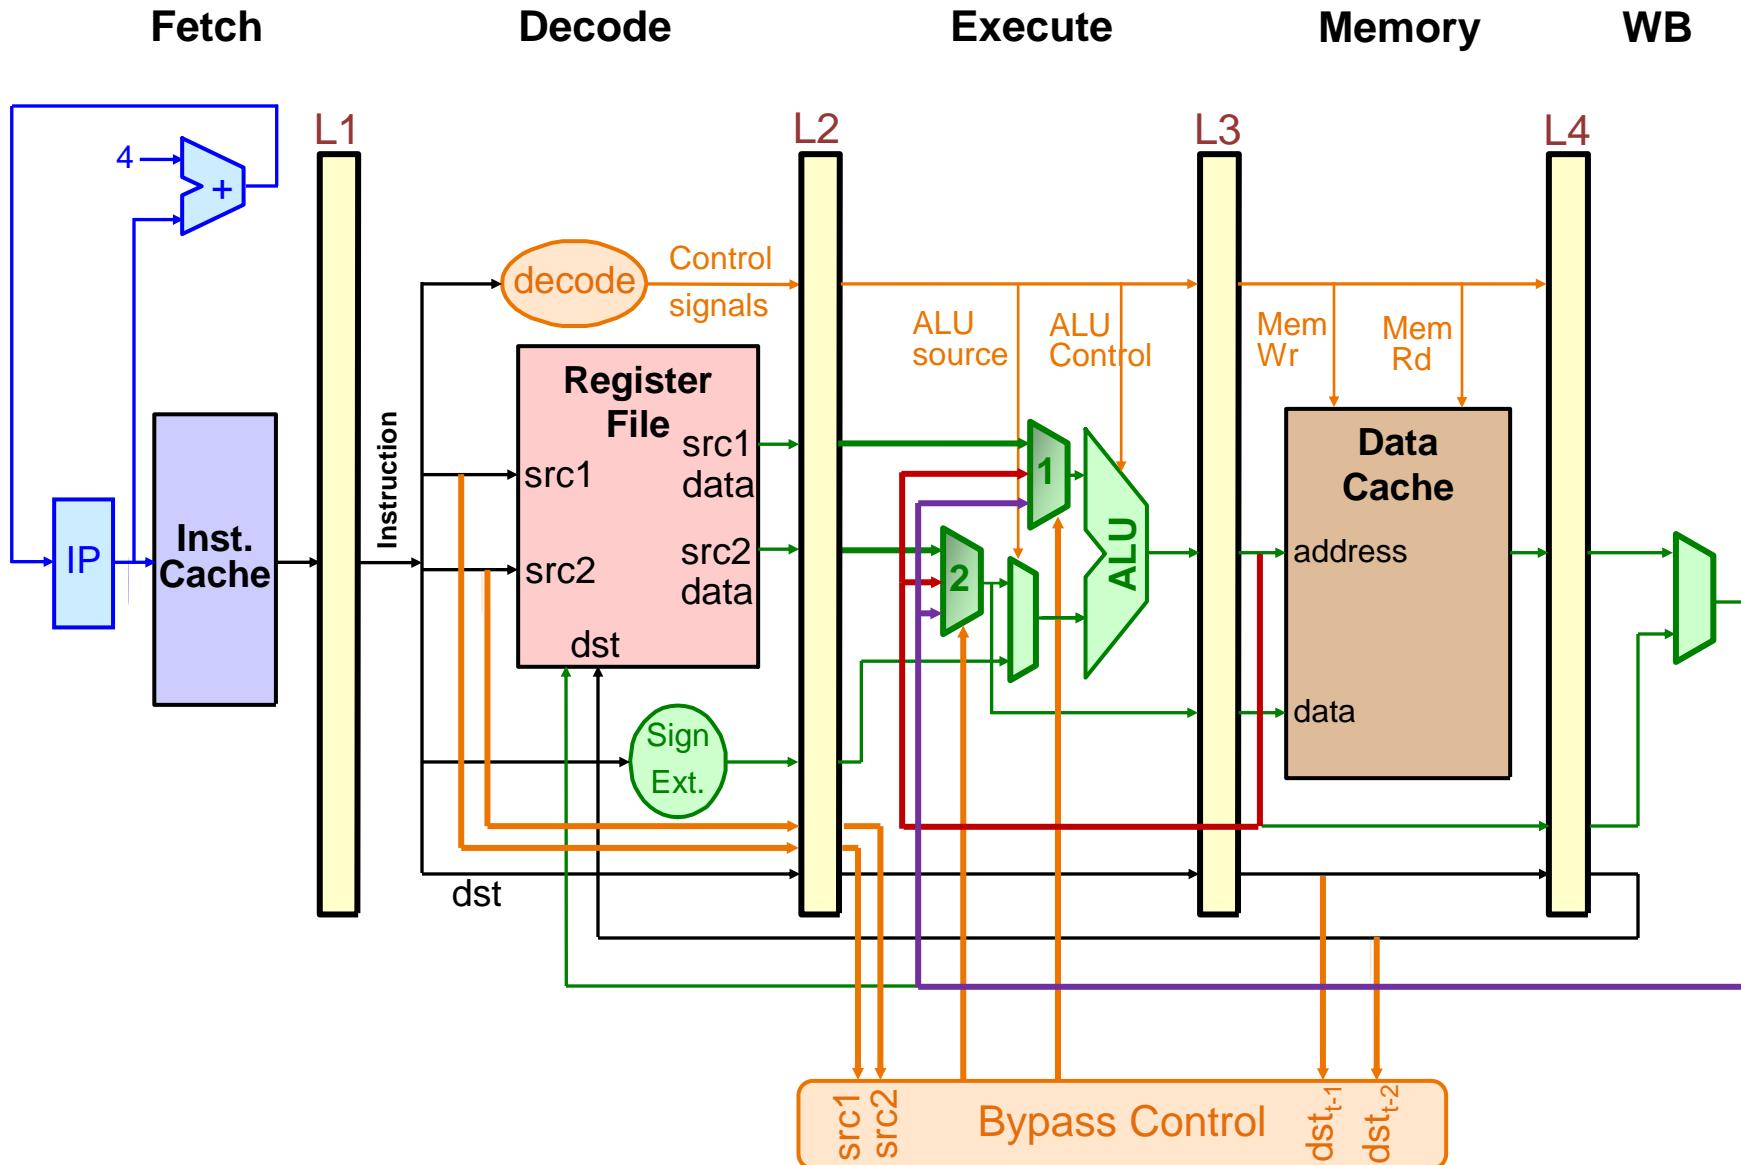
0  lw   R4 , 9 (R1)
4  sub  R5 , R2 , R3
8  and  R12 , R4 , R5
12 or   R13 , R6 , R7

```

# Bypass Hardware



# Bypass Hardware



# Bypass Control

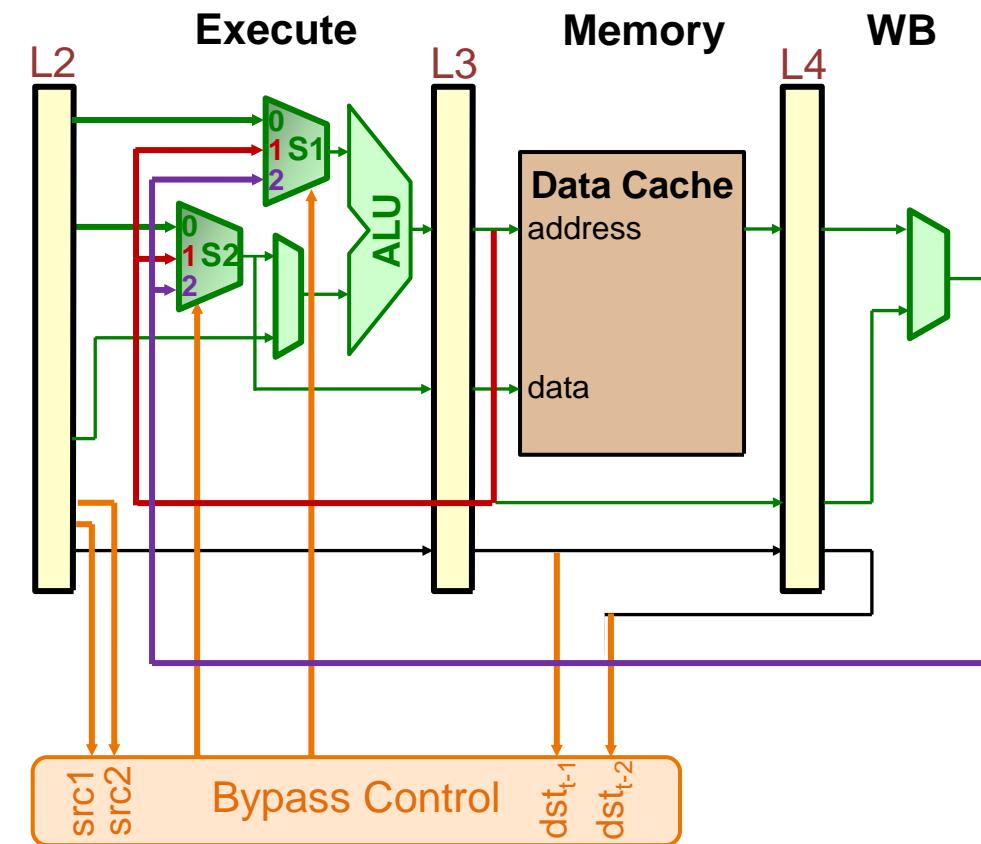
- ◆  $L3\_reg\_wr = L3.\text{RegWrite}$  and ( $L3.\text{opcode} \neq \text{lw}$ )

## ◆ Forwarding from EXE (L3)

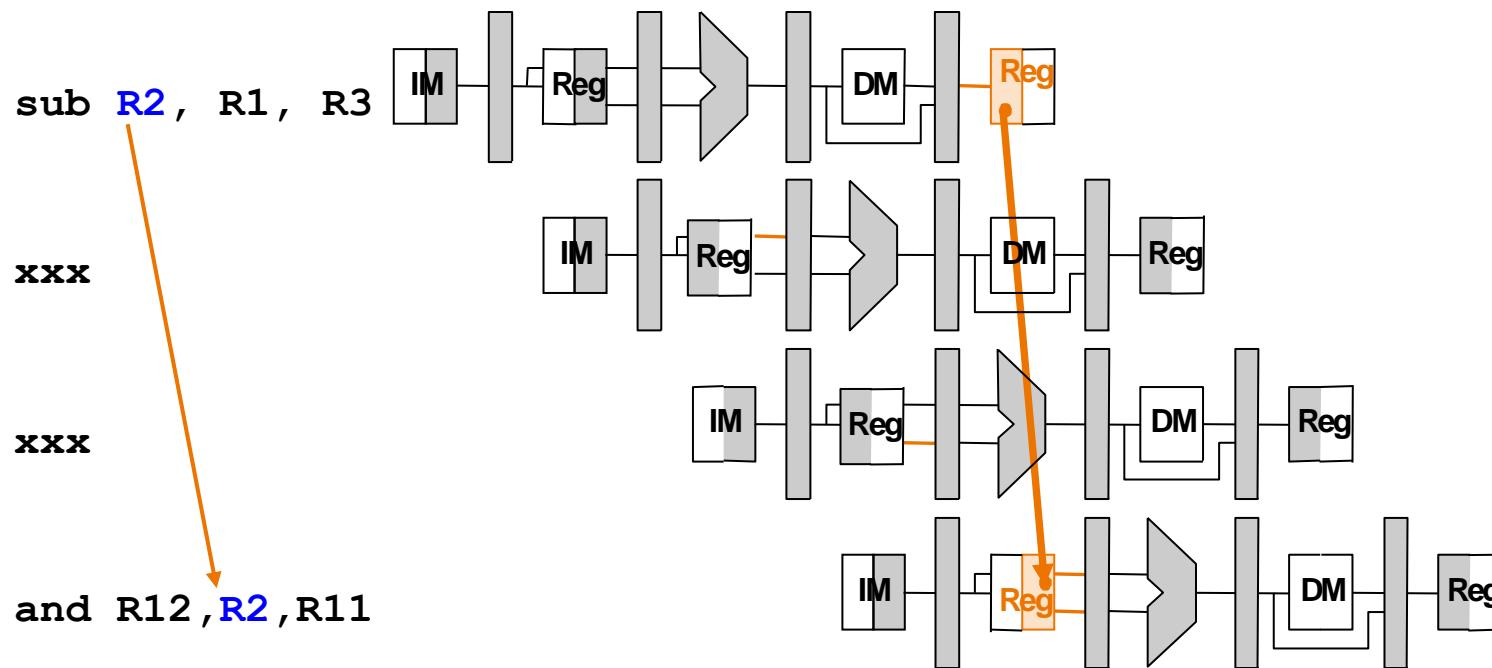
- ❖ if ( $L3.\text{reg\_wr}$  and ( $L3.\text{dst} == L2.\text{src1}$ ))  
 $\text{ALUSelA} = 1$
- ❖ if ( $L3.\text{reg\_wr}$  and ( $L3.\text{dst} == L2.\text{src2}$ ))  
 $\text{ALUSelB} = 1$

## ◆ Forwarding from MEM (L4)

- ❖ if ( $L4.\text{RegWrite}$  and  
((not  $L3.\text{reg\_wr}$ ) or ( $L3.\text{dst} \neq L2.\text{src1}$ )) and  
( $L4.\text{dst} = L2.\text{src1}$ )  $\text{ALUSelA} = 2$
- ❖ if ( $L4.\text{RegWrite}$  and  
((not  $L3.\text{reg\_wr}$ ) or ( $L3.\text{dst} \neq L2.\text{src2}$ )) and  
( $L4.\text{dst} = L2.\text{src2}$ )  $\text{ALUSelB} = 2$



# Register File Split

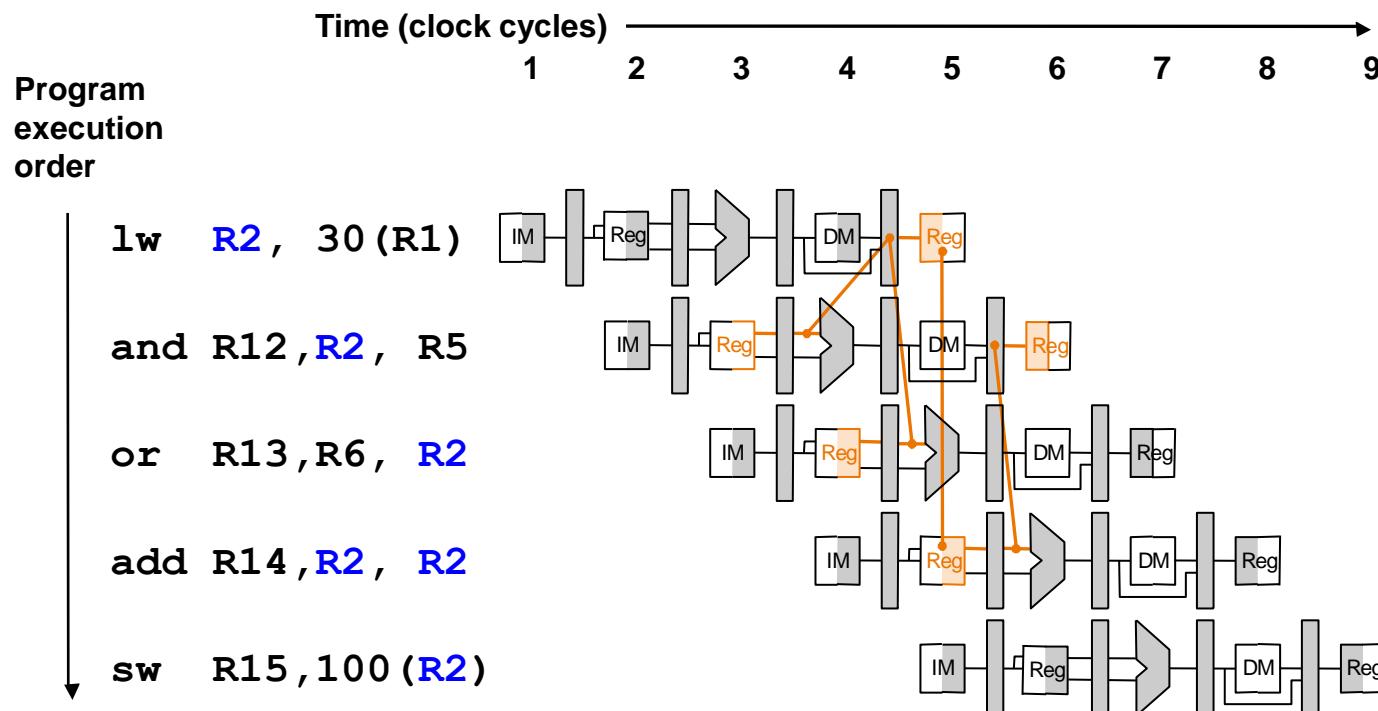


- ◆ Register file is written during first half of the cycle
- ◆ Register file is read during second half of the cycle
- ⇒ Register file is written before it is read ⇒ returns the correct data

# Can't Always Bypass

- ◆ Load word can still cause a hazard

- ❖ an instruction tries to read a register following a load instruction that writes to the same register

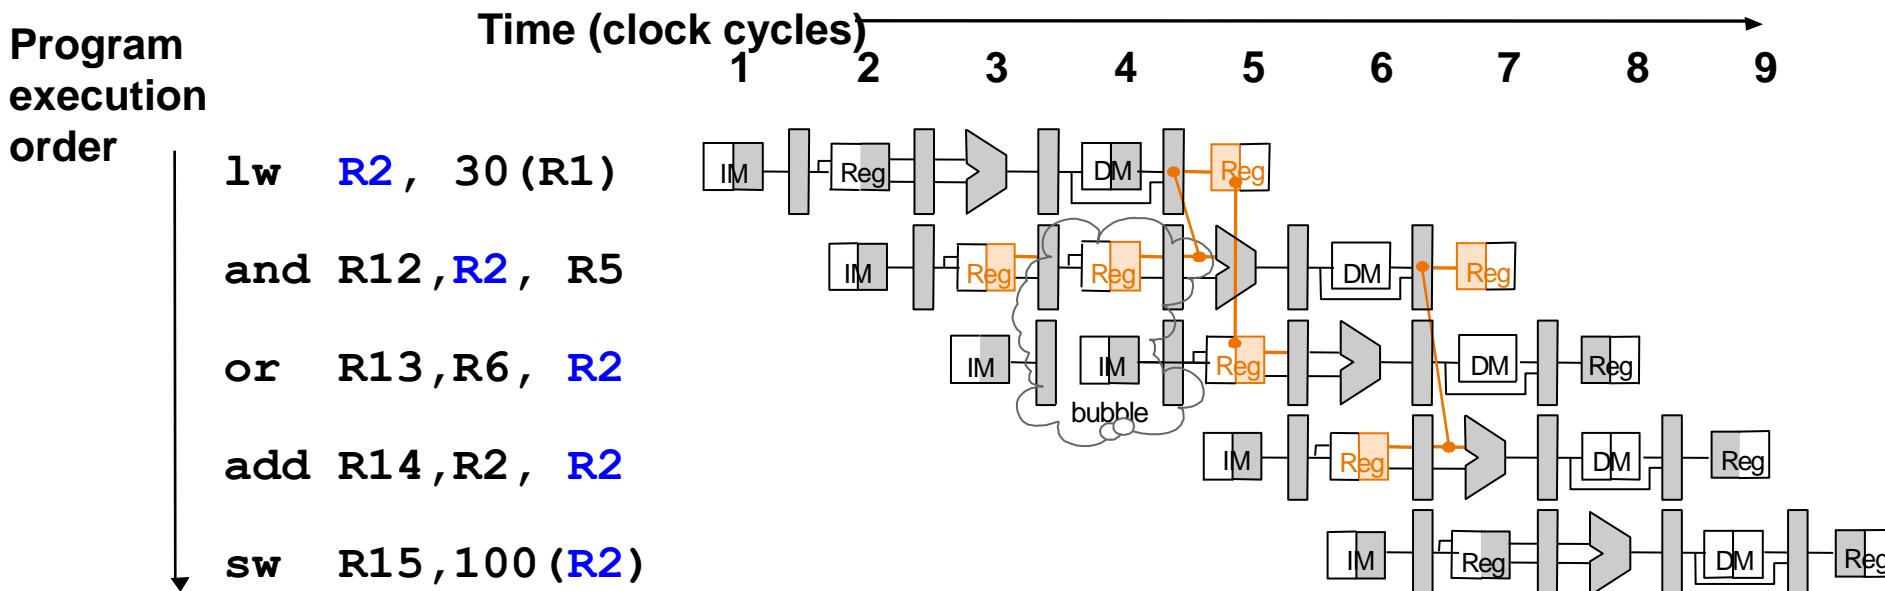


⇒ A hazard detection unit is needed to “stall” the load instruction

# Stall If Cannot Bypass

if (L2.RegWrite and (L2.opcode == lw) and ( (L2.dst == L1.src1) or (L2.dst == L1.src2) ) then stall

- ◆ De-assert the enable to the L1 latch, and to the IP
  - ❖ The dependent instruction (**and**) stays another cycle in L1
- ◆ Issue a NOP into the L2 latch (instead of the stalled inst.)
- ◆ Allow the stalling instruction (**lw**) to move on



# Software Scheduling to Avoid Load Hazards

Example: code for (assume all variables are in memory):

```
a = b + c;
```

```
d = e - f;
```

## Slow code

```
LW Rb,b
```

```
LW Rc,c
```

**Stall** ADD Ra,Rb,Rc

```
SW a,Ra
```

```
LW Re,e
```

```
LW Rf,f
```

**Stall** SUB Rd,Re,Rf

```
SW d,Rd
```

## Fast code

```
LW Rb,b
```

```
LW Rc,c
```

LW Re,e

ADD Ra,Rb,Rc

```
LW Rf,f
```

SW a,Ra

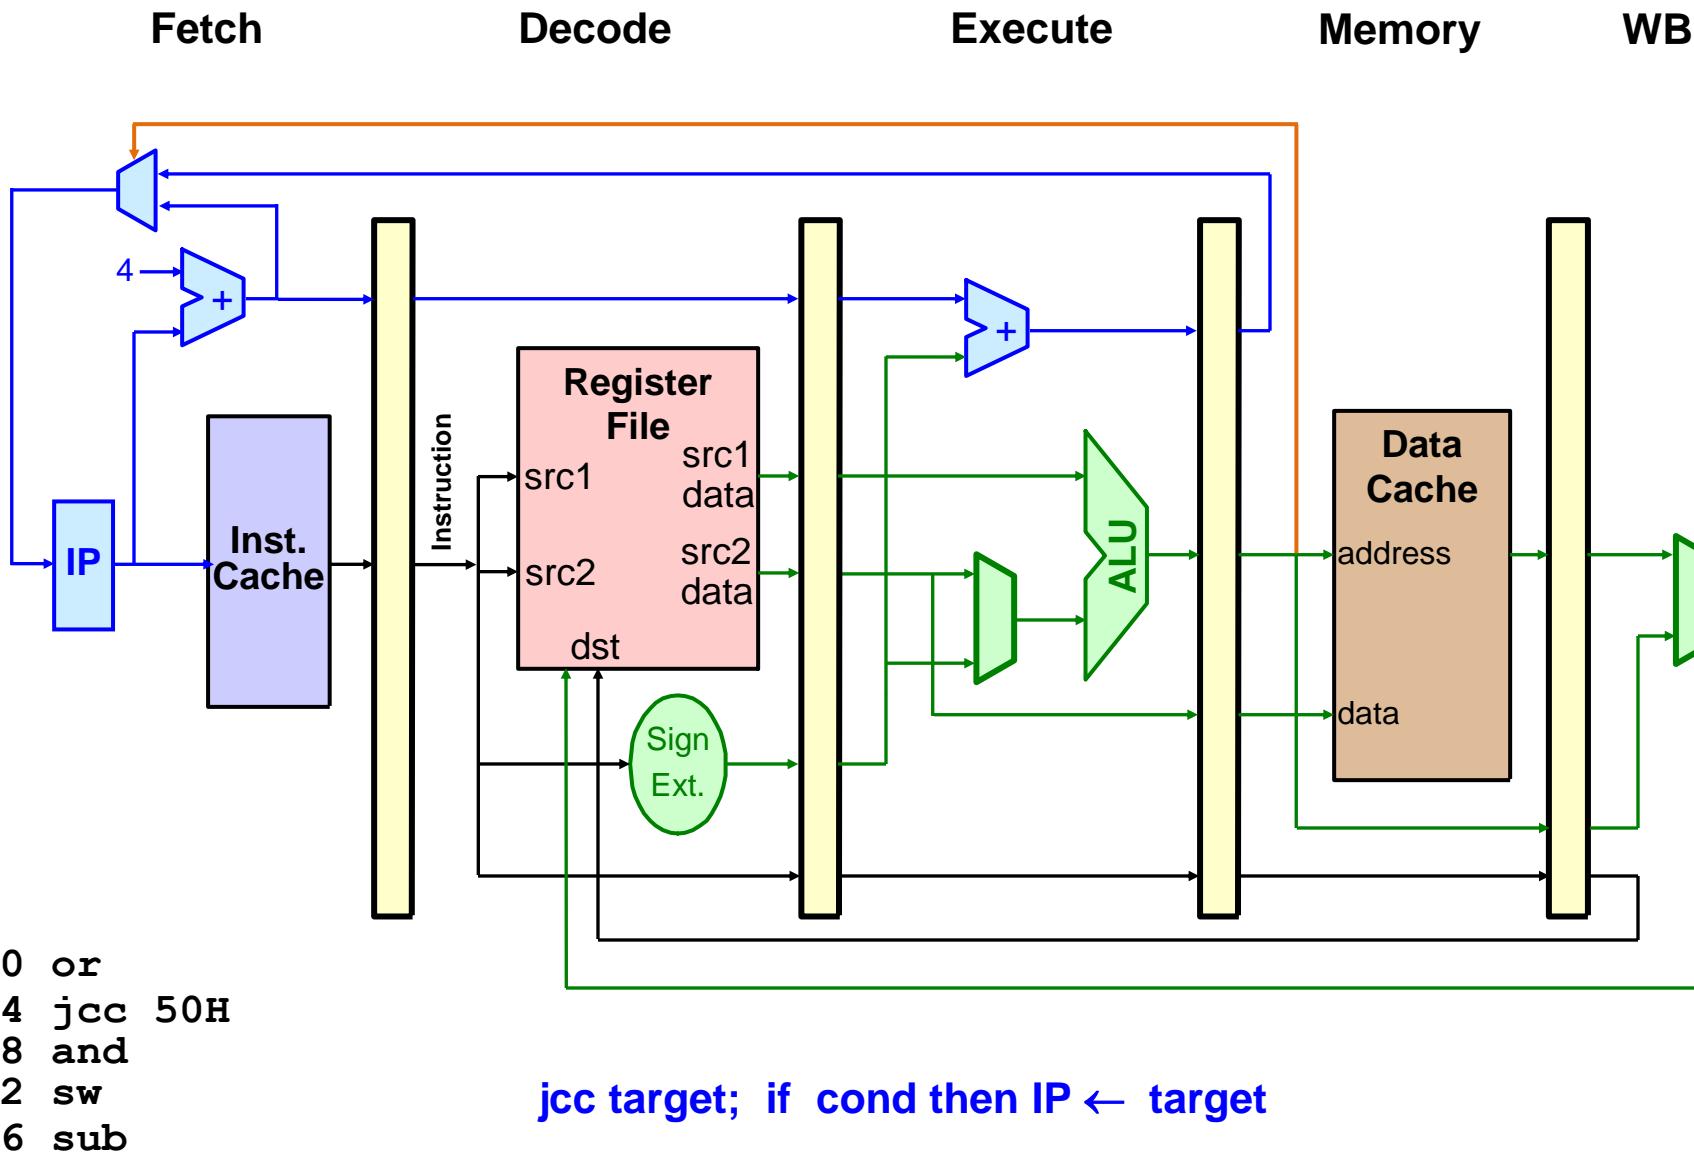
SUB Rd,Re,Rf

```
SW d,Rd
```

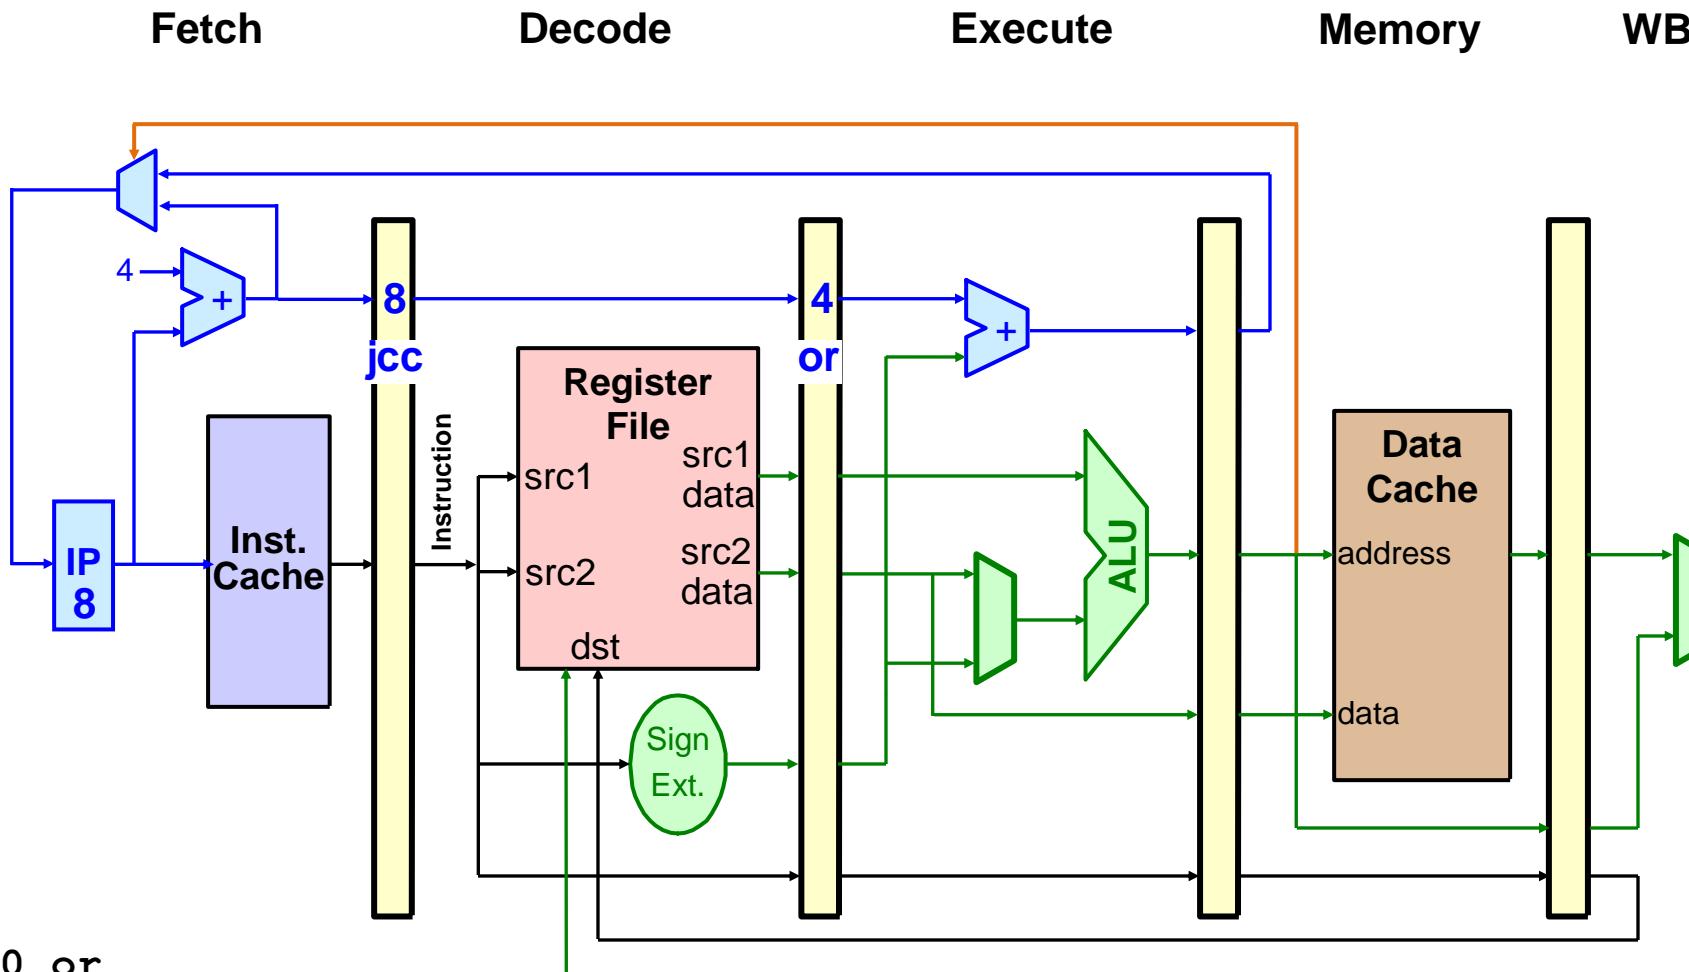
Instruction order can be changed as long as the correctness is kept

# Control Hazards

# Control Hazard on Branches

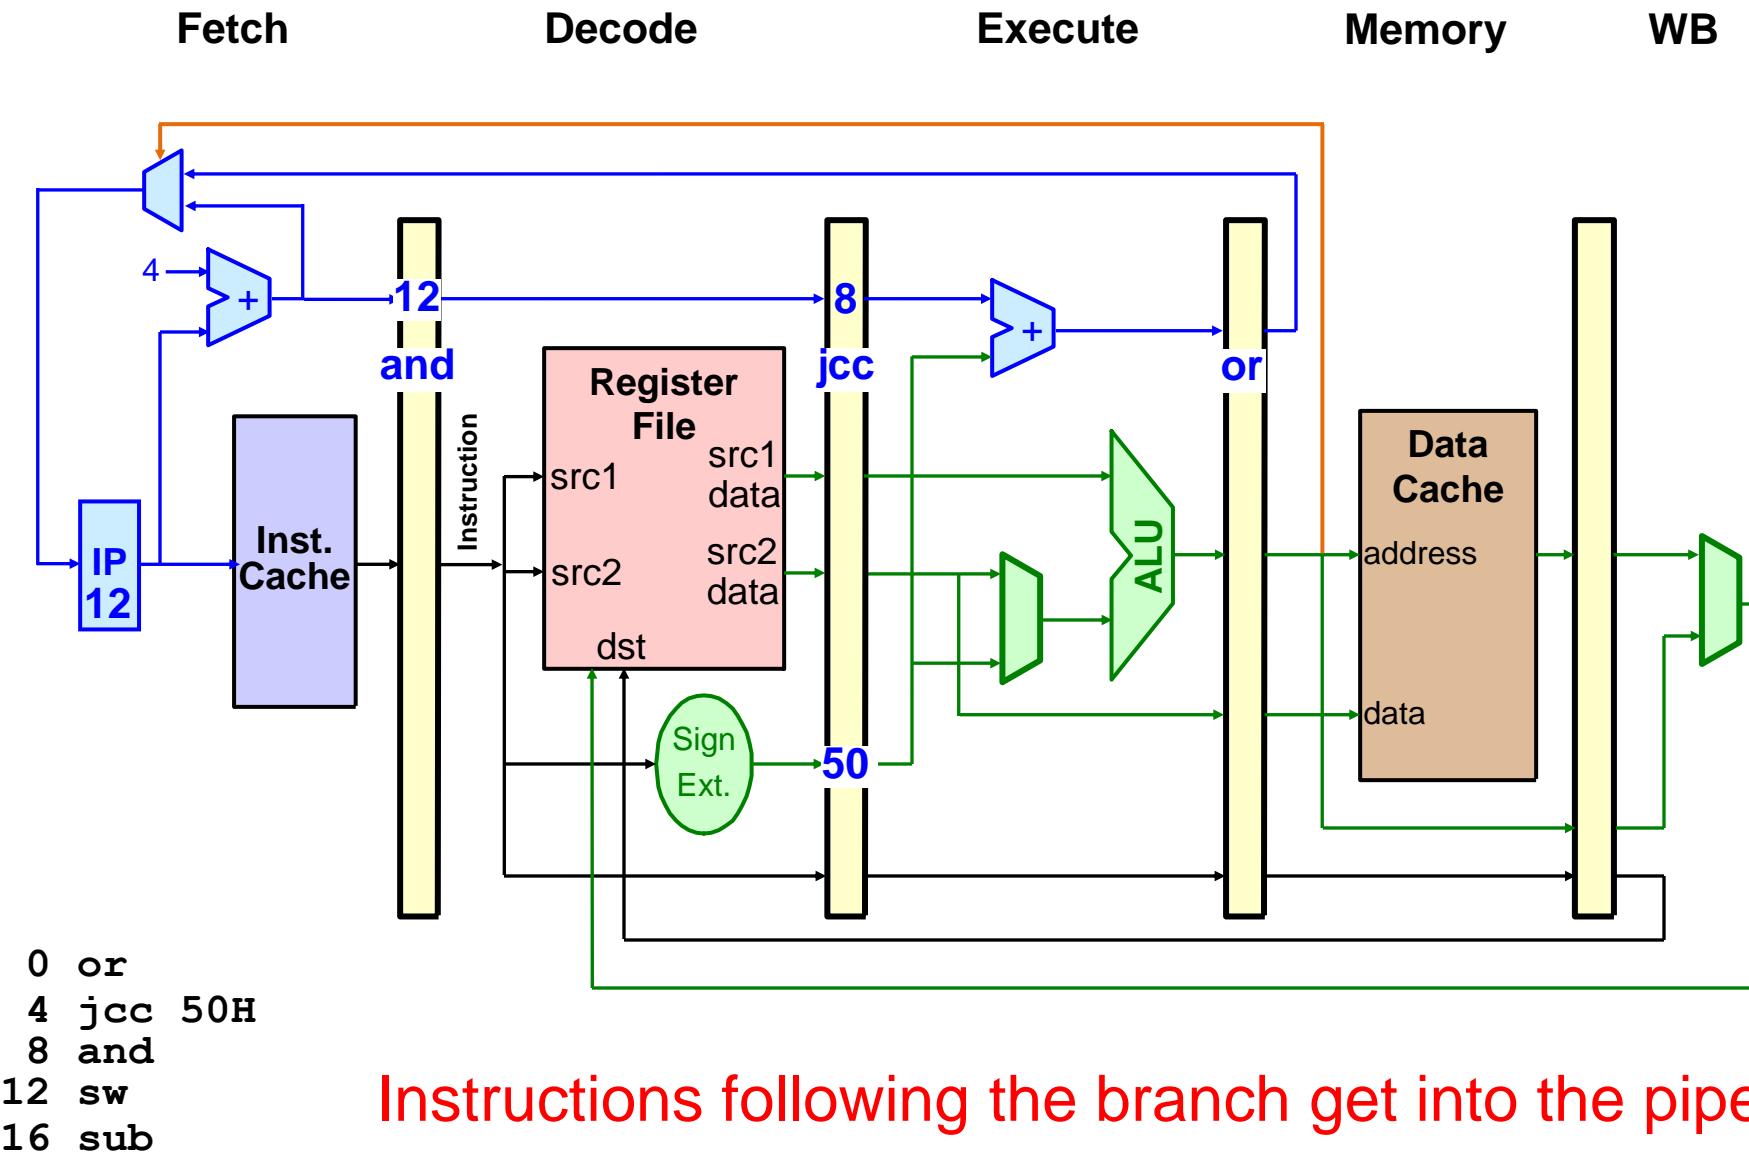


# Control Hazard on Branches

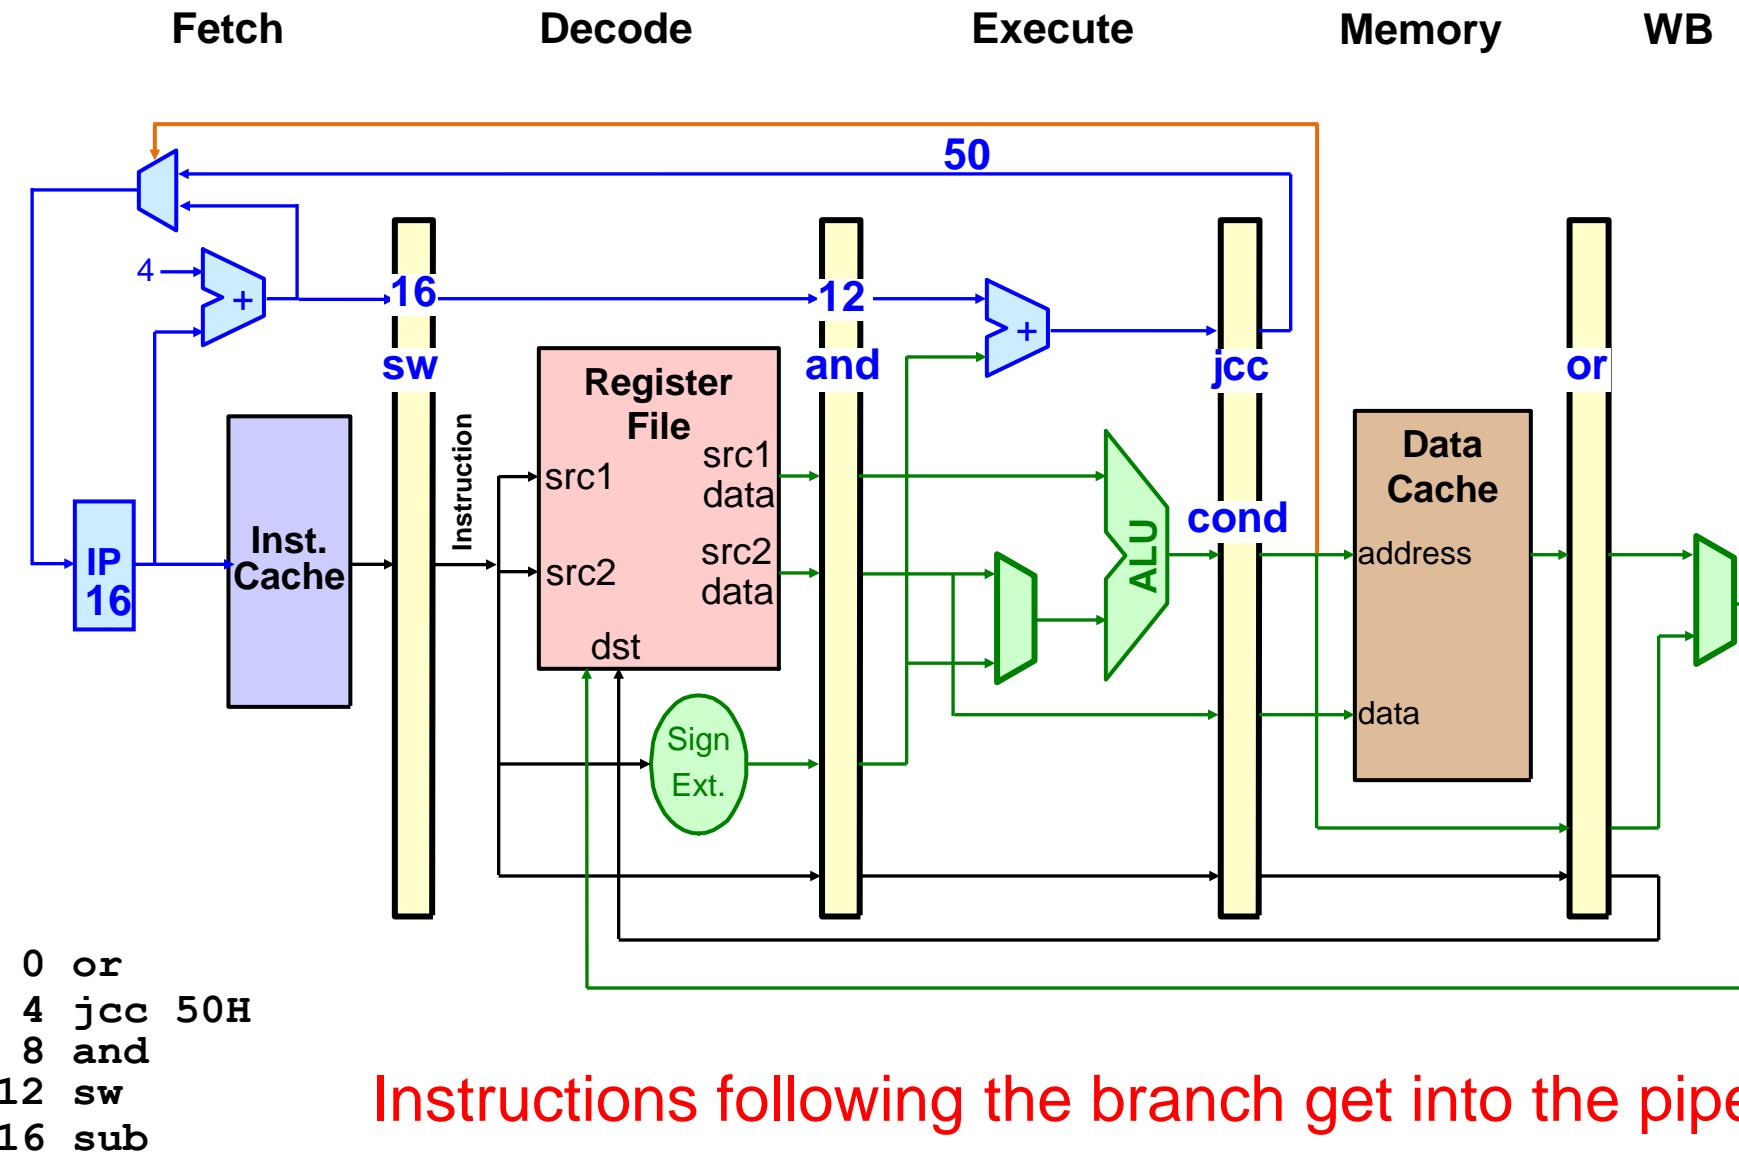


0 or  
 4 jcc 50H  
 8 and  
 12 sw  
 16 sub

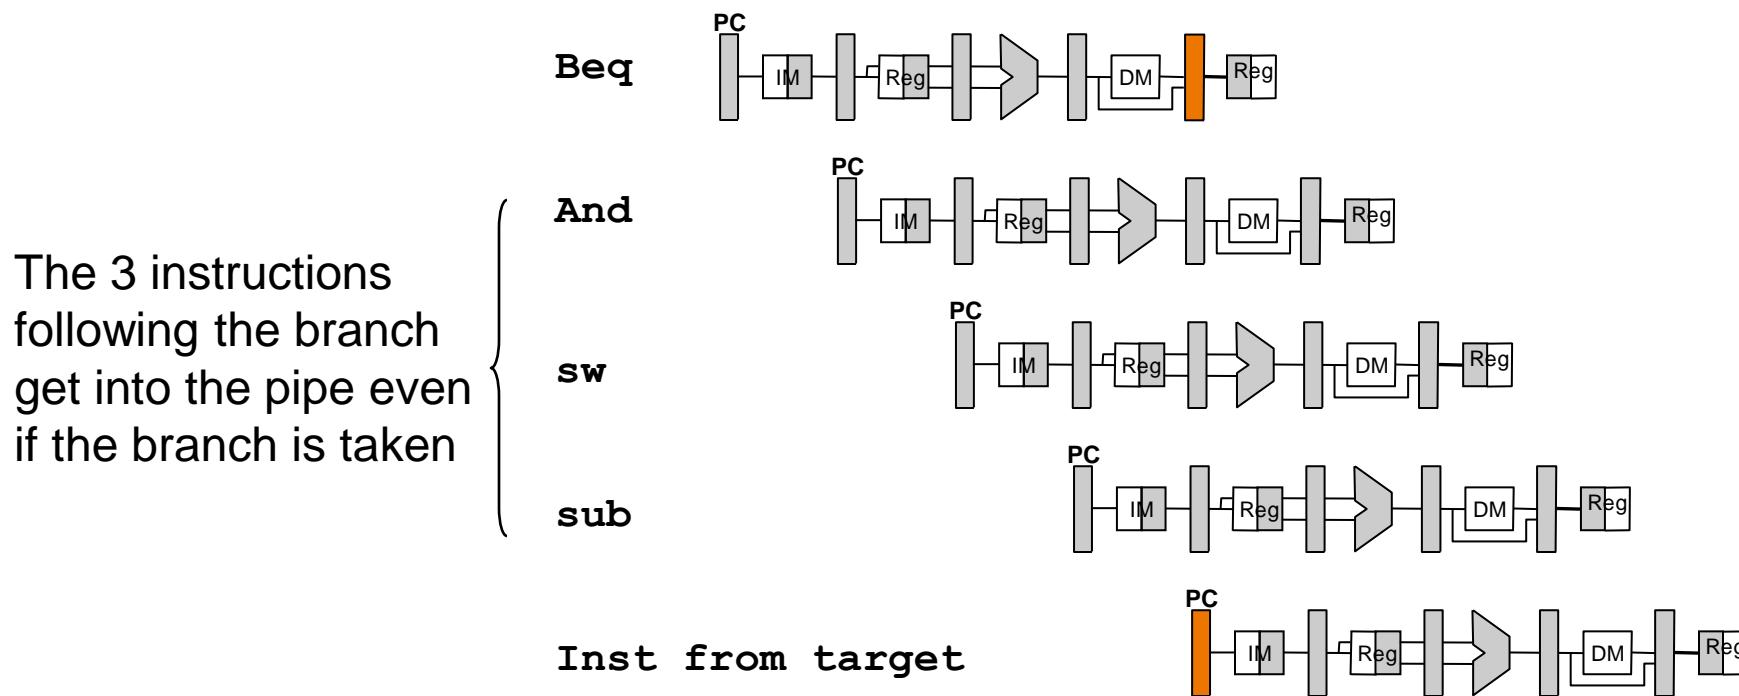
# Control Hazard on Branches



# Control Hazard on Branches



# Control Hazard on Branches



# Control Hazard: Stall

- ◆ Stall pipe when branch is encountered until resolved
  - ◆ Stall impact: assumptions
    - ❖ CPI = 1
    - ❖ 20% of instructions are branches
    - ❖ Stall 3 cycles on every branch
- ⇒  $CPI_{new} = 1 + 0.2 \times 3 = 1.6$

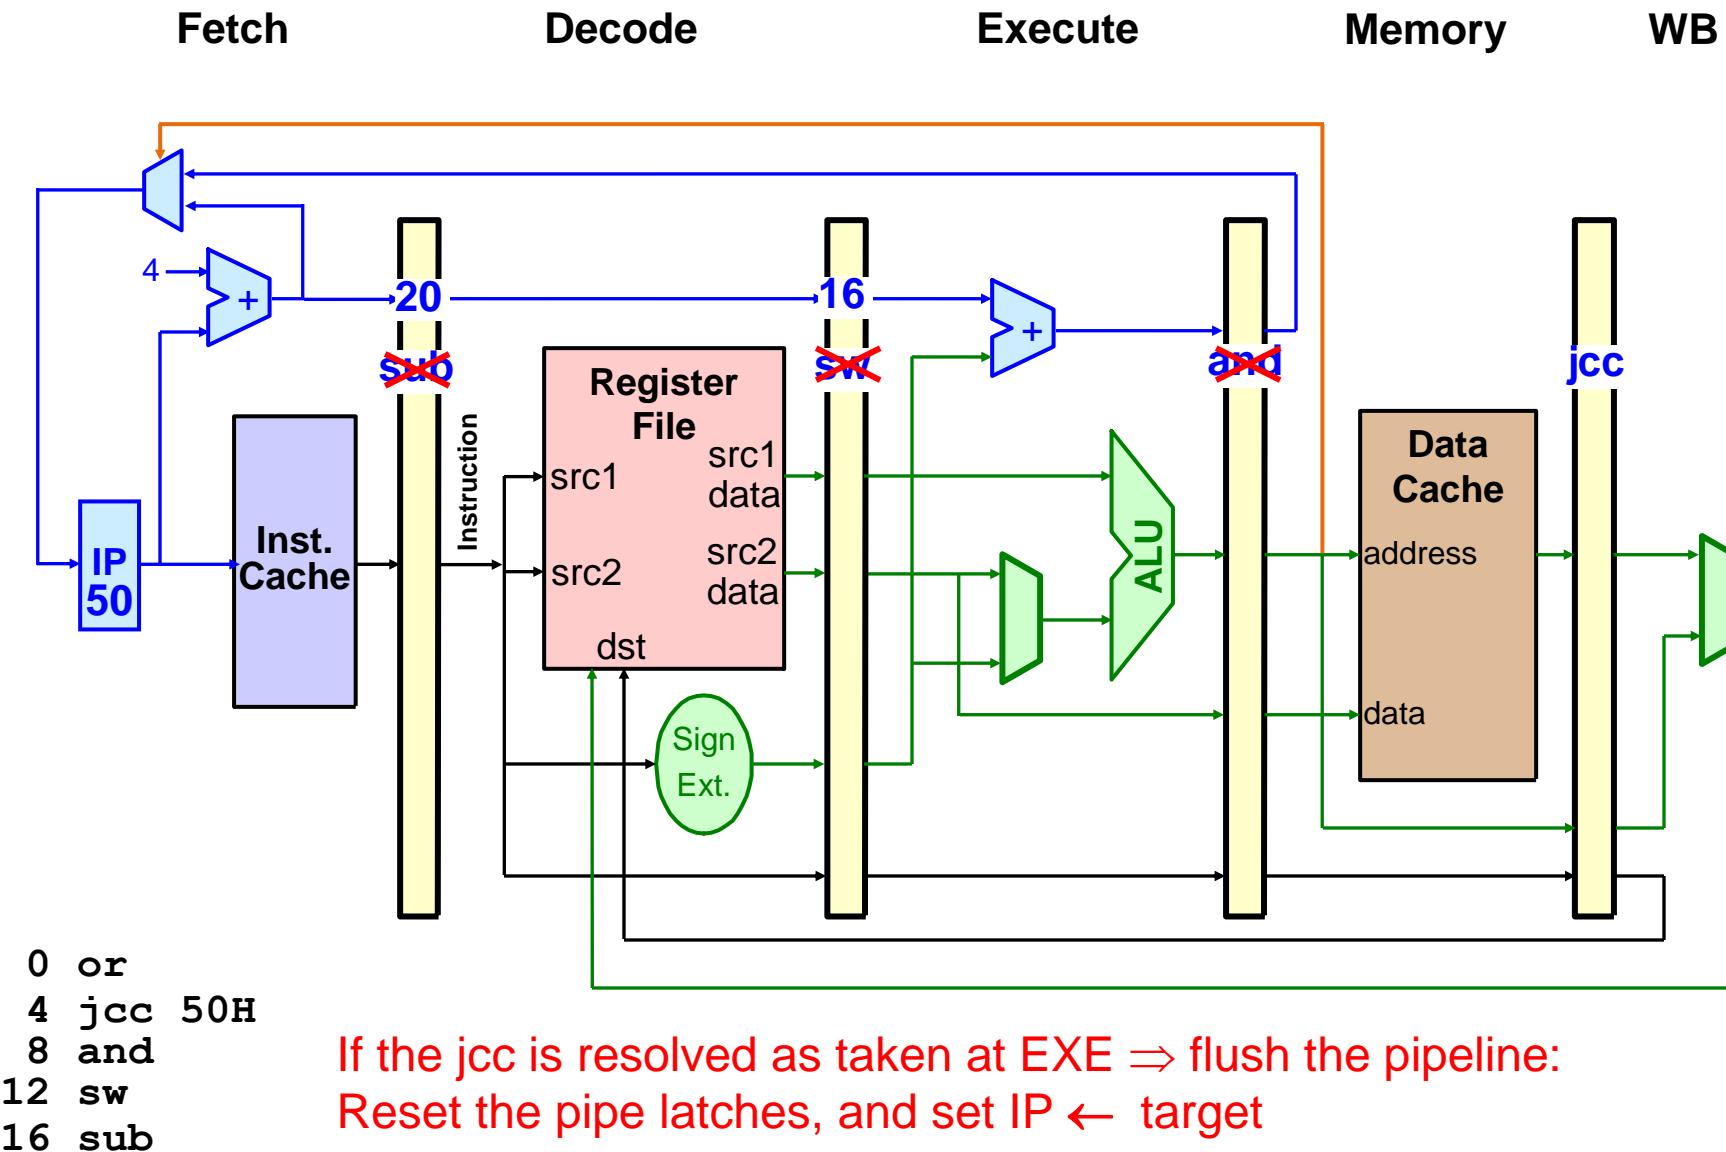
$$(CPI_{new} = CPI_{ideal} + \text{avg. stall cycles / instr.})$$

We loose 60% of the performance

# Control Hazard: Predict Not Taken

- ◆ Execute instructions from the fall-through (not-taken) path
  - ❖ As if there is no branch
  - ❖ If the branch is not-taken (~50%), no penalty is paid
- ◆ If branch actually taken
  - ❖ Flush the fall-through path instructions before they change the machine state (memory / registers)
  - ❖ Fetch the instructions from the correct (taken) path
- ◆ Assuming 20% of instructions are branches and ~50% branches not taken on average  
$$\text{CPI new} = 1 + (0.2 \times 0.5) \times 3 = 1.3$$

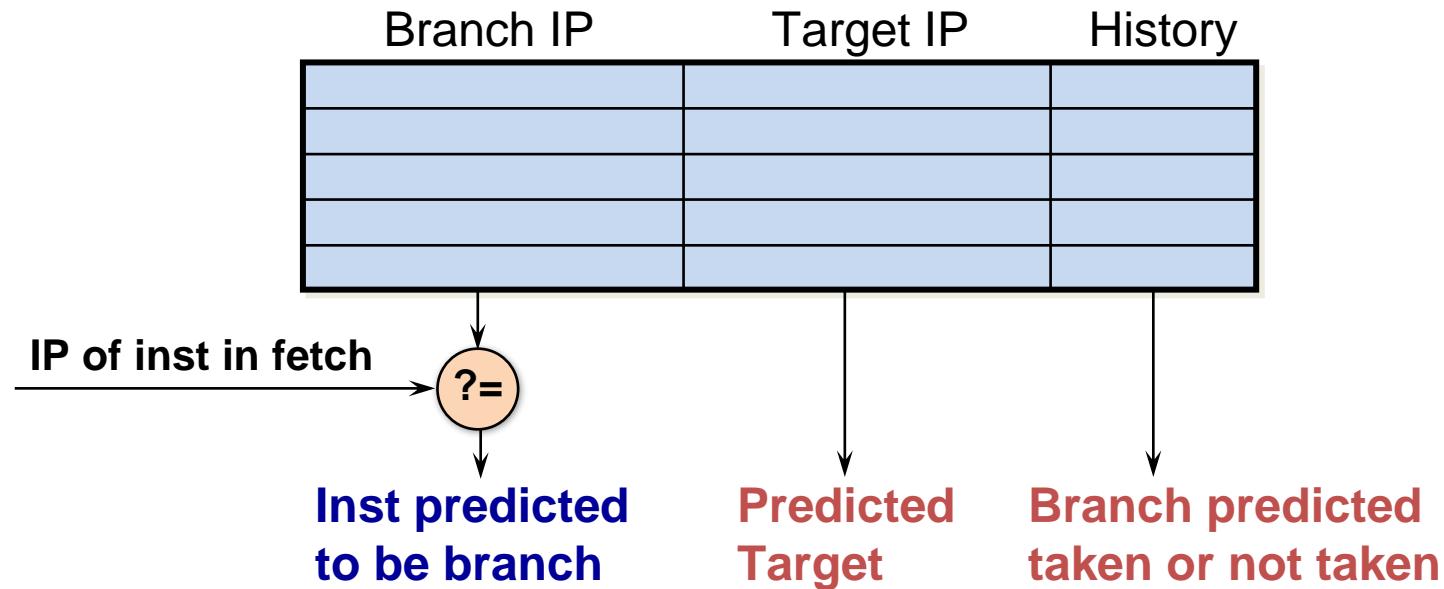
# Control Hazard on Branches



# Dynamic Branch Prediction

- ◆ Add a **Branch Target Buffer (BTB)** that predicts (at fetch)

- ❖ Instruction is a branch
- ❖ Branch taken / not-taken
- ❖ Taken branch target



- ◆ **BTB allocated at execute – after all branch info is known**
- ◆ **BTB is looked up at instruction fetch**

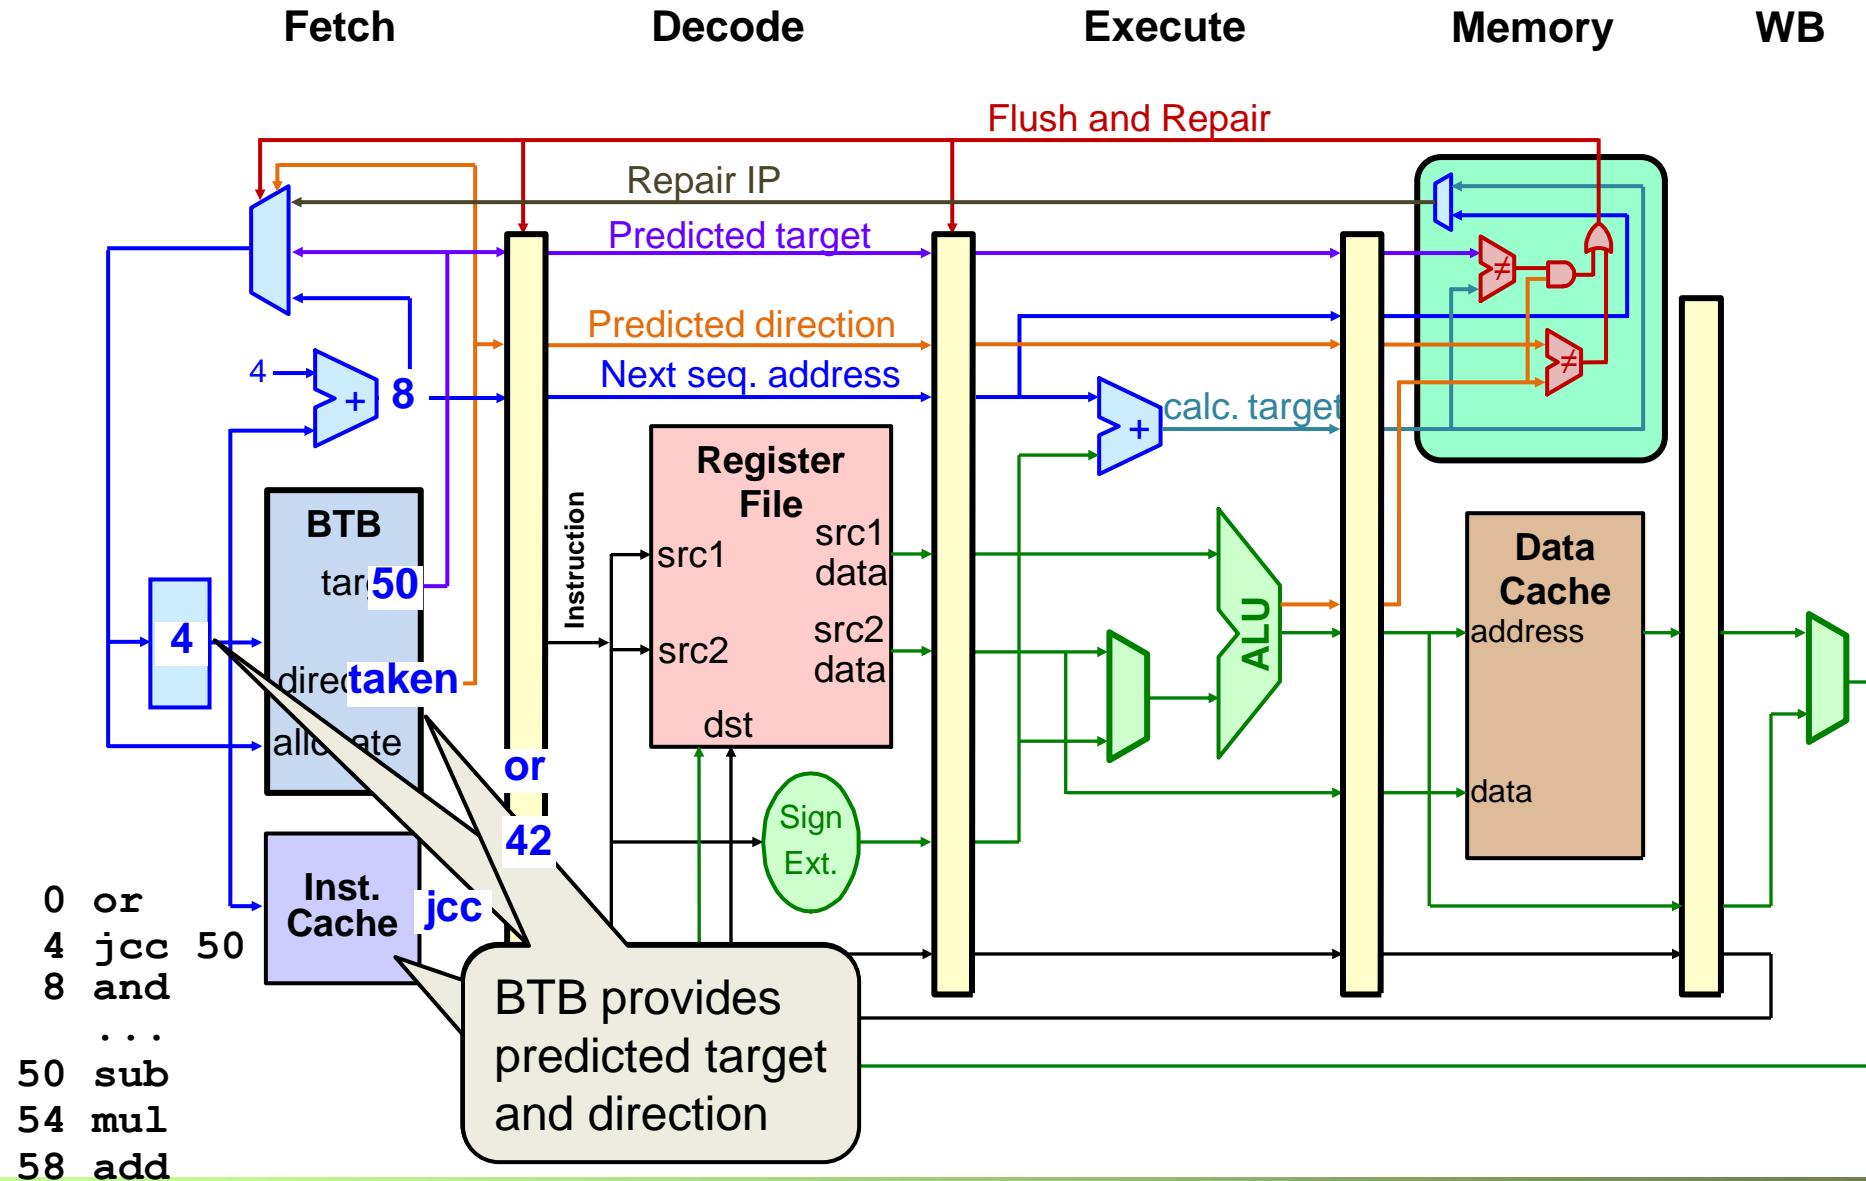
# BTB

- ◆ **Allocation – at Decode / EXE**
  - ❖ Allocate instructions identified as branches (after decode)
    - Both conditional and unconditional branches are allocated
  - ❖ Not taken branches need not be allocated
    - BTB miss implicitly predicts not-taken
- ◆ **Prediction – at Fetch**
  - ❖ BTB lookup is done parallel to IC lookup
  - ❖ BTB provides
    - Indication that the instruction is a branch (BTB hits)
    - Branch predicted target
    - Branch predicted direction
    - Branch predicted type (e.g., conditional, unconditional)
- ◆ **Update (when branch outcome is known – at EXE)**
  - ❖ Branch target
  - ❖ Branch history (taken / not-taken)

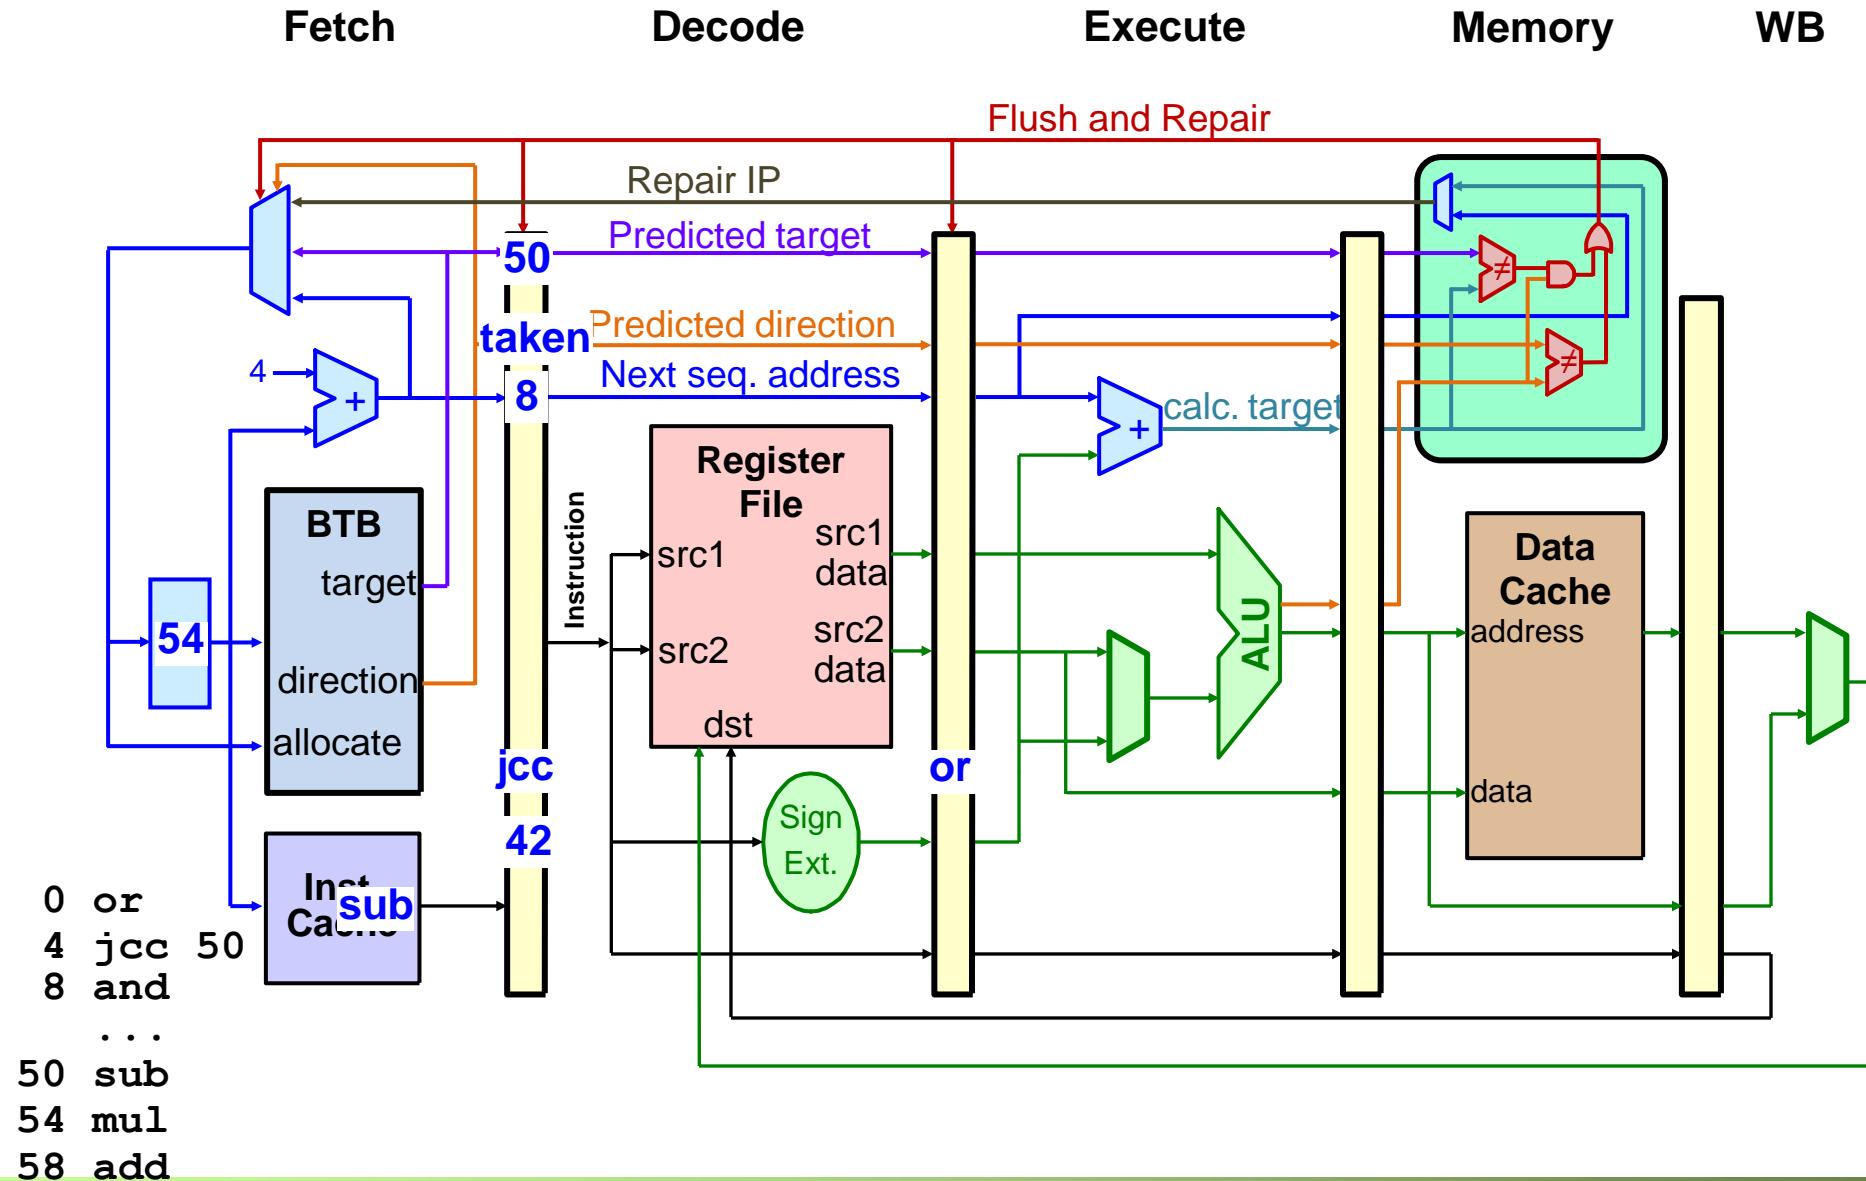
# BTB (cont.)

- ◆ **Wrong prediction**
  - ❖ Predict not-taken, actual taken
  - ❖ Predict taken, actual not-taken, or actual taken but wrong target
- ◆ **In case of wrong prediction – flush the pipeline**
  - ❖ Reset latches (same as making all instructions to be NOPs)
  - ❖ Select the PC source to be from the correct path
    - Need get the fall-through with the branch
  - ❖ Start fetching instruction from correct path
- ◆ **Assuming P% correct prediction rate**
$$\text{CPI new} = 1 + (0.2 \times (1-P)) \times 3$$
  - ❖ For example, if P=0.7
$$\text{CPI new} = 1 + (0.2 \times 0.3) \times 3 = 1.18$$

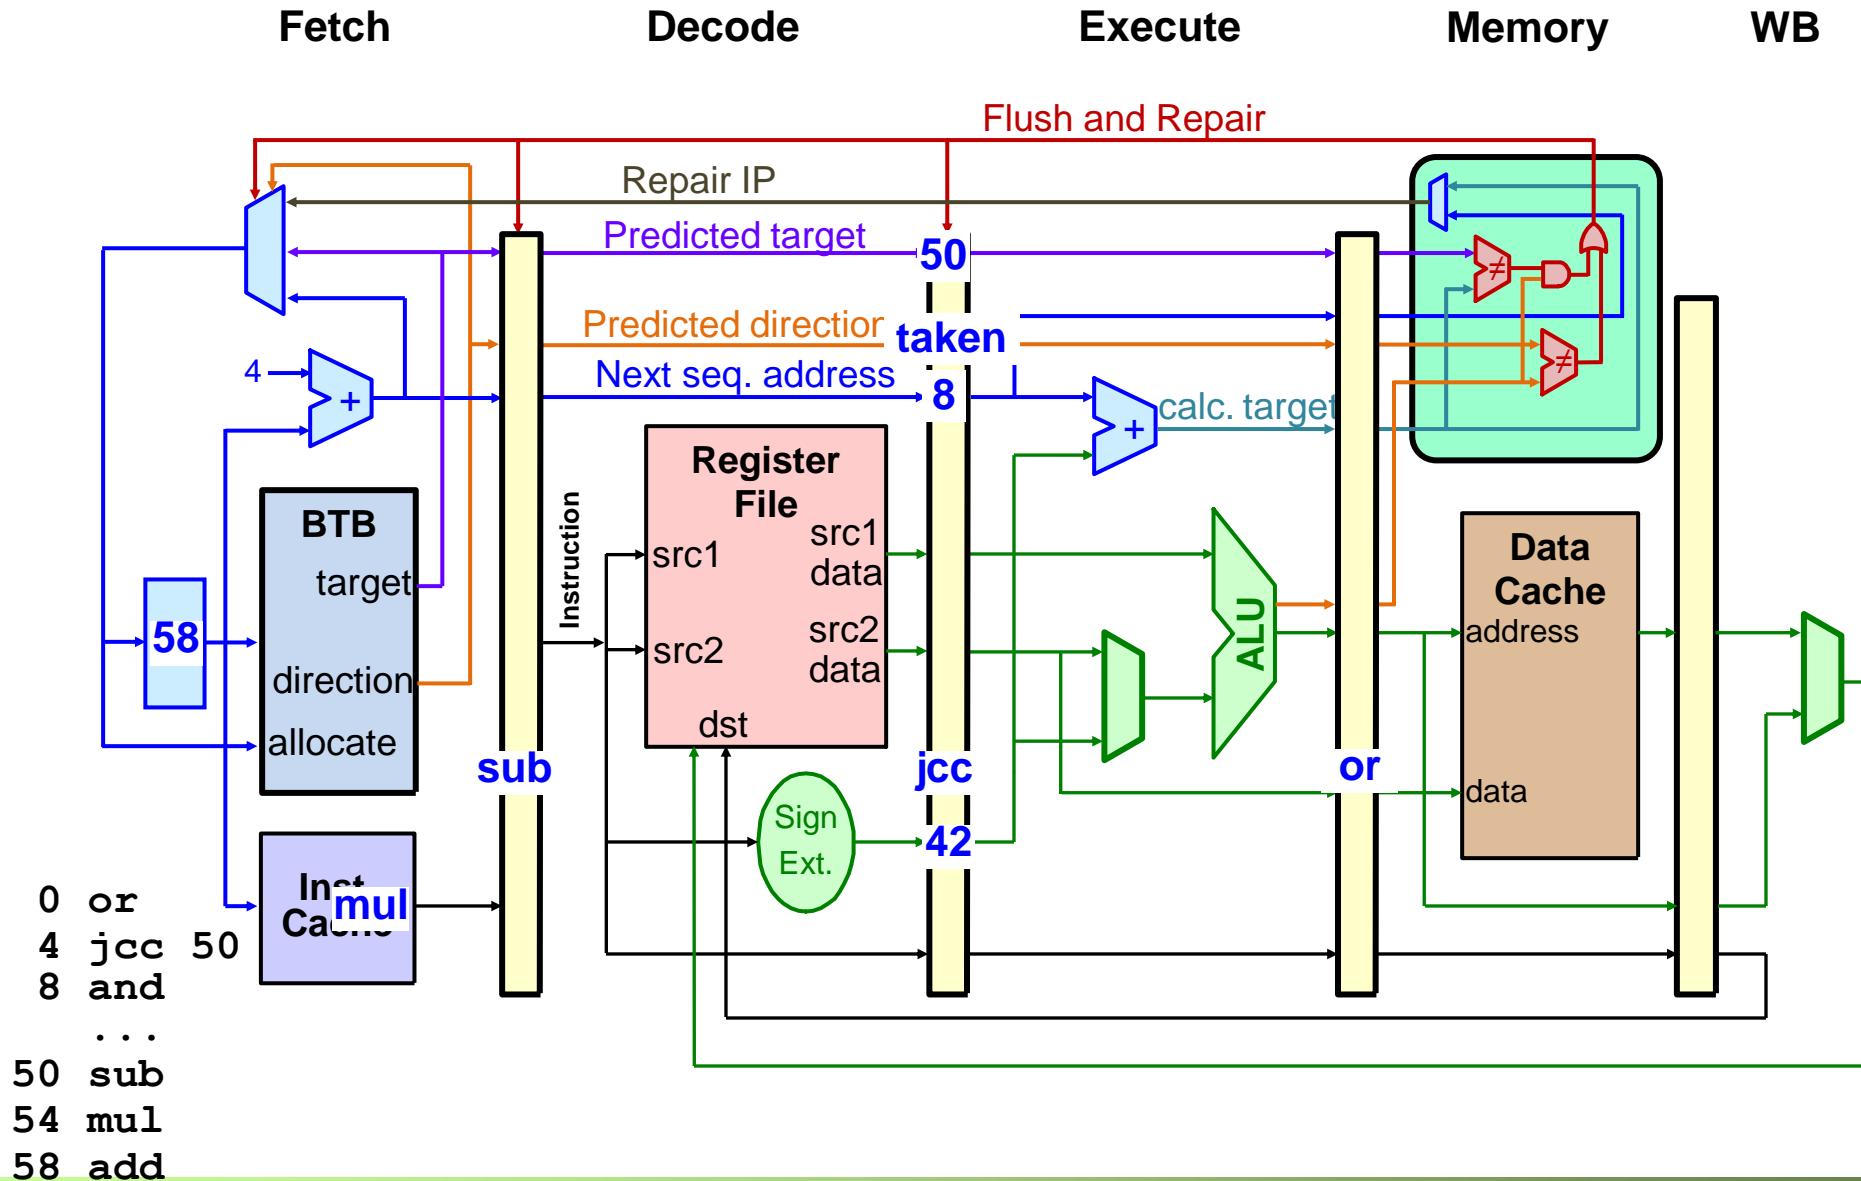
# Adding a BTB to the Pipeline



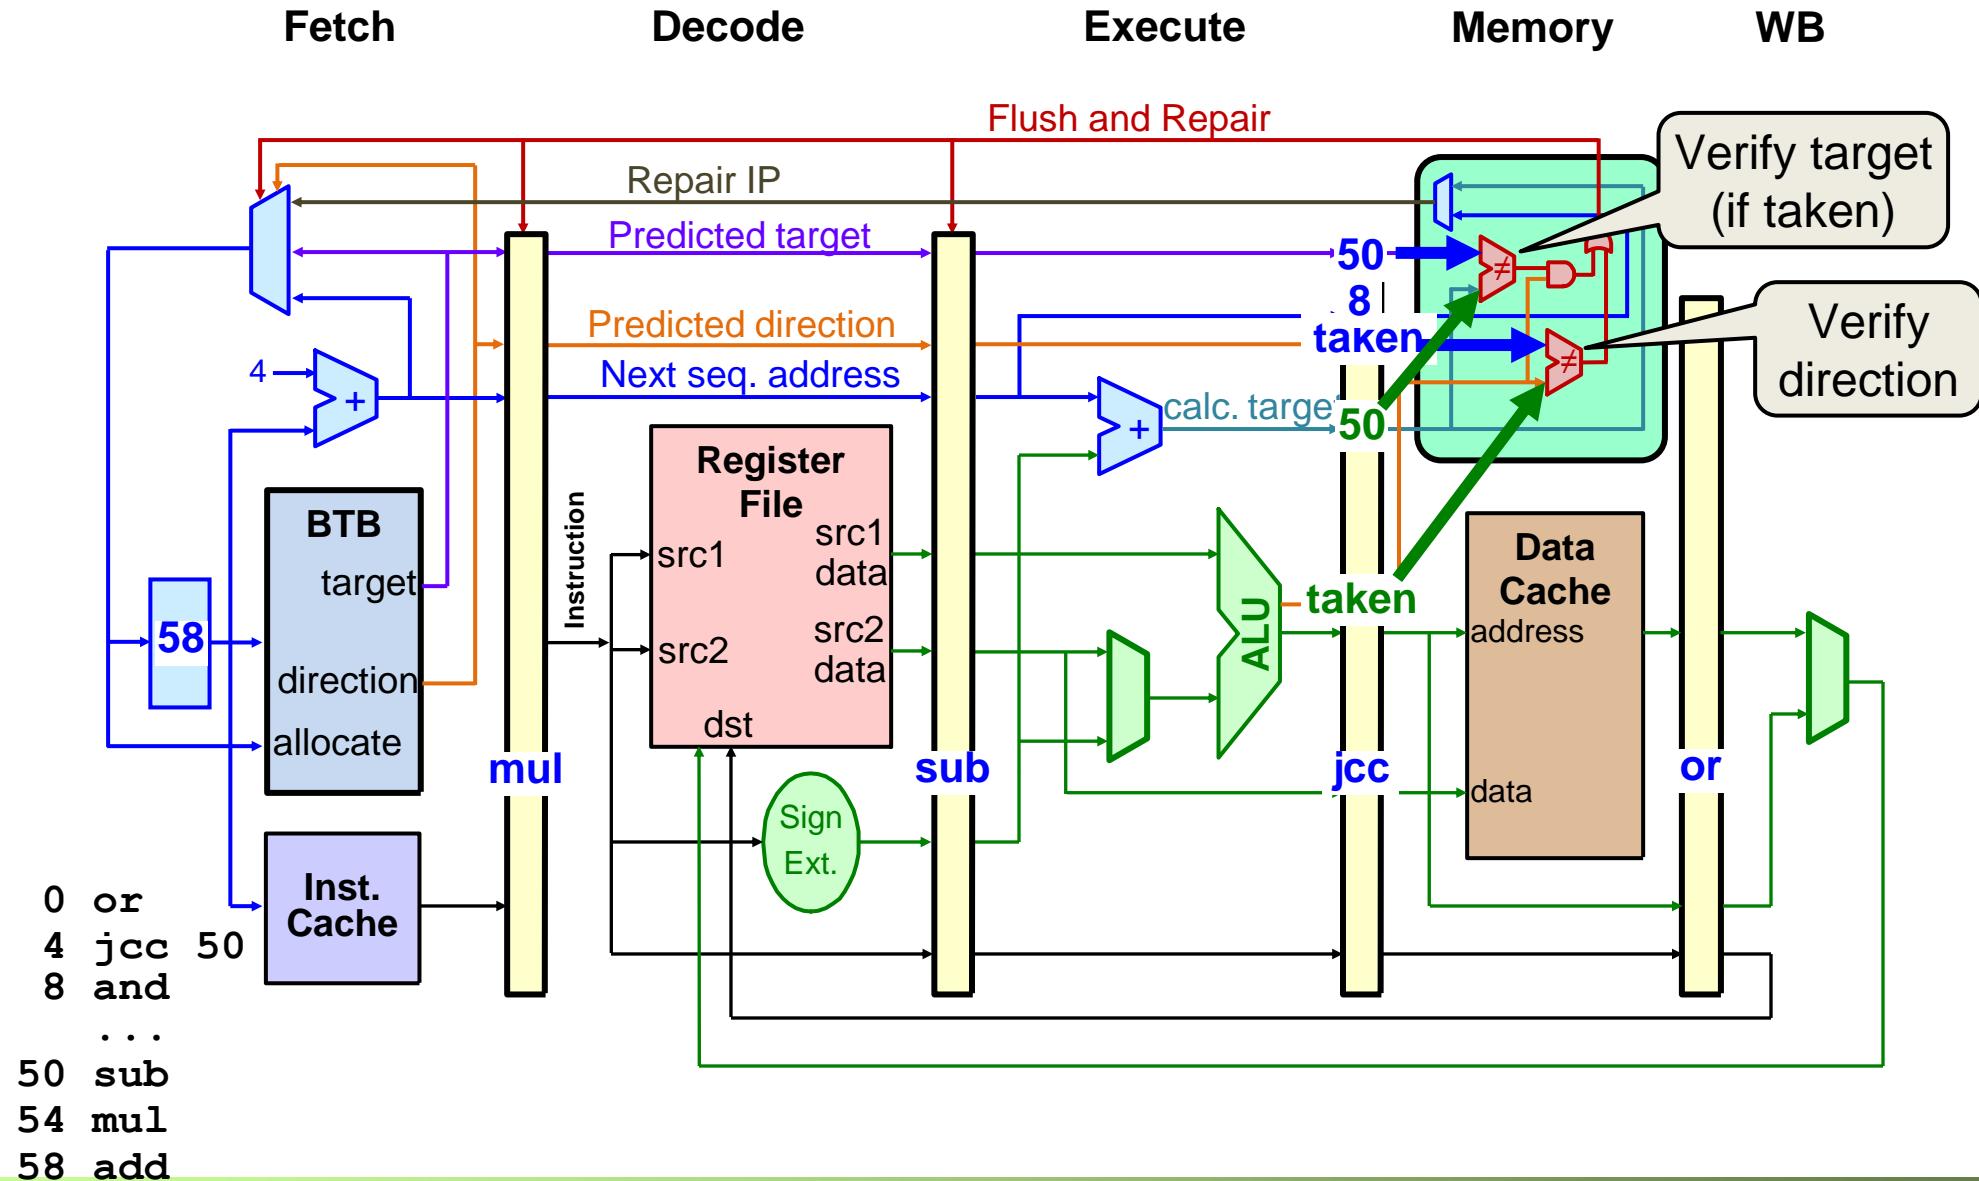
# Adding a BTB to the Pipeline



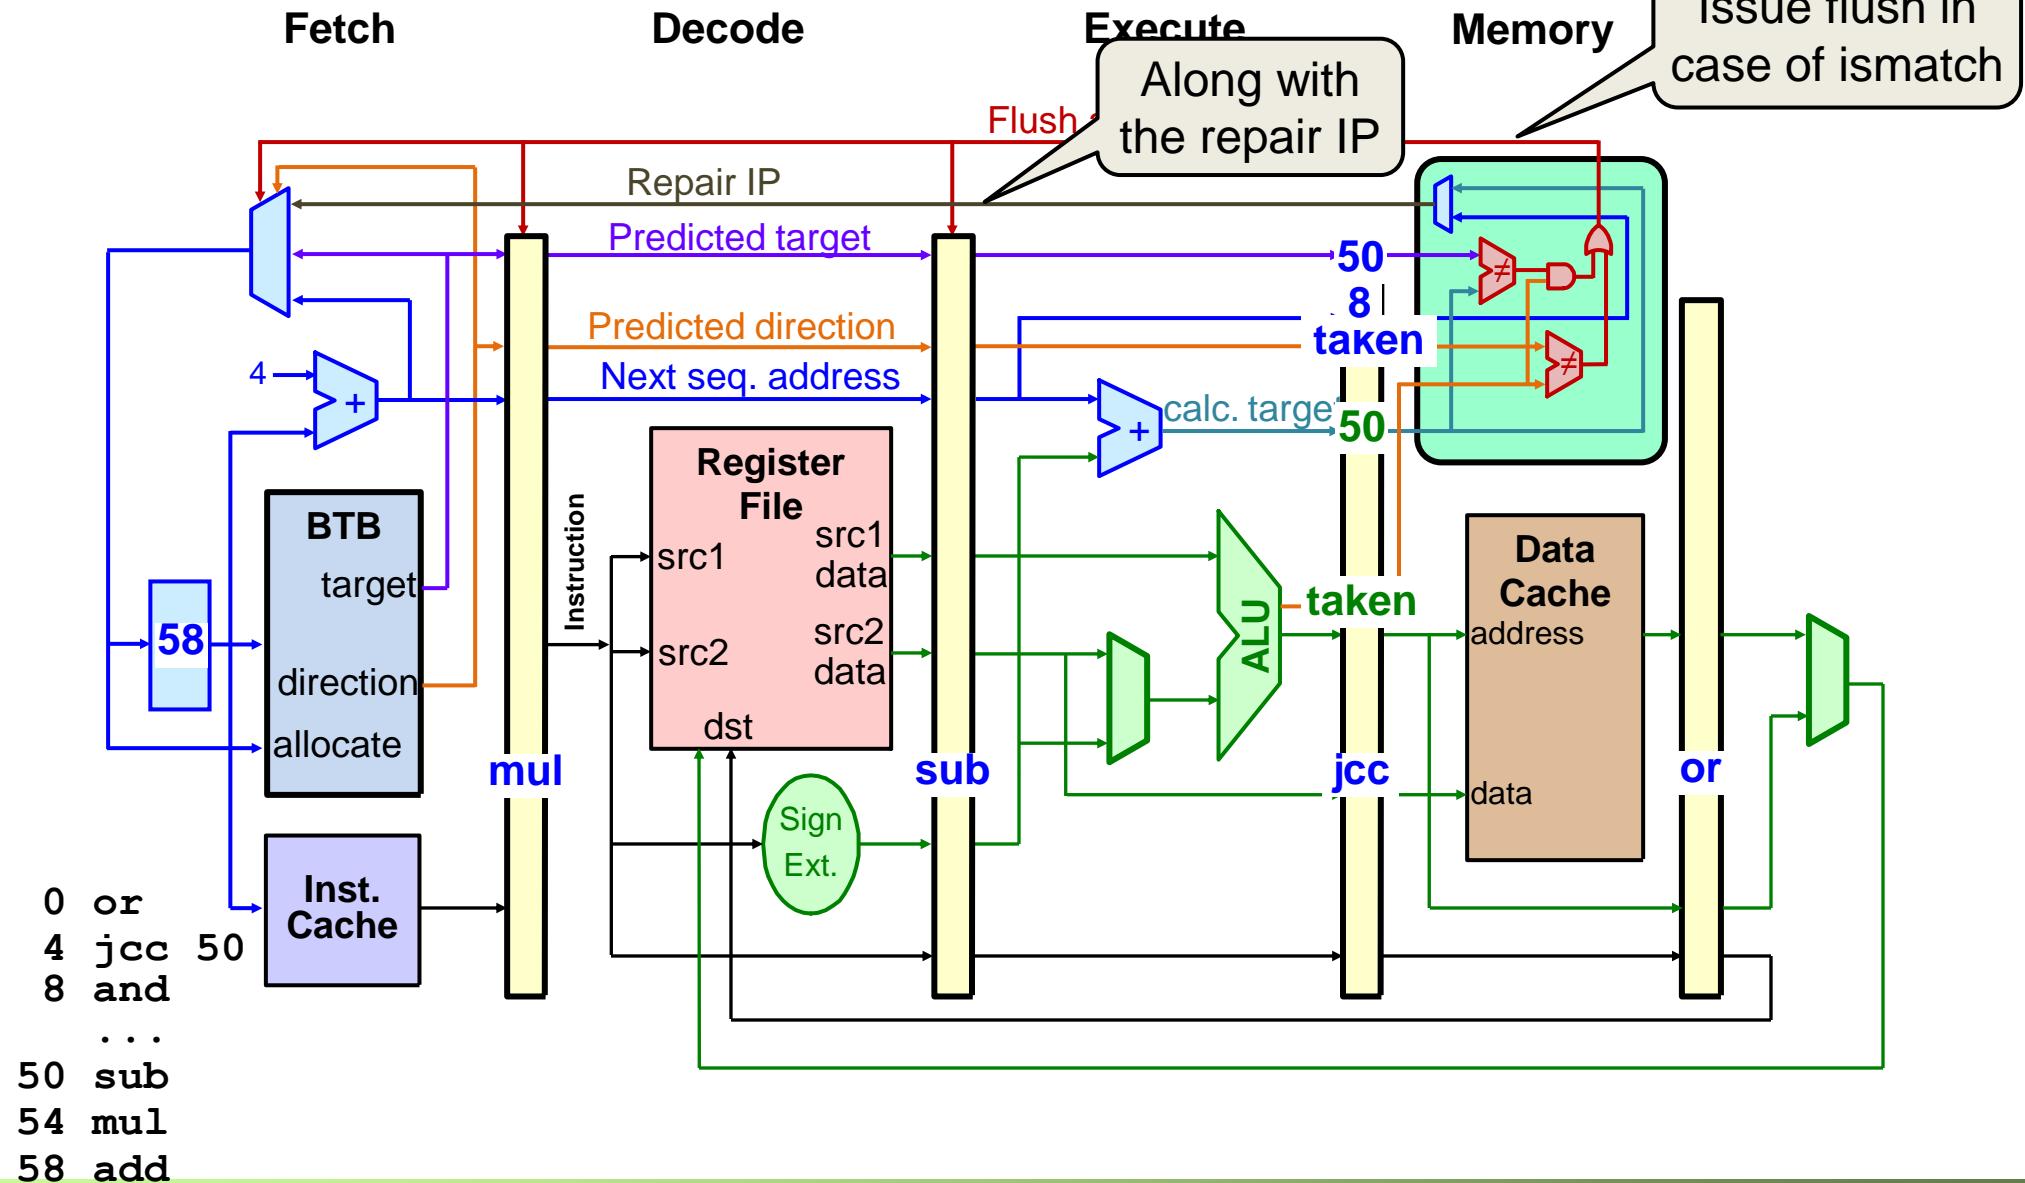
# Adding a BTB to the Pipeline



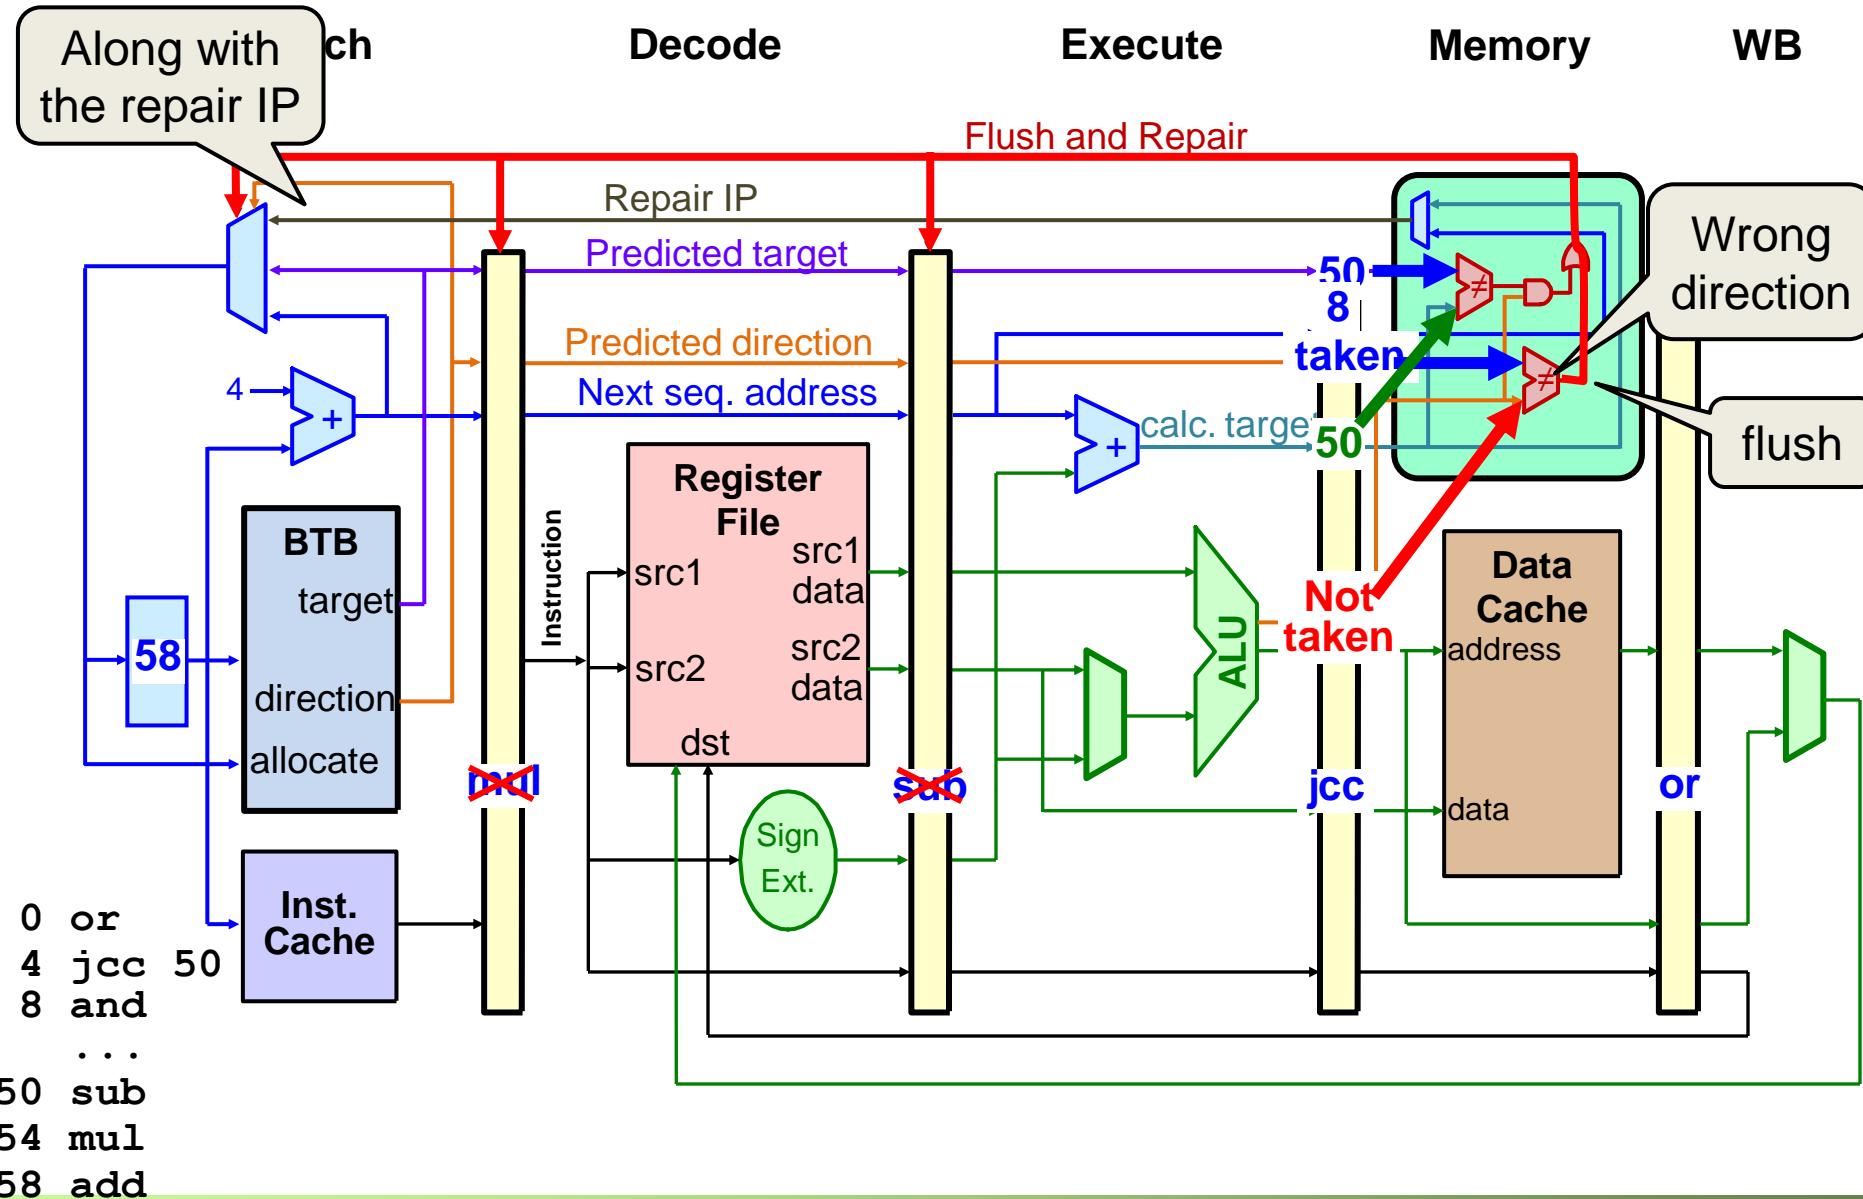
# Adding a BTB to the Pipeline



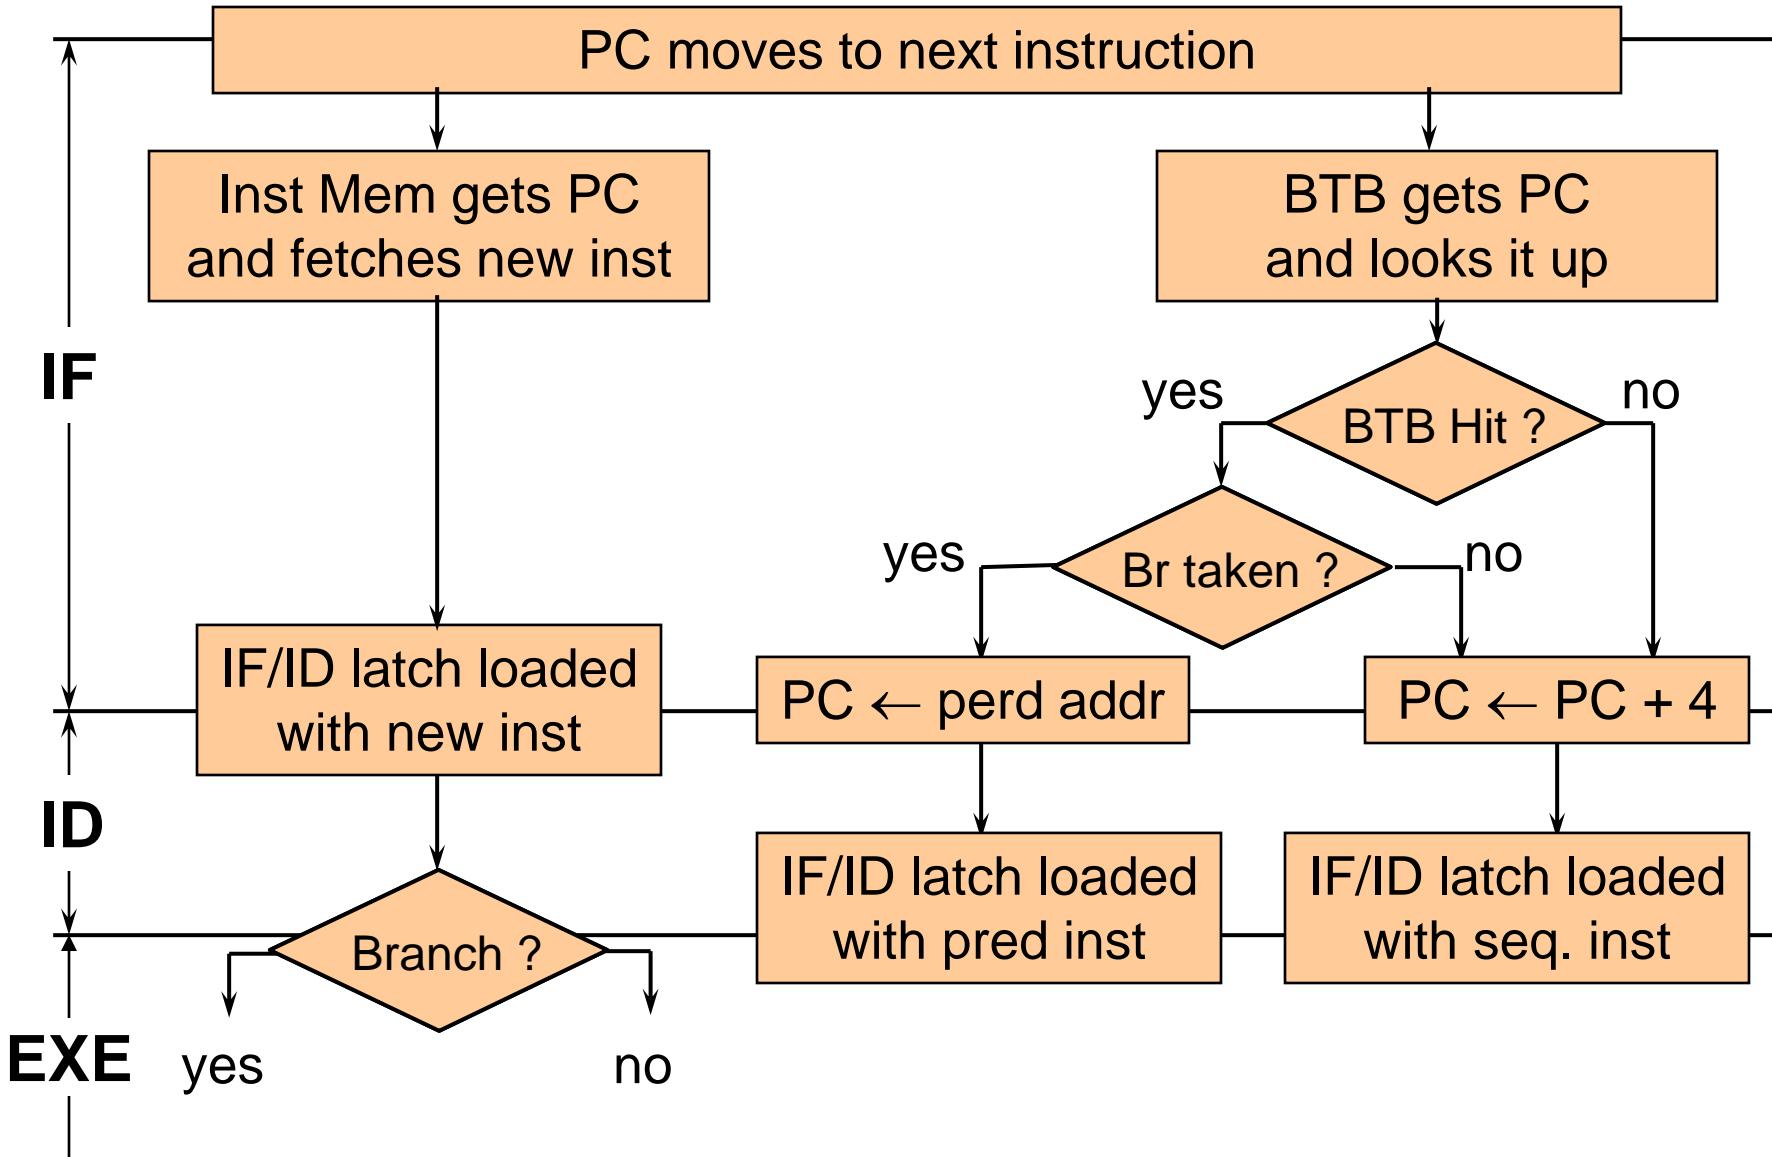
# Adding a BTB to the Pipeline



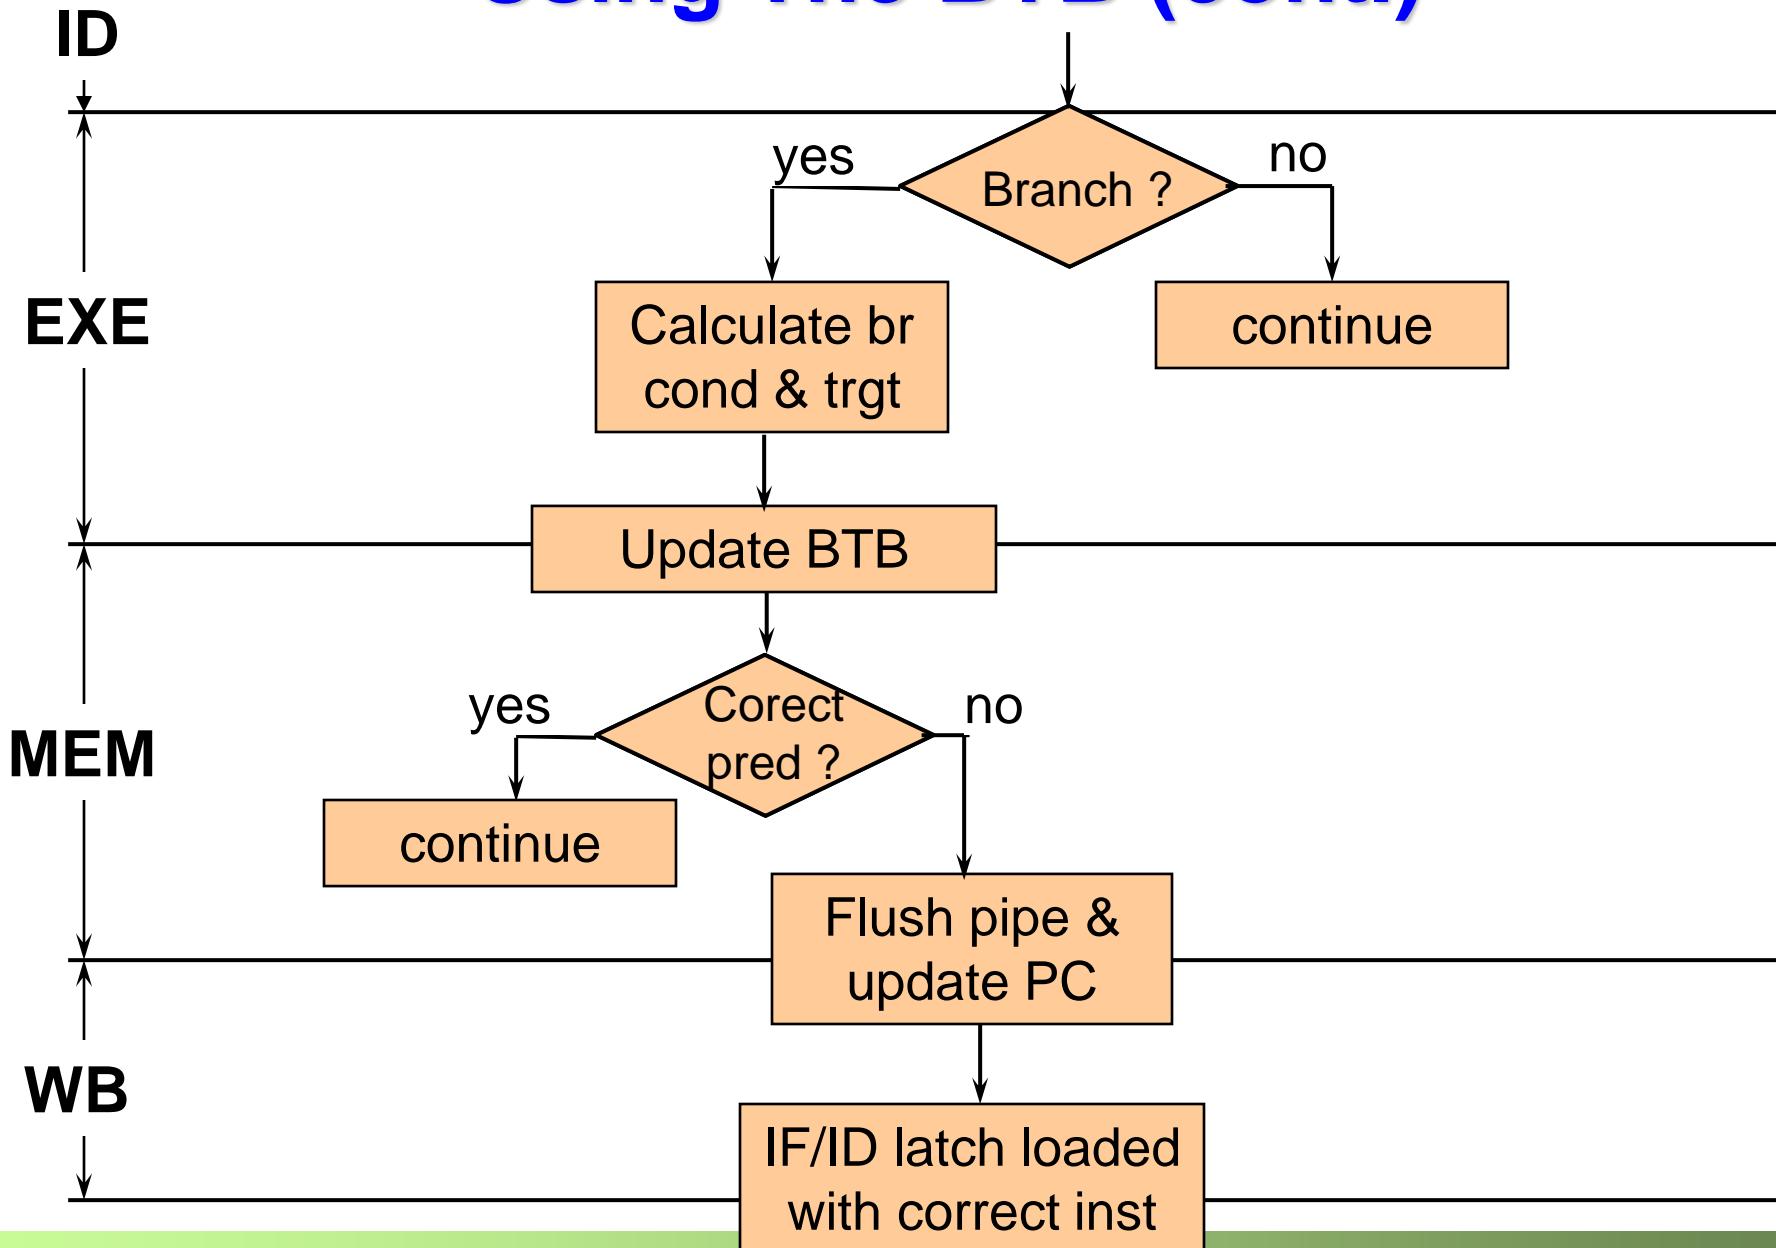
# Predicted Taken Actual Not Taken



# Using The BTB



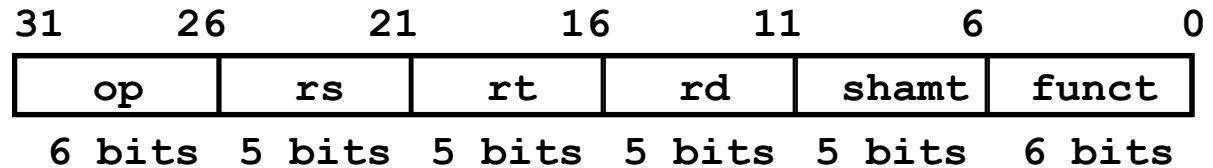
# Using The BTB (cont.)



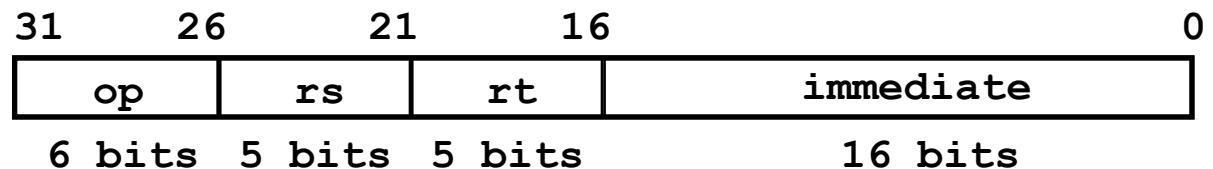
# Backup

# MIPS Instruction Formats

- R-type  
(register insts)



- I-type (Load, Store, Branch,  
inst's w/imm data)



- J-type (Jump)



**op:** operation of the instruction

**rs, rt, rd:** the source and destination register specifiers

**shamt:** shift amount

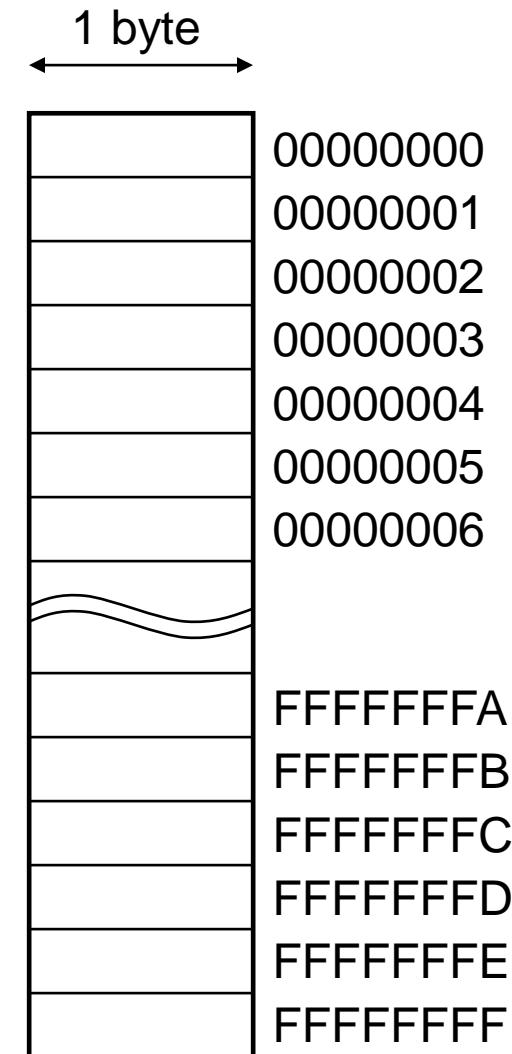
**funct:** selects the variant of the operation in the “op” field

**address / immediate:** address offset or immediate value

**target address:** target address of the jump instruction

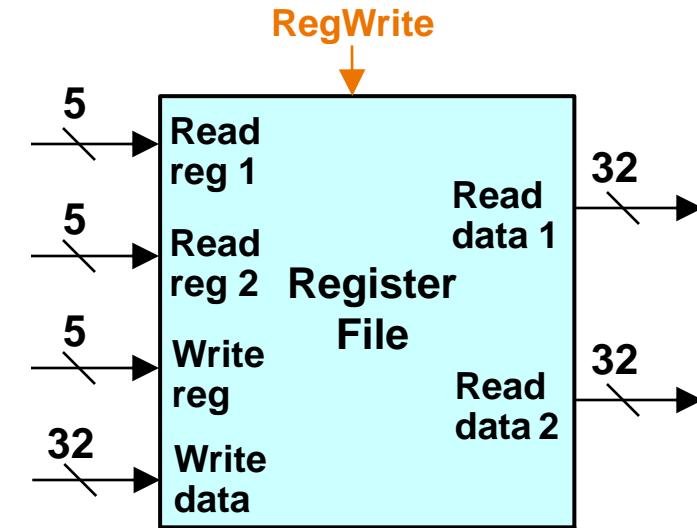
# MIPS Memory Space

- ◆ **Each memory location**
  - ❖ is 8 bit = 1 byte wide
  - ❖ has an **address**
- ◆ **We assume 32 byte address**
  - ❖ An address space of  $2^{32}$  bytes
- ◆ **Memory stores both instructions and data**
  - ❖ Each instruction is 32 bit wide  $\Rightarrow$  stored in 4 consecutive bytes in memory
  - ❖ Various data types have different width



# Register File

- ◆ The Register File holds 32 registers
- ◆ Each register is 32 bit wide
- ◆ The RF supports parallel
  - ❖ reading any two registers and
  - ❖ writing any register
- ◆ Inputs
  - ❖ **Read reg 1/2**: #register whose value will be output on *Read data 1/2*
  - ❖ **RegWrite**: write enable
  - ❖ **Write reg** (relevant when *RegWrite*=1)
    - #register to which the value in *Write data* is written to
  - ❖ **Write data** (relevant when *RegWrite*=1)
    - data written to *Write reg*
- ◆ Outputs
  - ❖ **Read data 1/2**: data read from *Read reg 1/2*



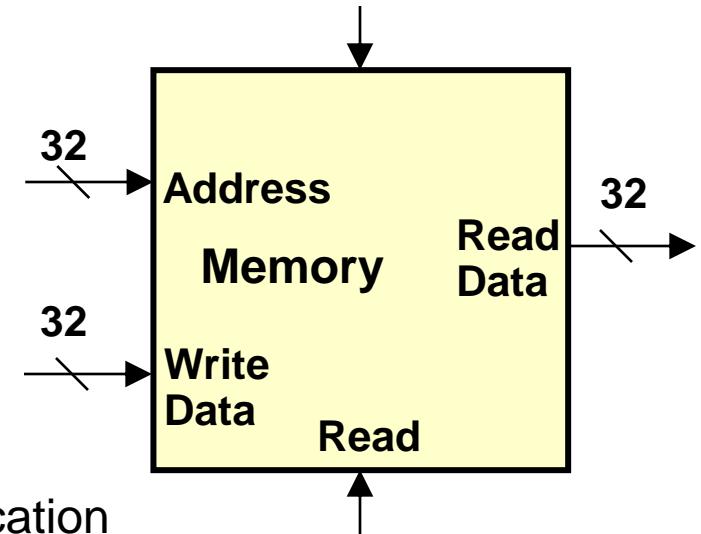
# Memory Components

## ◆ Inputs

- ❖ **Address**: address of the memory location we wish to access
- ❖ **Read**: read data from location
- ❖ **Write**: write data into location
- ❖ **Write data** (relevant when Write=1) data to be written into specified location

## ◆ Outputs

- ❖ **Read data** (relevant when Read=1) data read from specified location



# Cache

## ◆ Memory components are slow relative to the CPU

- ❖ A **cache** is a fast memory which contains only small part of the memory
- ❖ **Instruction cache** stores parts of the memory space which hold code
- ❖ **Data Cache** stores parts of the memory space which hold data

# The Program Counter (PC)

- ◆ Holds the address (in memory) of the next instruction to be executed
  - ◆ After each instruction, advanced to point to the next instruction

- ❖ If the current instruction is not a taken branch,  
the next instruction resides right after the current instruction

**PC**  $\leftarrow$  **PC** + 4

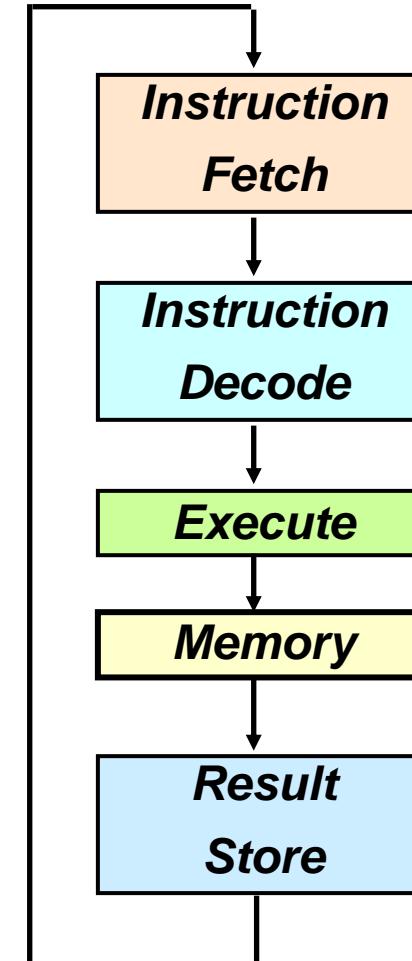
- ❖ If the current instruction is a taken branch,  
the next instruction resides at the branch target

**PC  $\leftarrow$  target**      (absolute jump)

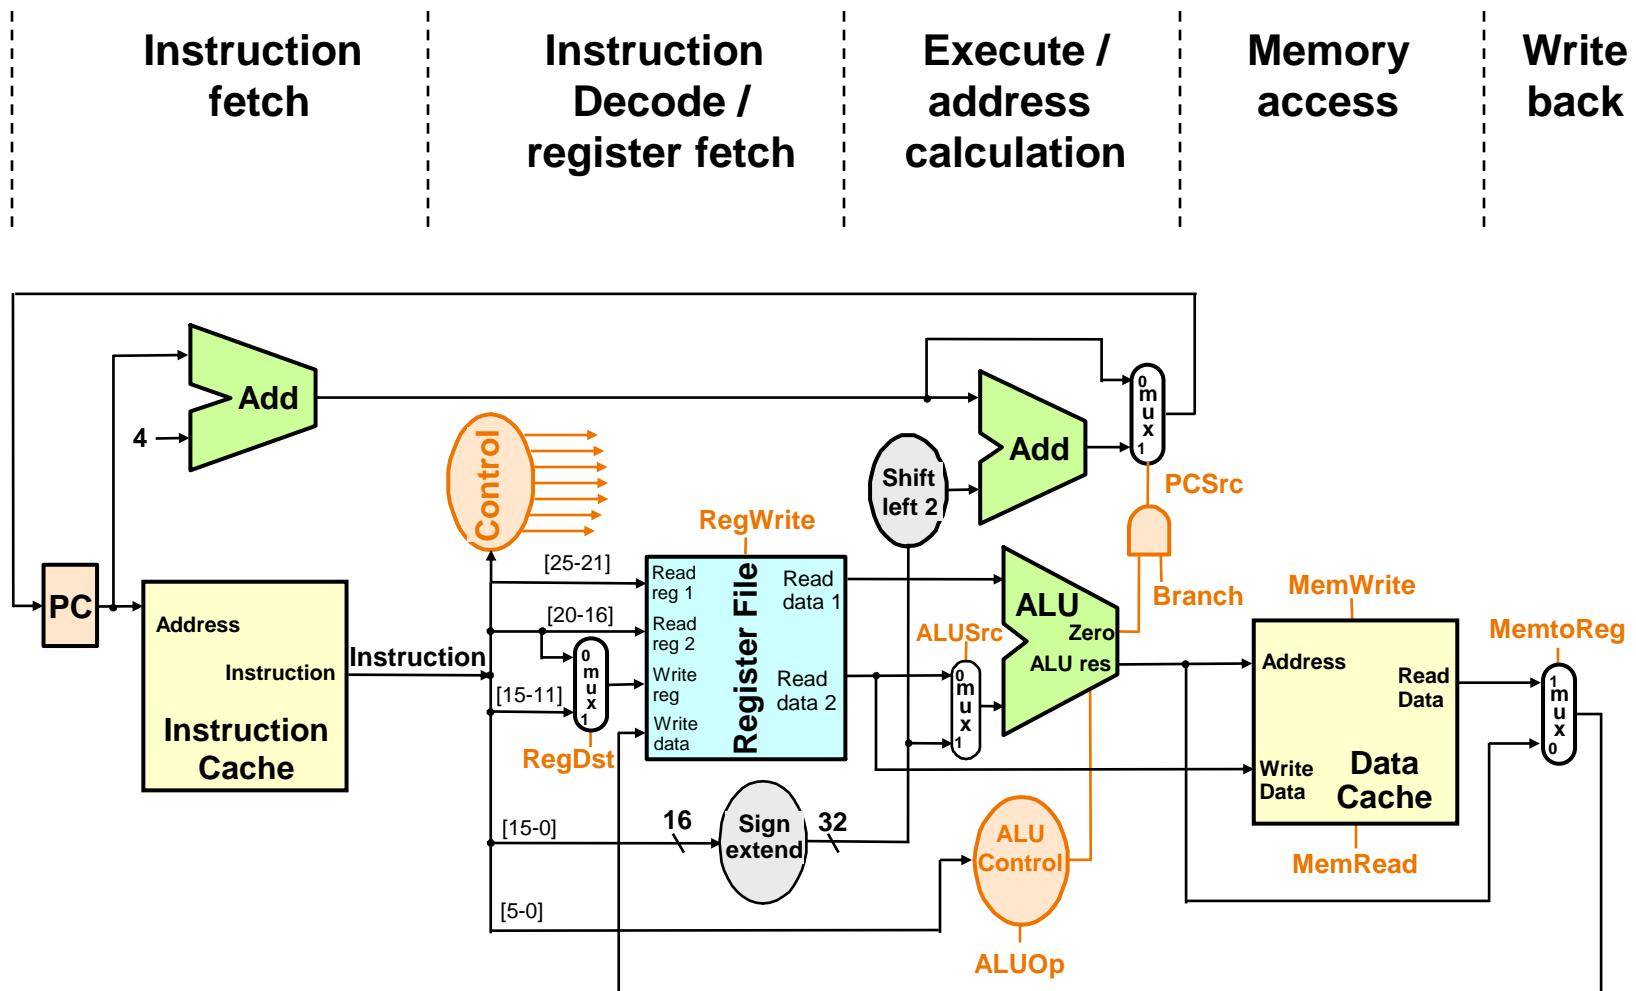
**PC  $\leftarrow$  PC + 4 + offset $\times$ 4 (relative jump)**

# Instruction Execution Stages

- ◆ **Fetch**
  - ❖ Fetch instruction pointed by PC from I-Cache
- ◆ **Decode**
  - ❖ Decode instruction (generate control signals)
  - ❖ Fetch operands from register file
- ◆ **Execute**
  - ❖ For a memory access: calculate effective address
  - ❖ For an ALU operation: execute operation in ALU
  - ❖ For a branch: calculate condition and target
- ◆ **Memory Access**
  - ❖ For load: read data from memory
  - ❖ For store: write data into memory
- ◆ **Write Back**
  - ❖ Write result back to register file
  - ❖ update program counter

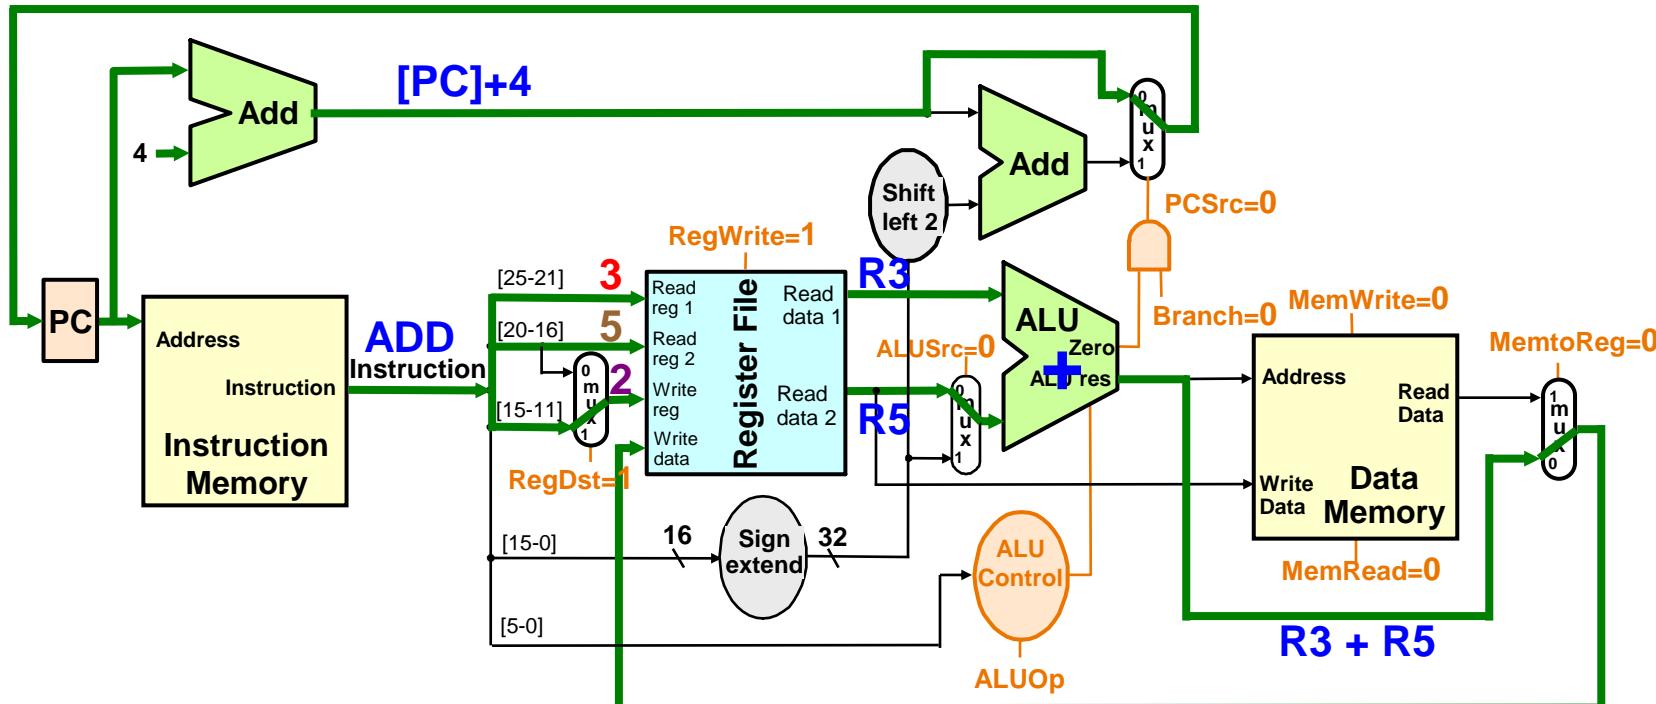
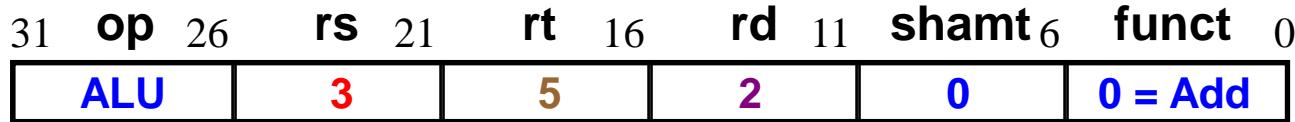


# The MIPS CPU



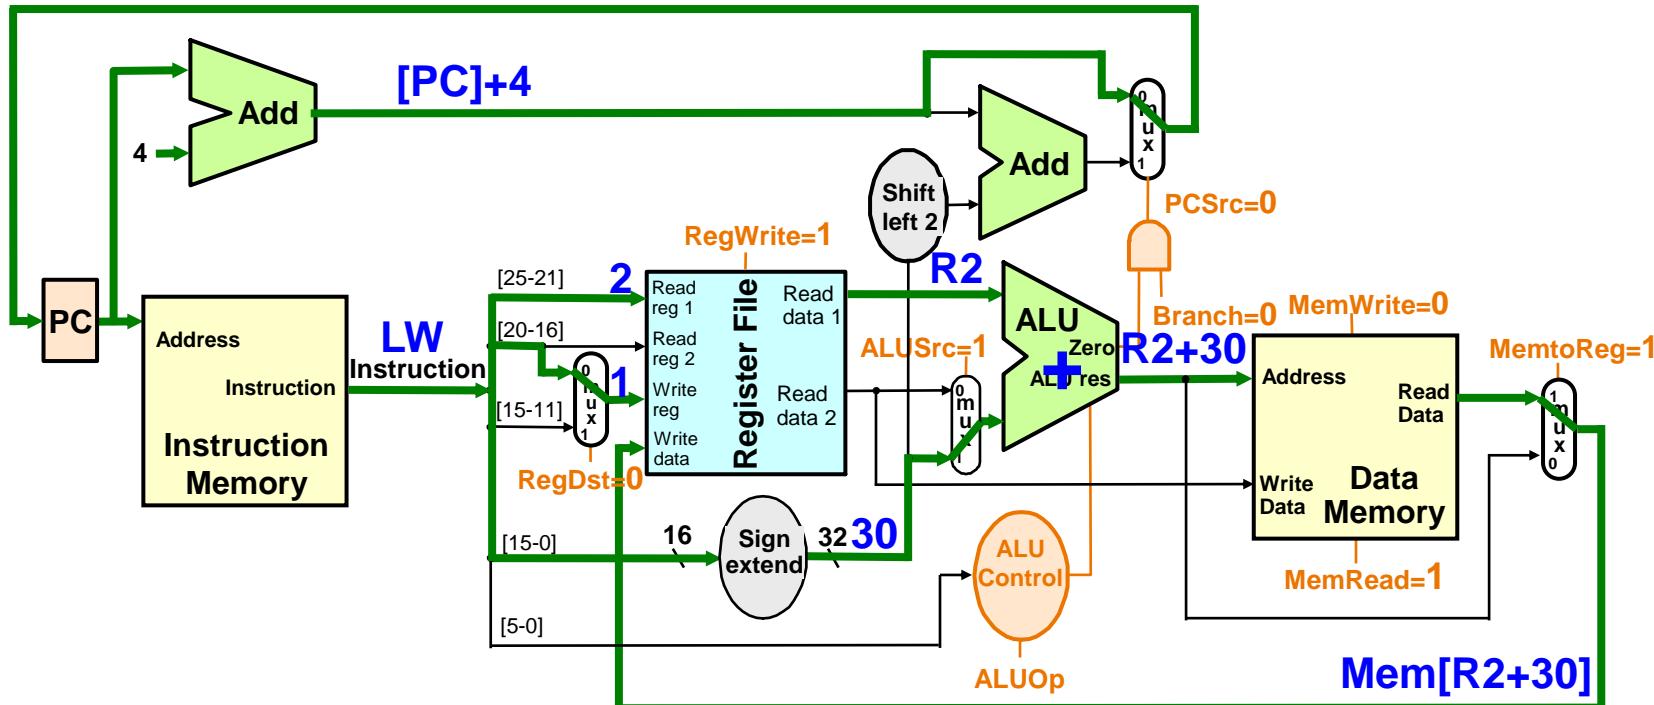
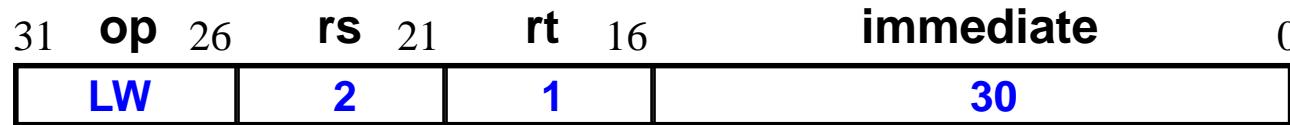
# Executing an Add Instruction

Add R2, R3, R5 ;  $R2 \leftarrow R3 + R5$



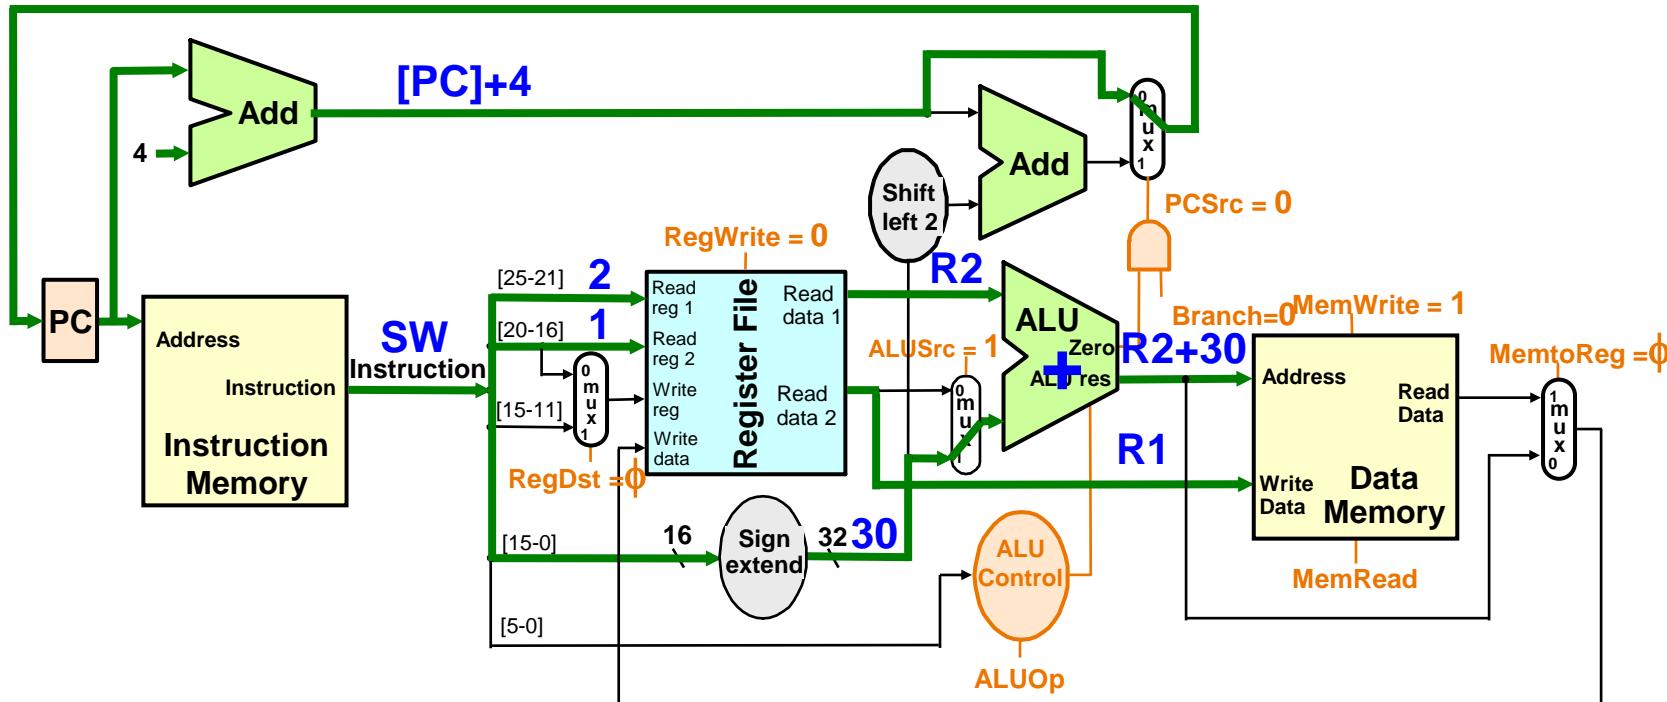
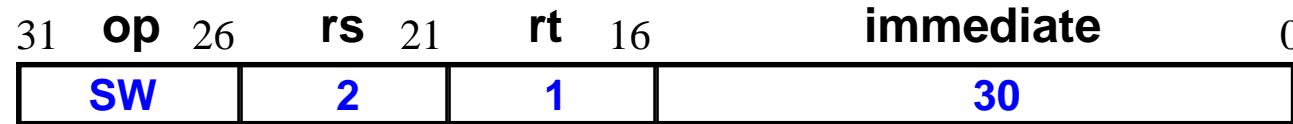
# Executing a Load Instruction

**LW R1, (30)R2 ; R1  $\leftarrow$  Mem[R2+30]**



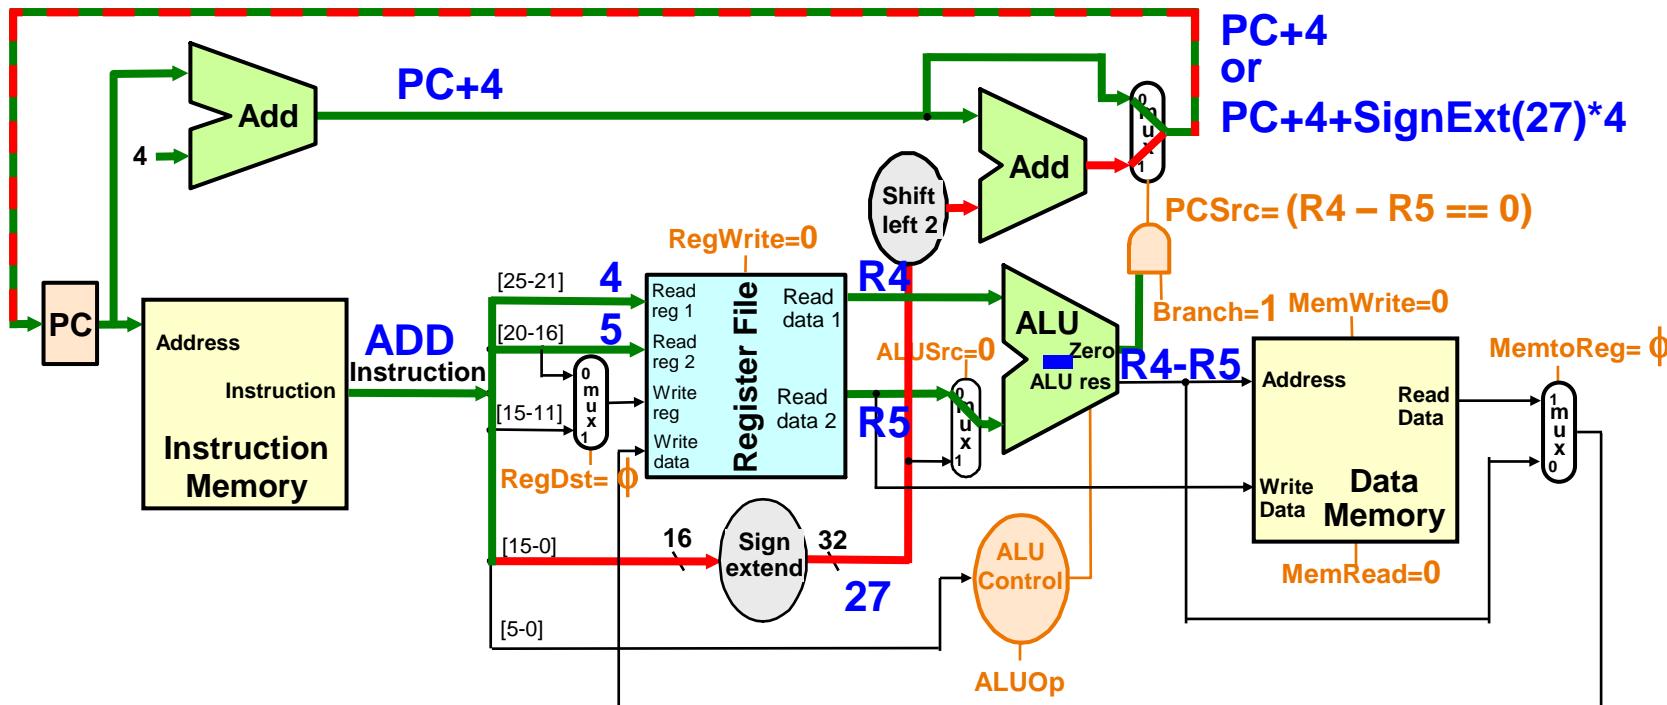
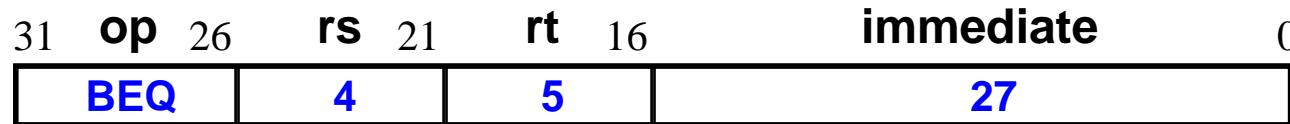
# Executing a Store Instruction

**SW R1, (30)R2 ; Mem[R2+30] ← R1**



# Executing a BEQ Instruction

**BEQ R4, R5, 27 ; if  $(R4-R5=0)$  then  $PC \leftarrow PC+4+\text{SignExt}(27)*4$  ;  
else  $PC \leftarrow PC+4$**

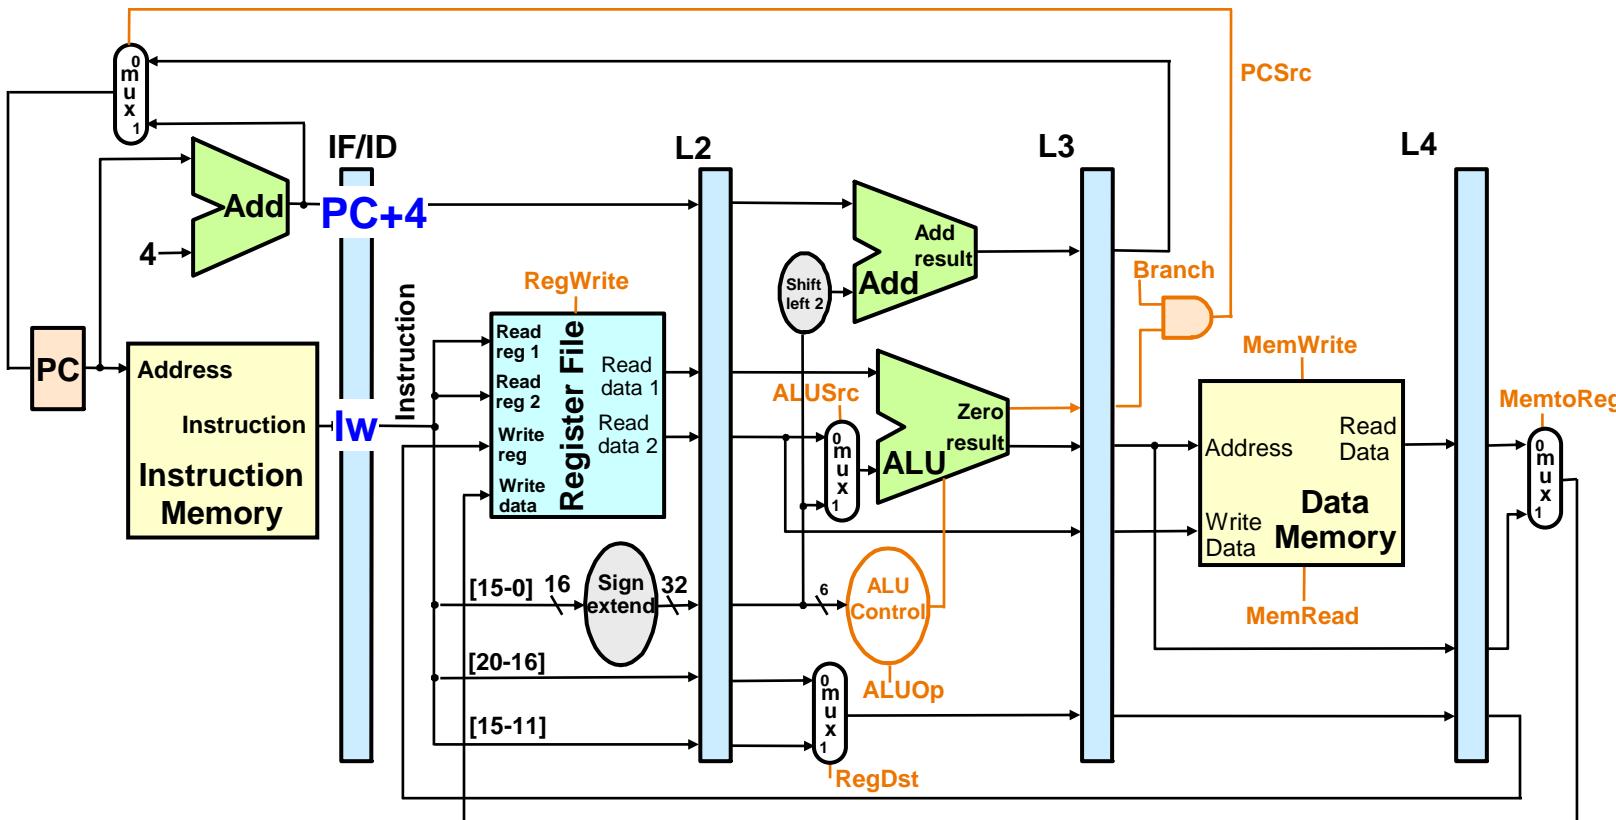
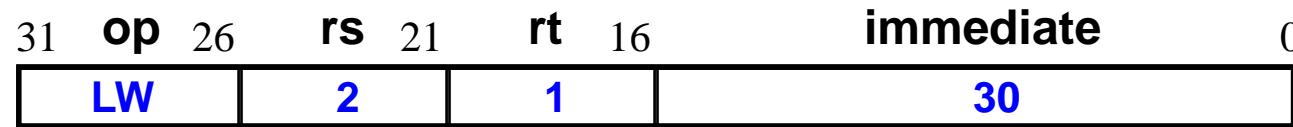


# Control Signals

func	10 0000	10 0010	Don't Care				
op	00 0000	00 0000	00 1101	10 0011	10 1011	00 0100	00 0010
	add	sub	ori	lw	sw	beq	jump
<b>RegDst</b>	1	1	0	0	x	x	x
<b>ALUSrc</b>	0	0	1	1	1	0	x
<b>MemtoReg</b>	0	0	0	1	x	x	x
<b>RegWrite</b>	1	1	1	1	0	0	0
<b>MemWrite</b>	0	0	0	0	1	0	0
<b>Branch</b>	0	0	0	0	0	1	x
<b>Jump</b>	0	0	0	0	0	0	1
<b>ALUctr&lt;2:0&gt;</b>	Add	Subtract	Or	Add	Add	Subtract	xxx

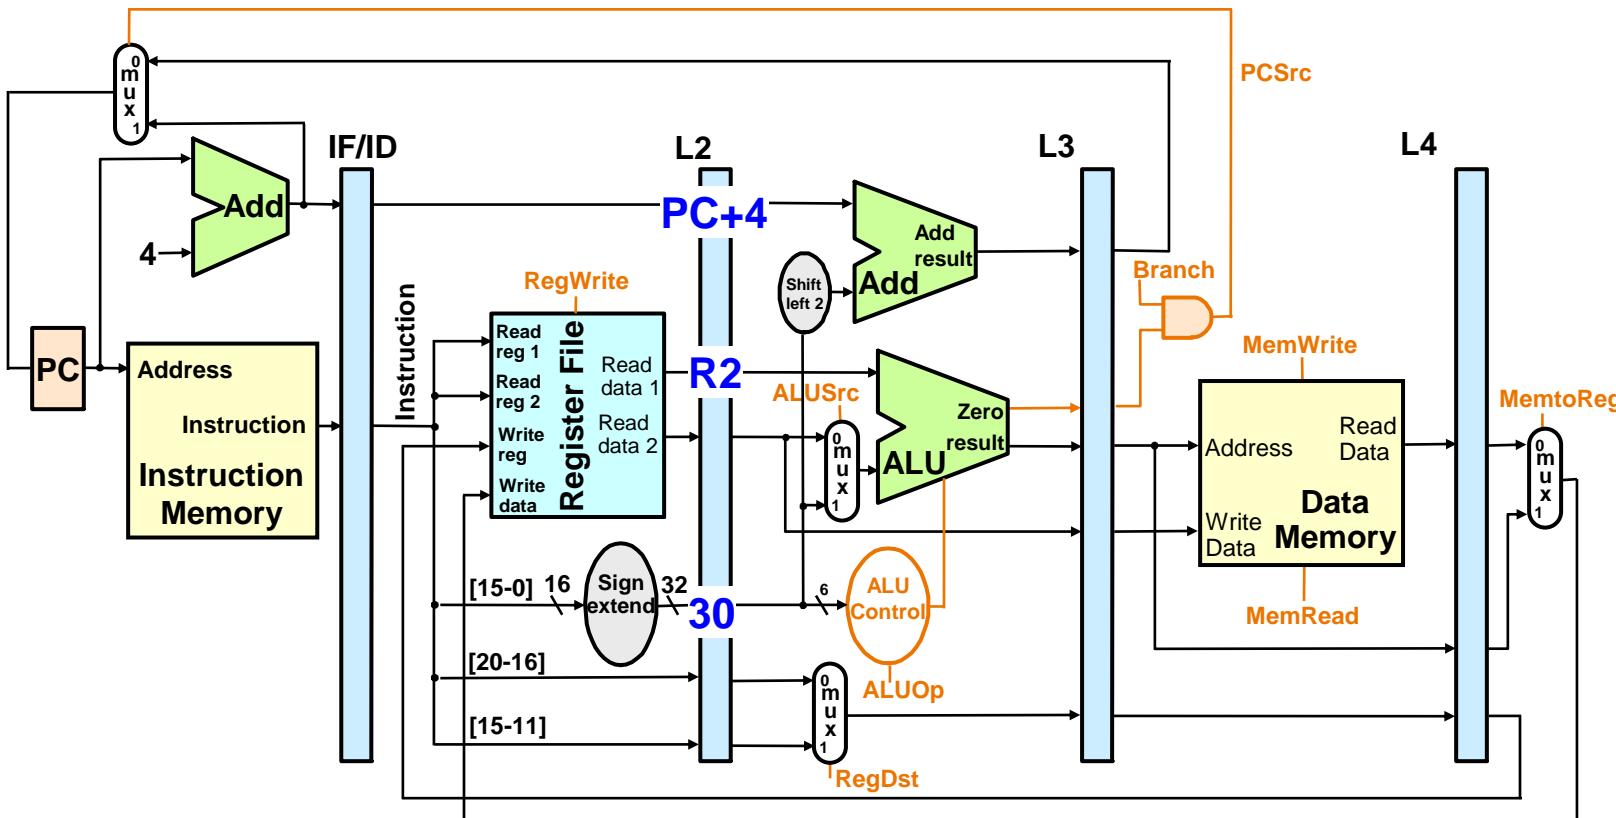
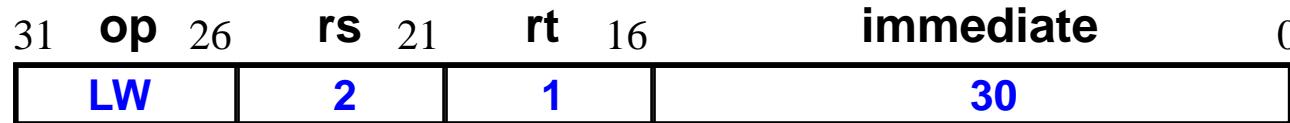
# Pipelined CPU: Load (cycle 1 – Fetch)

**LW R1, (30)R2 ; R1  $\leftarrow$  Mem[R2+30]**



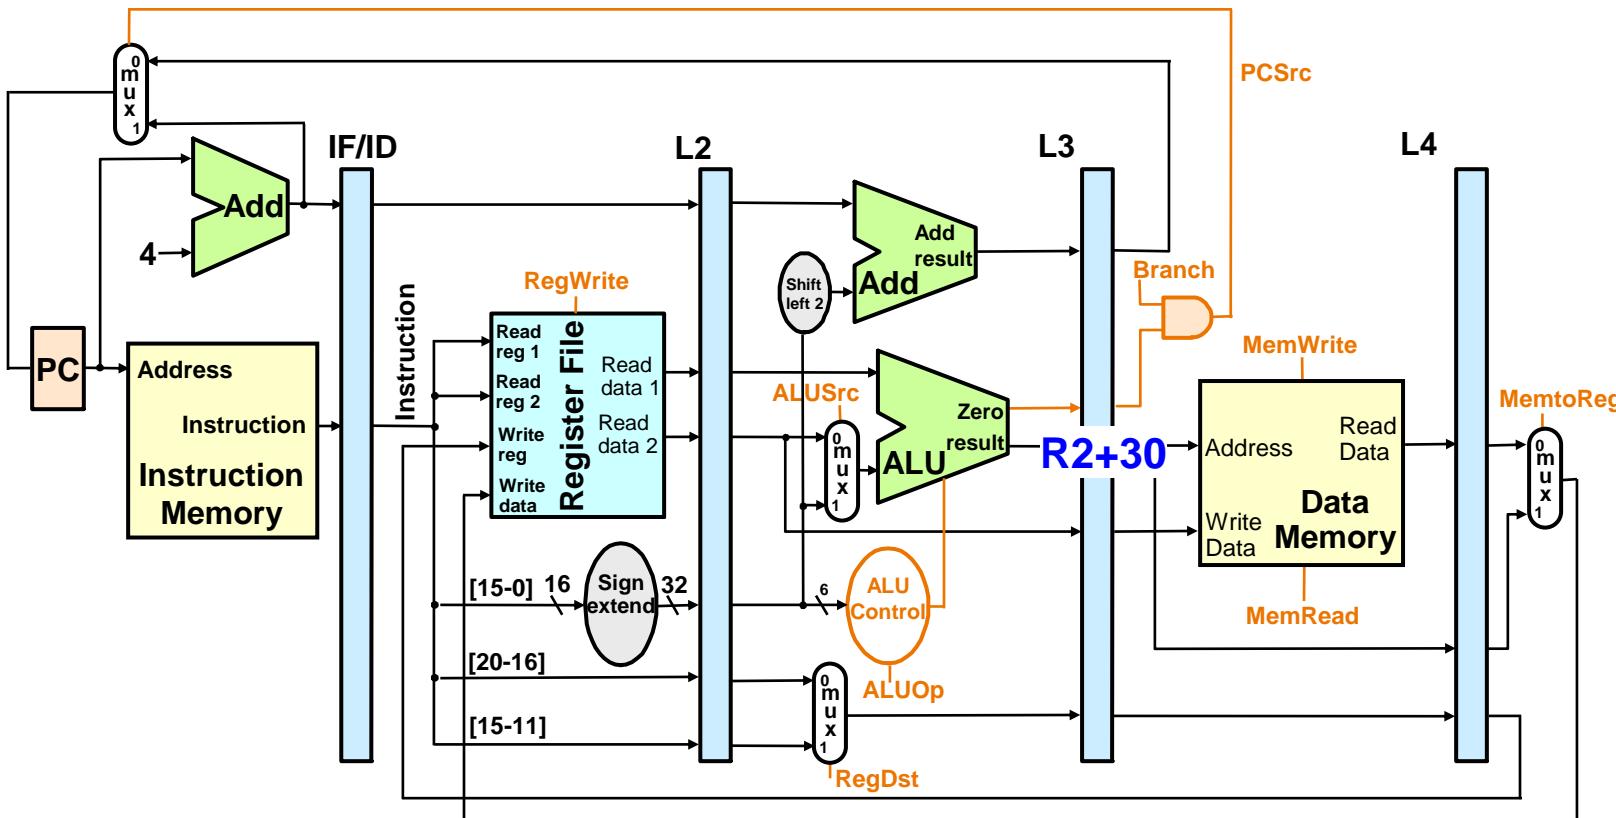
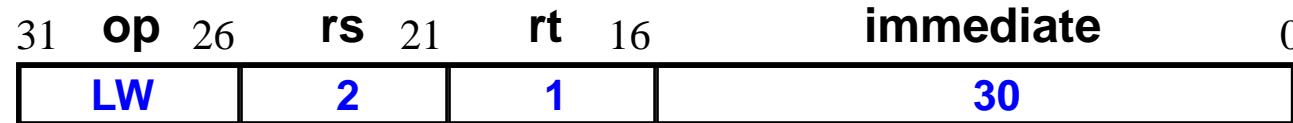
# Pipelined CPU: Load (cycle 2 – Dec)

LW R1, (30)R2 ;  $R1 \leftarrow \text{Mem}[R2+30]$



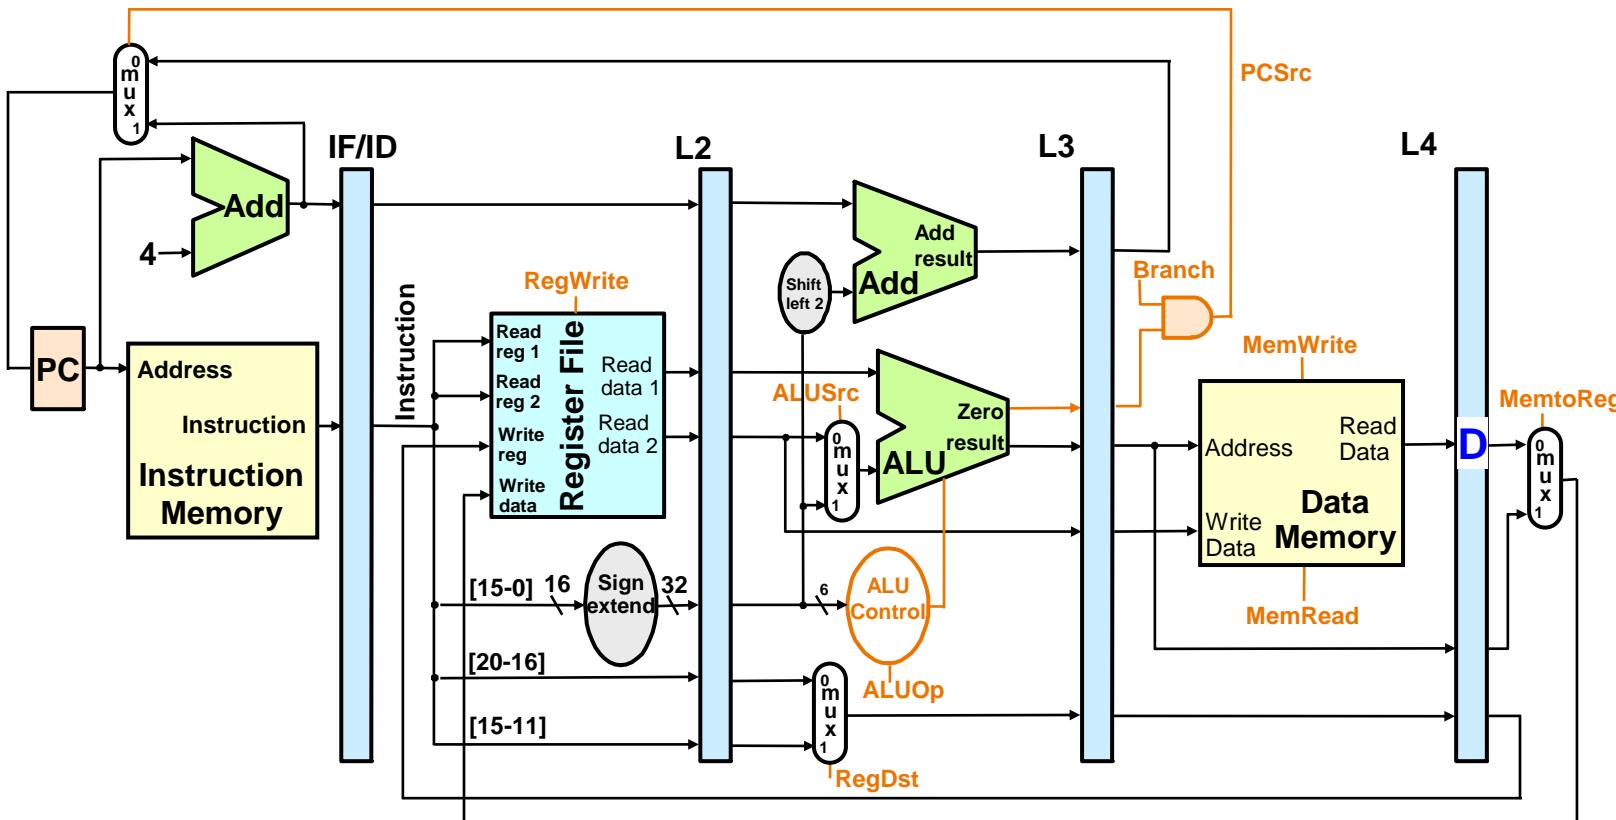
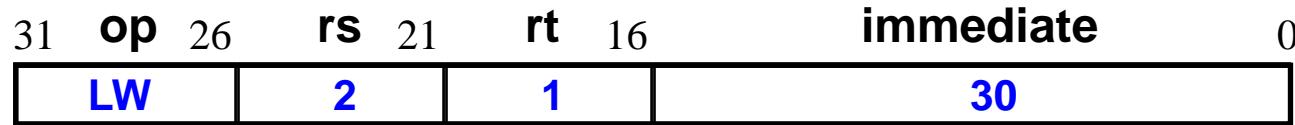
# Pipelined CPU: Load (cycle 3 – Exe)

**LW R1, (30)R2 ; R1  $\leftarrow$  Mem[R2+30]**



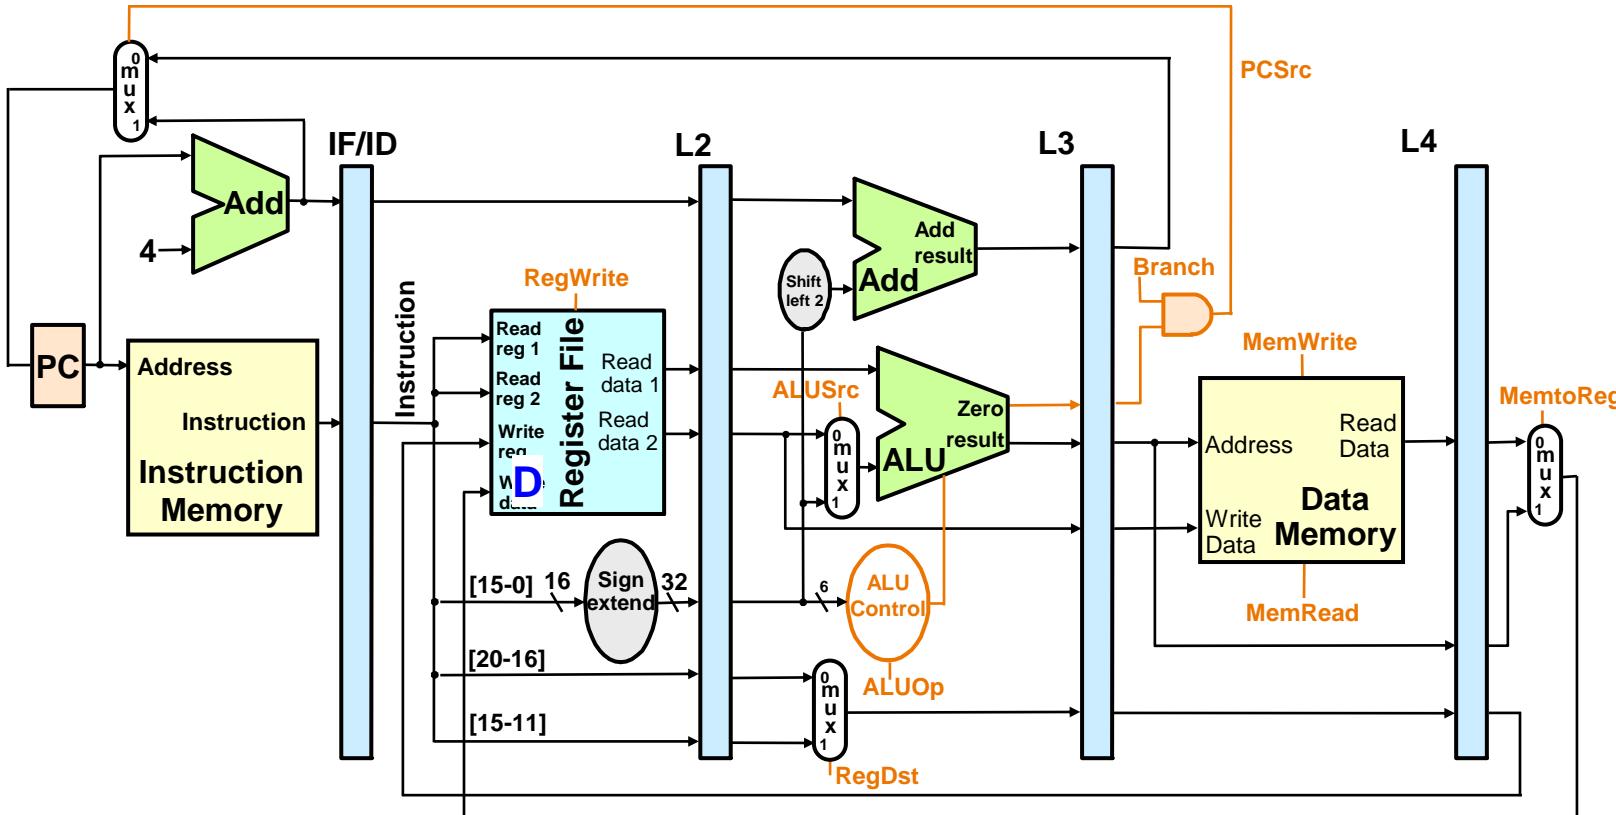
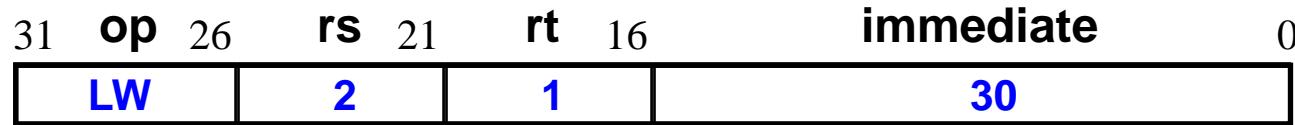
# Pipelined CPU: Load (cycle 4 – Mem)

LW R1, (30)R2 ;  $R1 \leftarrow \text{Mem}[R2+30]$

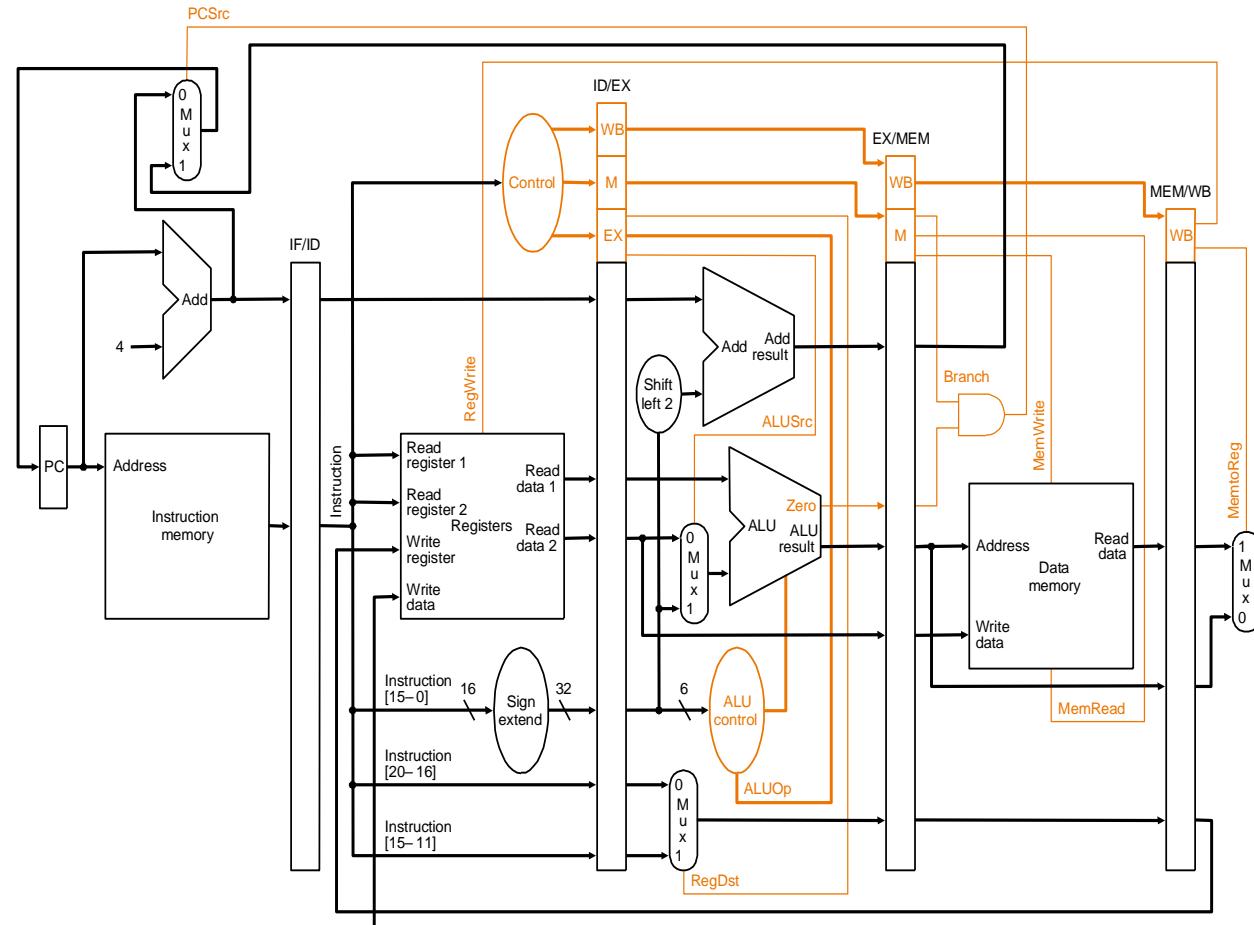


# Pipelined CPU: Load (cycle 5 – WB)

**LW R1, (30)R2 ; R1  $\leftarrow$  Mem[R2+30]**



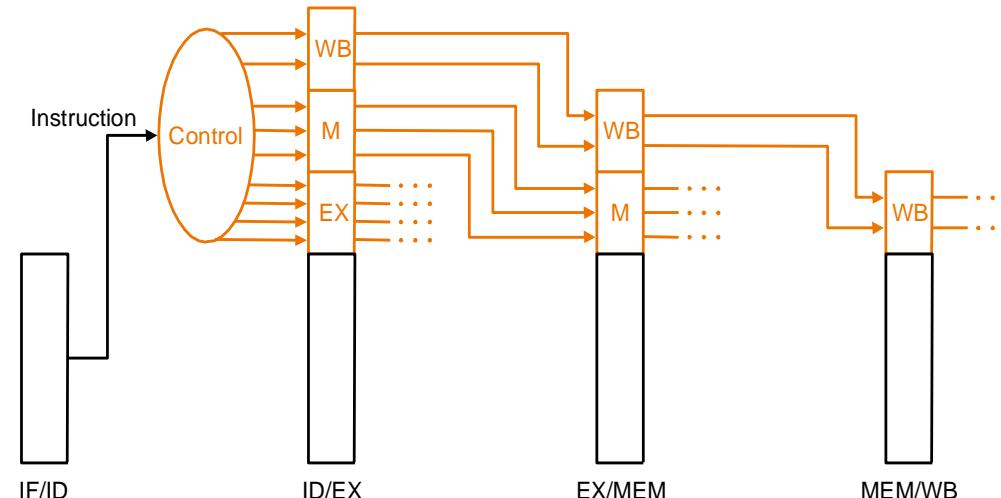
# Datapath with Control



# Multi-Cycle Control

- ◆ Pass control signals along just like the data

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X



# Five Execution Steps

## ◆ Instruction Fetch

- ❖ Use PC to get instruction and put it in the Instruction Register.
- ❖ Increment the PC by 4 and put the result back in the PC.

IR = Memory[PC];

PC = PC + 4;

## ◆ Instruction Decode and Register Fetch

- ❖ Read registers rs and rt
- ❖ Compute the branch address

A = Reg [IR[25-21]] ;

B = Reg [IR[20-16]] ;

ALUOut = PC + (sign-extend(IR[15-0]) << 2) ;

- ❖ We aren't setting any control lines based on the instruction type  
(we are busy "decoding" it in our control logic)

# Five Execution Steps (cont.)

- Execution

ALU is performing one of three functions, based on instruction type:

❖ **Memory Reference: effective address calculation.**

$$\text{ALUOut} = A + \text{sign-extend}(\text{IR}[15-0]);$$

❖ **R-type:**

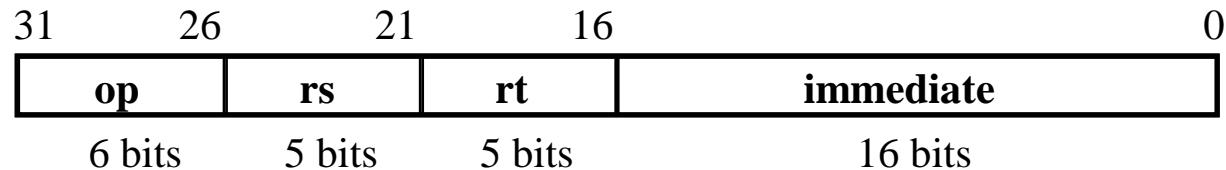
$$\text{ALUOut} = A \text{ op } B;$$

❖ **Branch:**

$$\text{if } (A==B) \text{ PC} = \text{ALUOut};$$

- Memory Access or R-type instruction completion
- Write-back step

# The Store Instruction



◆sw rt, rs, imm16

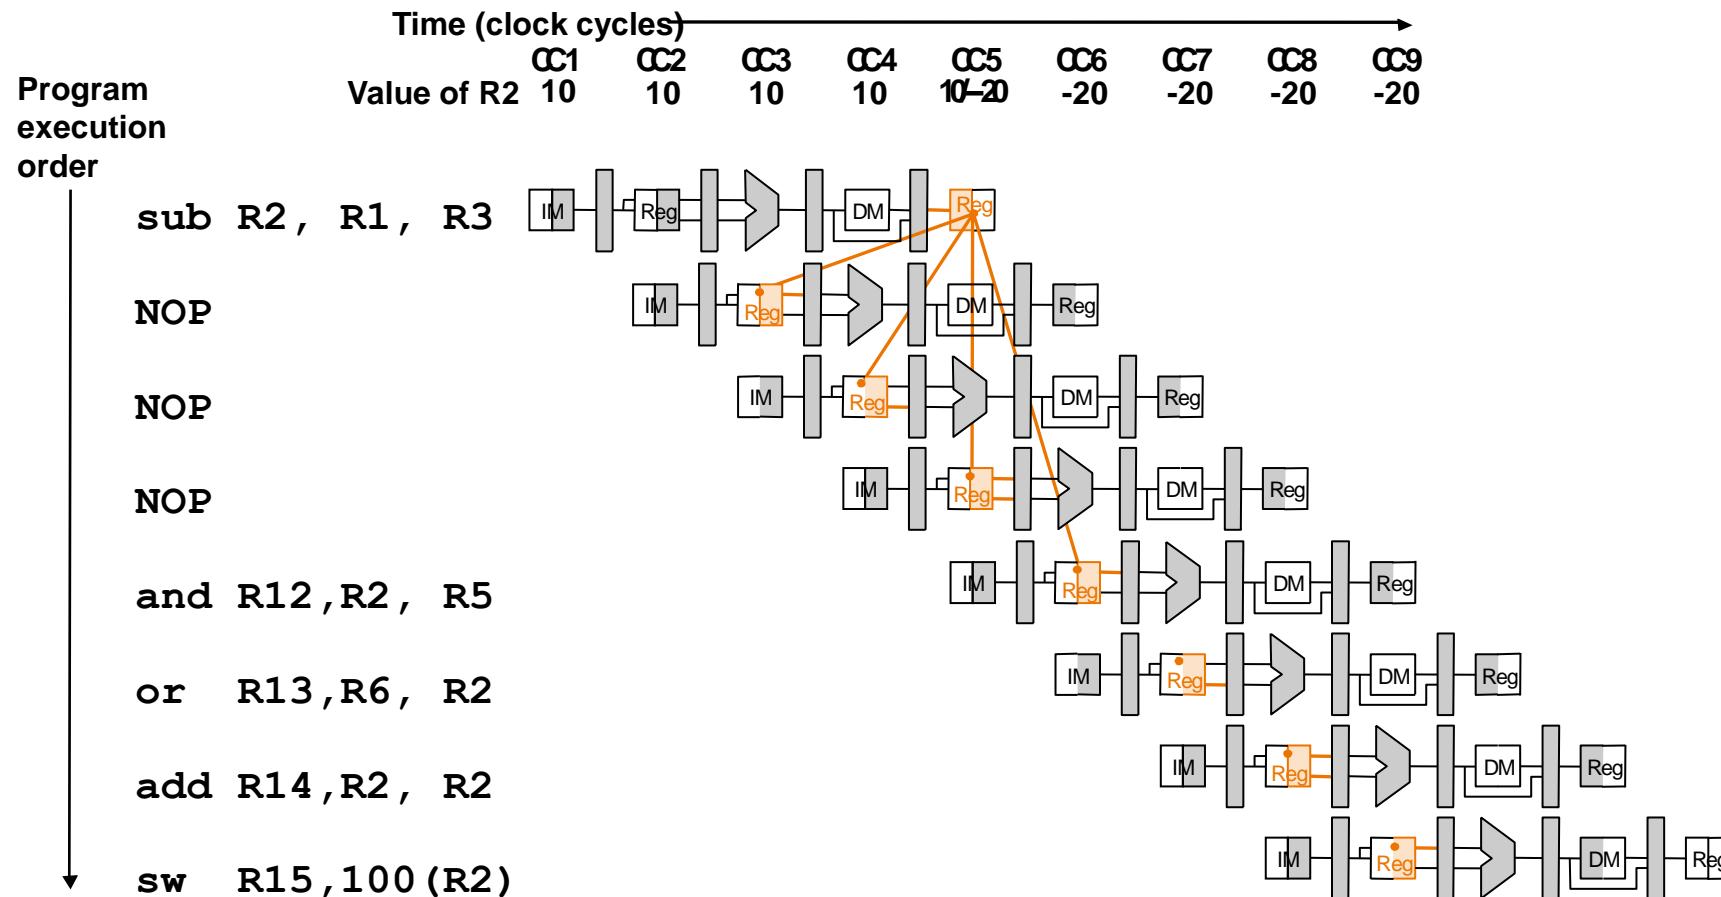
- ❖ mem[PC] Fetch the instruction from memory
- ❖ Addr  $\leftarrow$  R[rs] + SignExt(imm16) Calculate the memory address
- ❖ Mem[Addr]  $\leftarrow$  R[rt] Store the register into memory
- ❖ PC  $\leftarrow$  PC + 4 Calculate the next instruction's address

Mem[Rs + SignExt[imm16]]  $\leftarrow$  Rt

Example: sw rt, rs, imm16

# RAW Hazard: SW Solution

- Have compiler avoid hazards by adding NOP instructions



- Problem: this really slows us down!

# Delayed Branch

- ◆ Define branch to take place AFTER  $n$  following instruction
  - ❖ HW executes  $n$  instructions following the branch regardless of branch is taken or not
- ◆ SW puts in the  $n$  slots following the branch instructions that need to be executed regardless of branch resolution
  - ❖ Instructions that are before the branch instruction, or
  - ❖ Instructions from the converged path after the branch
- ◆ If cannot find independent instructions, put NOP

## Original Code

```
r3 = 23  
R4 = R3+R5  
If (r1==r2) goto x  
R1 = R4 + R5  
x: R7 = R1
```



## New Code

```
If (r1==r2) goto x  
r3 = 23  
R4 = R3 +R5  
NOP  
R1 = R4 + R5  
x: R7 = R1
```

# Delayed Branch Performance

- ◆ Filling 1 delay slot is easy, 2 is hard, 3 is harder
- ◆ Assuming we can effectively fill  $d\%$  of the delayed slots

$$\text{CPI}_{\text{new}} = 1 + 0.2 \times (3 \times (1-d))$$

- ◆ For example, for  $d=0.5$ , we get  $\text{CPI}_{\text{new}} = 1.3$
- ◆ Mixing architecture with micro-arch
  - ❖ New generations requires more delay slots
  - ❖ Cause computability issues between generations