


Introduction to AI – 236501

Problem Solving and Blind Search

[Chapter 3]



Oren Salzman and Sarah Keren

Slides adapted from ai.berkeley.edu by


Dan Klein, Pieter Abbeel and Anca Dragan

and from Shaul Markovitz @ Technion



Today


- ▶ Agents that Plan Ahead
- ▶ Search Problems





Designing Rational Agents


- ▶ An **agent** is an entity that **perceives** and **acts**.
- ▶ A **rational** agent selects actions that **maximize** its (expected) **utility**.
- ▶ Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions





How do we design the rational agent's **decision procedure**?

- ▶ An agent's **decision procedure** maps a world's **state** and a desired **goal** to an **action**
- ▶ We can hardcode the mapping
 $\text{State} \times \text{Goal} \rightarrow \text{Action}$
 - May be an infinite table
- ▶ We can hardcode the mapping
 $\text{Percept} \times \text{Goal} \rightarrow \text{Action}$
 - May (still) be an infinite table
 - This is called a **reflex agent**






Reflex Agents

► Reflex agents:


- Choose action based on current percept (and maybe memory)
- May have memory or a model of the world's current state
- Do not consider the future consequences of their actions
- Consider how the world IS





Planning Agents

- ▶ Planning agents:
 - Ask “what if”
 - Decisions based on (hypothesized) consequences of actions
 - Must have a model of how the world evolves in response to actions
 - Must formulate a goal (test)
 - Consider how the world **WOULD BE**
- ▶ Optimal vs. complete planning
- ▶ Planning vs. replanning






Video of Demo Reflex - Optimal

▶ Reflex agent

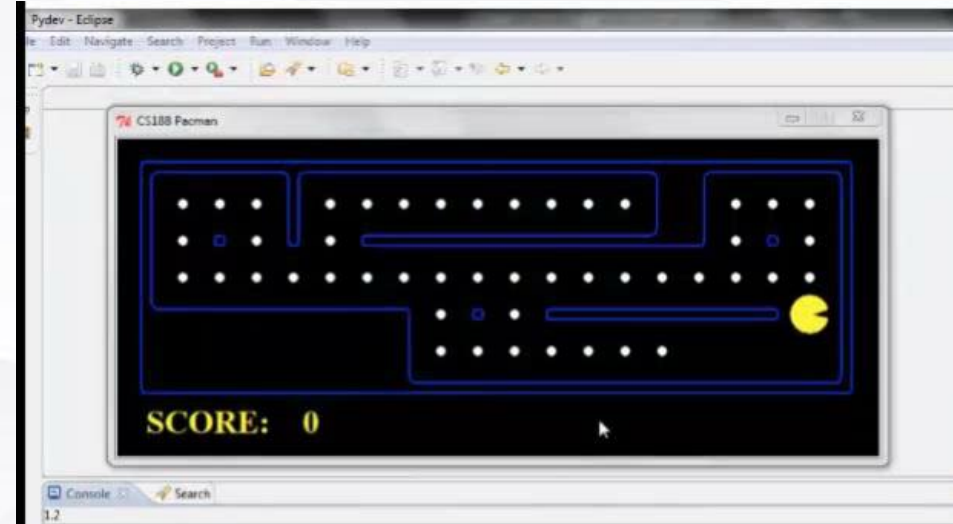
- Move in the direction of the nearest uneaten dot
- You can think ahead all you want, still end up doing the same thing
- Can be rational





Video of Demo Reflex – incomplete


- ▶ Reflex agent
 - Move in the direction of the nearest uneaten dot
 - Can be incomplete
 - No planning
- ▶ Reflex agents not generally optimal, but they can be depending on circumstance.





Video of Demo replaning reflex agent – sub optimal

- ▶ Replanning reflex agent
 - Move in the direction of the nearest uneaten dot
 - Clears board
 - Non-optimal
 - Fast






Video of Demo mastermind– Optimal

► Mastermind agent

- Thinking, thinking, thinking...
- Not just clear NEXT dot, but all dots
- Know to end with dead end so it doesn't have to back up



Problem-Solving Agents



How do we formalize a search problem?







Search Problems

► A **search problem** consists of:

■ A **state space**



■ A **successor** function
(with actions, costs)
or a set of **operators**



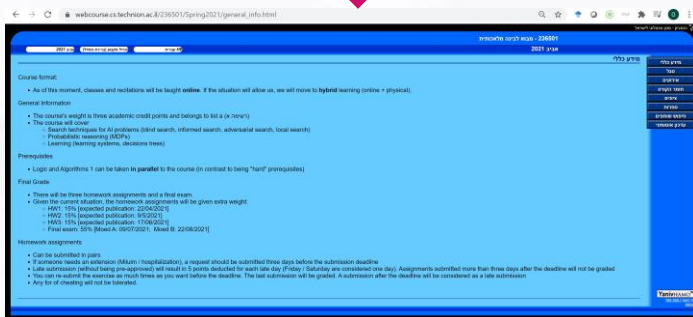
■ A **start state** and a **goal test**

► A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state



Example: searching the internet

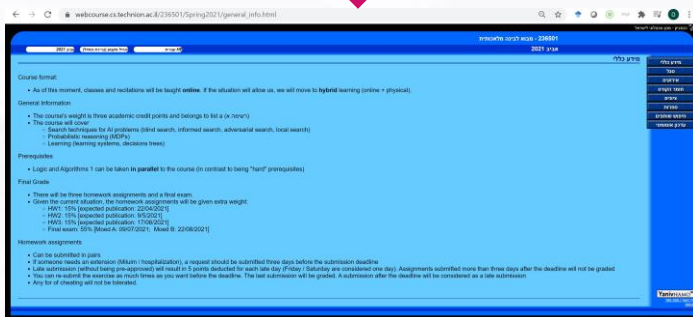
- **Problem:** Given the Technion's homepage, find the series of clicks that will get you to the homepage of Intro To AI





Example: searching the internet

- ▶ **Problem:** Given the Technion's homepage, find the series of clicks that will get you to the homepage of Intro To AI
- ▶ **State space:**





Example: searching the internet

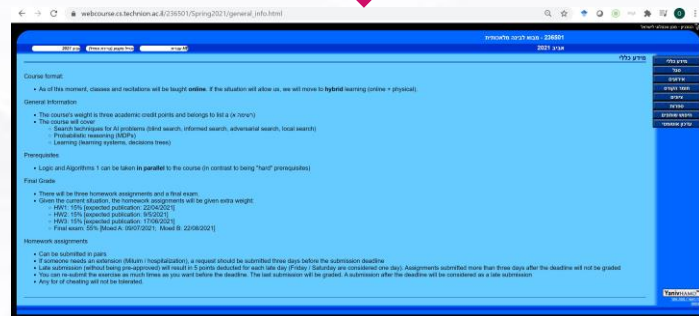
- ▶ **Problem:** Given the Technion's homepage, find the series of clicks that will get you to the homepage of Intro To AI
- ▶ **State space:** all URLs





Example: searching the internet

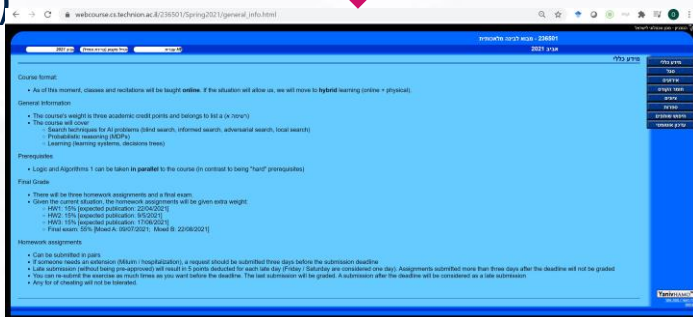
- ▶ **Problem:** Given the Technion's homepage, find the series of clicks that will get you to the homepage of Intro To AI
- ▶ **State space:** all URLs
- ▶ **Successor function:**





Example: searching the internet

- ▶ **Problem:** Given the Technion's homepage, find the series of clicks that will get you to the homepage of Intro To AI
- ▶ **State space:** all URLs
- ▶ **Successor function:** for each page the set of successors is defined by the links available at that page (cost = 1)





Example: searching the internet

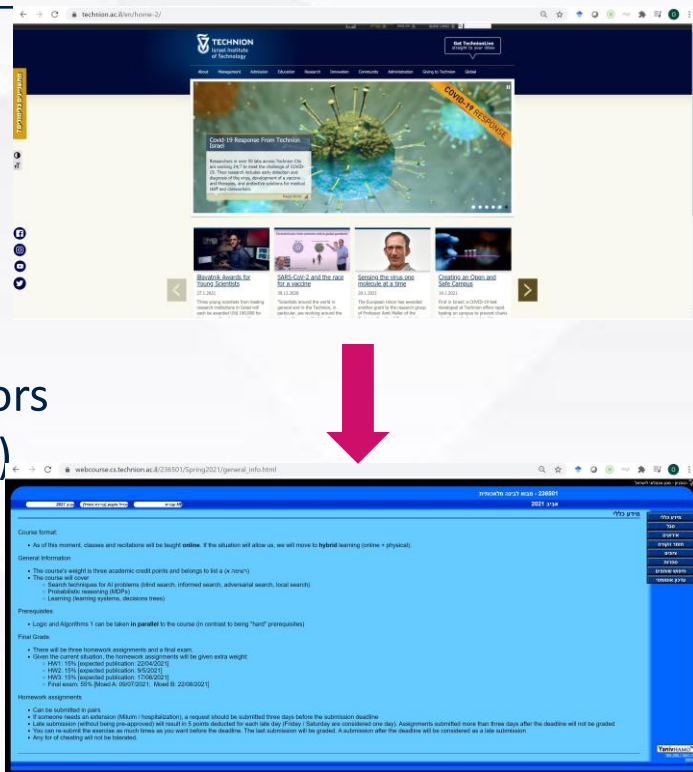
- ▶ **Problem:** Given the Technion's homepage, find the series of clicks that will get you to the homepage of Intro To AI
- ▶ **State space:** all URLs
- ▶ **Successor function:** for each page the set of successors is defined by the links available at that page (cost = 1)
- ▶ **Start state:**
- ▶ **Goal state:**





Example: searching the internet


- ▶ **Problem:** Given the Technion's homepage, find the series of clicks that will get you to the homepage of Intro To AI
- ▶ **State space:** all URLs
- ▶ **Successor function:** for each page the set of successors is defined by the links available at that page (cost = 1)
- ▶ **Start state:** <https://www.cs.technion.ac.il/>
- ▶ **Goal state:** <https://webcourse.cs.technion.ac.il/236501/>





Search Problems Are Models

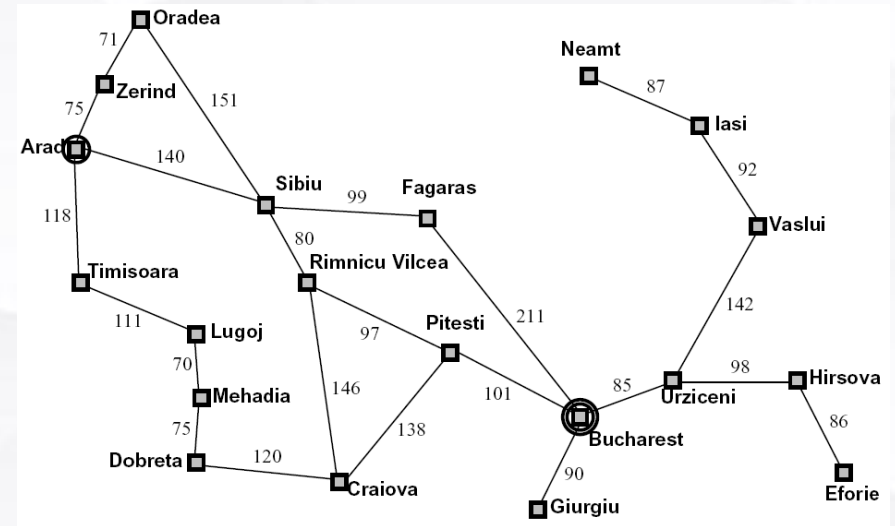
- ▶ Models aren't perfect
 - Too detailed, you can't solve
 - Not detailed enough, doesn't solve





Example: Traveling in Romania


- ▶ State space:
 - Cities
- ▶ Successor function:
 - Roads: Go to adjacent city with cost = distance
- ▶ Start state:
 - Arad
- ▶ Goal test:
 - Is state == Bucharest?
- ▶ Solution?





What's in a State Space?

The **world state** includes every last detail of the environment



A **search state** keeps only the details needed for planning (abstraction)

► **Problem: Path finding**

- States: (x,y) location
- Actions: NSEW
- Successor: update location only
- Goal test: is $(x,y)=\text{END}$

► **Problem: Eat-All-Dots**

- States: $\{(x,y), \text{dot booleans}\}$
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false




State Space Sizes?

World state:

- Agent positions: $12 \times 10 = 120$
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW

How many

- World states?
 $120 \times (2^{30}) \times (12^2) \times 4$
- States for path finding?
120
- States for eat-all-dots?
 $120 \times (2^{30})$





Some assumptions

- ▶ Set of actions is **finite** and bounded by some constant
 - In practice, they are often parametrized and thus infinite (e.g., steering angle of a car)
 - In such settings we can **discretize** the set of actions
- ▶ State of agent is **observable**
 - i.e., agent knows its state
- ▶ Actions are **deterministic**
 - i.e., agent knows what will happen when it does something
- ▶ We will relax some of these assumptions along the course

State Space Search Problem : Classical Planning

► Formal definition:


- a finite and discrete **state space** S
- a known **initial state** $s_0 \in S$
- a non-empty set of goal states $S_G \subseteq S$ (or a **goal test**)
- set of **actions/operators** $A(s) \in A$ applicable in each state $s \in S$
- A **successor** function – a **deterministic** state transition function $f(a, s)$ such that $s' = f(s, a)$ stands for the state resulting from applying an applicable action a at state s
- A **cost function** $c(a, s)$ associating a **positive cost** to applying action a in s

► Objective

- Find a plan, a sequence of applicable actions, that leads from the initial state to a goal state.
- Optimality criteria - typically interested in shortest/ cheapest plan (but other considerations are possible)

Factored State Space Definition: Using Features

- ▶ Typically, it is not possible to explicitly specify the **state space** S
- ▶ Instead, **feature set** X is used to represent the state space such that a state is described via a set of random variables $X = x_1, \dots, x_n$ such that each variable takes on values in some finite domain $\text{Dom}(x_i)$.
- ▶ Can be Boolean or multi-valued
- ▶ Example:
 - Position of block a
 - Position of block c
 - ...
 - Is gripper empty



Formalizing a search problem



Example: Missionaries and cannibals problem

► Problem:

- 3 missionaries and 3 cannibals must cross a river using a boat that carries at most 2 people
- If there are missionaries present on the bank, they cannot be outnumbered by cannibals (if they were, the cannibals would eat the missionaries)
- The boat cannot cross the river by itself with no people on board

► How do we represent the search problem?





Example: Missionaries and cannibals problem

► Take (1):

- A state is a 7-tuple $\langle L_1, L_2, L_3, L_4, L_5, L_6, B \rangle$
- $L_i \in \{\text{Left, Right, Boat}\}$ location of person i
- $B \in \{\text{Left, Crossing, Right}\}$ location of boat
- Start state: $\langle \text{Right, Right, Right, Right, Right, Right, Right} \rangle$
- Goal state: $\langle \text{Left, Left, Left, Left, Left, Left, Left} \rangle$

► How many states do we have?





Example: Missionaries and cannibals problem

► Take (1):

- A state is a 7-tuple $\langle L_1, L_2, L_3, L_4, L_5, L_6, B \rangle$
- $L_i \in \{\text{Left}, \text{Right}, \text{Boat}\}$ location of person i
- $B \in \{\text{Left}, \text{Crossing}, \text{Right}\}$ location of boat
- Start state: $\langle \text{Right}, \text{Right}, \text{Right}, \text{Right}, \text{Right}, \text{Right}, \text{Right} \rangle$
- Goal state: $\langle \text{Left}, \text{Left}, \text{Left}, \text{Left}, \text{Left}, \text{Left}, \text{Left} \rangle$

► How many states do we have? $3^7 = 2187$

► Unnecessary information

- We don't care **which** missionary / cannibal is on which side
- We don't care about the state of the boat while crossing



Example: Missionaries and cannibals problem

► Take (2):

- A state is a 3-tuple $\langle M_L, M_C, B \rangle$
- $M_L \in \{0, \dots, 3\}$ number of missionaries on left bank
- $M_C \in \{0, \dots, 3\}$ number of cannibals on left bank
- $B \in \{\text{Left}, \text{Right}\}$ location of boat
- Start state: $\langle 0, 0, \text{Right} \rangle$
- Goal state: $\langle 3, 3, \text{Left} \rangle$

► How many states do we have?





Example: Missionaries and cannibals problem

► Take (2):

- A state is a 3-tuple $\langle M_L, M_C, B \rangle$
- $M_L \in \{0, \dots, 3\}$ number of missionaries on left bank
- $M_C \in \{0, \dots, 3\}$ number of cannibals on left bank
- $B \in \{\text{Left}, \text{Right}\}$ location of boat
- Start state: $\langle 0, 0, \text{Right} \rangle$
- Goal state: $\langle 3, 3, \text{Left} \rangle$


► How many states do we have? $4 \times 4 \times 2 = 32$



Search algorithms - preliminaries




State Space Graphs and Search Trees






State-Space Graphs

- ▶ State-space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- ▶ In a state-space graph, each state occurs only once!
- ▶ We can rarely build this full graph in memory (it's too big), but it's a useful idea
 - The graph may be **infinite**
 - The graph is given **implicitly**





Search Trees




► A search tree:

- A “what if” tree of plans and their outcomes
- The start state is the root node
- Children correspond to successors
- Nodes show states, but correspond to PLANS that achieve those states
- For most problems, we can never actually build the whole tree



State Space Graphs vs. Search Trees


State Space Graph



Each **NODE** in the search tree is an entire **PATH** in the state space graph.

We construct both on demand – and we construct as little as possible.


Search Tree





State Space Graphs vs. Search Trees

Consider this 4-state graph:




How big is its search tree (from S)?






State Space Graphs vs. Search Trees

Consider this 4-state graph:




How big is its search tree (from S)?



Important: Lots of repeated structure in the search tree!

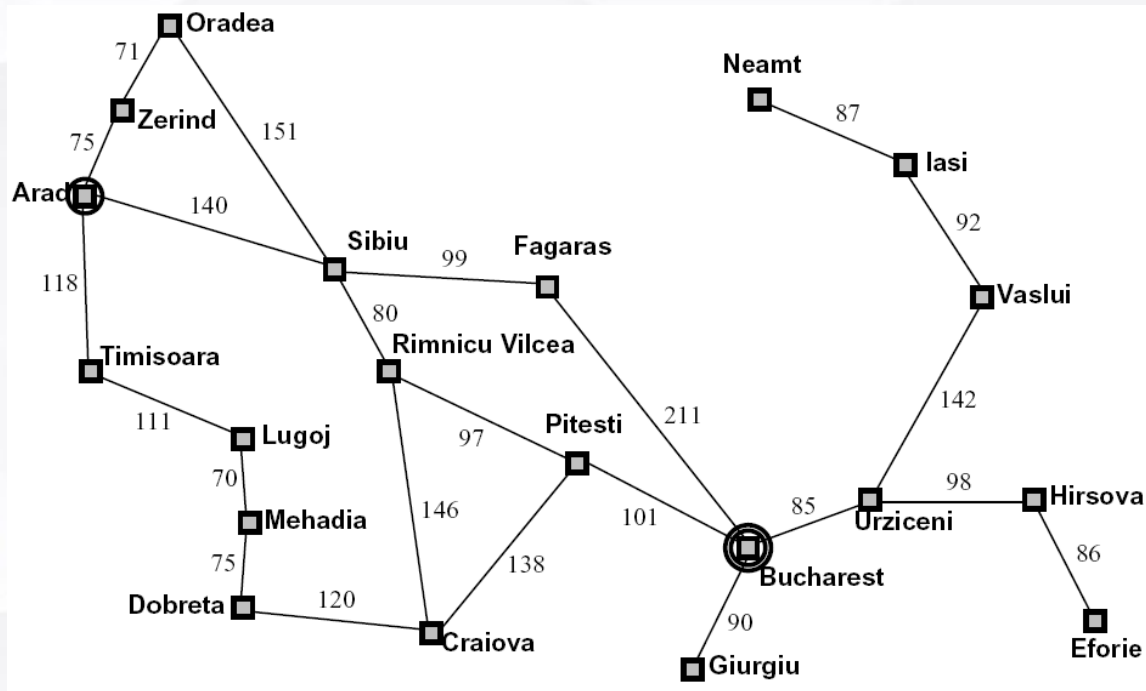


Tree Search



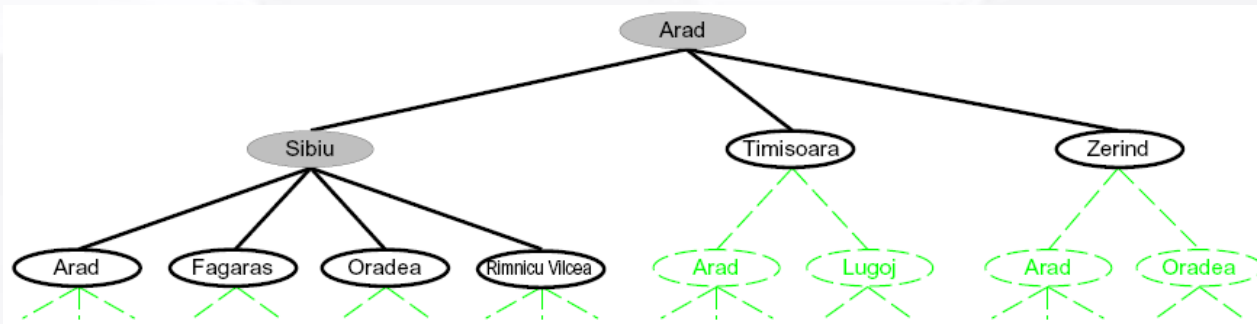


Search Example: Romania





Searching with a Search Tree



► Search:

- Expand out potential plans (tree nodes)
- Maintain a **fringe** of partial plans under consideration
- Try to expand as few tree nodes as possible



General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```


► Important ideas:

- Fringe
- Expansion
- Exploration strategy

► Main question: which fringe nodes to explore?





Example: Tree Search





Example: Tree Search



~~S~~
~~S → d~~
S → e
S → p
S → d → b
S → d → c
~~S → d → e~~
S → d → e → h
~~S → d → e → r~~
~~S → d → e → r → f~~
S → d → e → r → f → c
~~S → d → e → r → f → G~~



Measuring performance of a search algorithm

- ▶ We care about two main (conflicting) **metrics**
 - Algorithm's **resources**: running time, memory consumption
 - **Quality** of solution: Optimal, bounded suboptimal, unbounded
- ▶ Typically, the more resources we invest, the higher quality the solution we can obtain



Algorithm's resources

▶ Running time

- What we actually care about
- Problematic metric – sensitive to hardware, compiler, optimization, implementation language etc.
- Alternative measure: The number of **nodes expanded**



Algorithm's resources

▶ Running time

- What we actually care about
- Problematic metric – sensitive to hardware, compiler, optimization, implementation language etc.
- Alternative measure: The number of **nodes expanded**

▶ Memory consumption

- **Critical resource** – may deem a search algorithm worthless
- An algorithm that expands **1,000,000 nodes / second** and uses **10 bytes / node** can run out of memory in less than **2 minutes**



Quality of solution

- ▶ We will want to minimize the **cost** of a solution
- ▶ E.g., plan the path for a robot; action costs correspond to the robot's energy consumption; minimize total energy consumption





General Tree Search (recap)


```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

► Important ideas:

- Fringe
- Expansion
- Exploration strategy

► Main question: which fringe nodes to explore?

Search Algorithm Properties



Search Algorithm Properties

▶ **Complete:** Guaranteed to find a solution if one exists?

▶ **Optimal:** Guaranteed to find the least cost path?

▶ Time complexity?


▶ Space complexity?

▶ Cartoon of search tree:

- **b** is the branching factor
- **m** is the maximum depth
- solutions at various depths

▶ Number of nodes in entire tree?

- $1 + b + b^2 + \dots + b^m = O(b^m)$



Breadth-First Search







Breadth-First Search

Strategy: expand a shallowest node first

Implementation: Fringe is a FIFO queue





Breadth-First Search (BFS) Properties

What nodes does BFS expand?

- Processes all nodes above shallowest solution
- Let depth of shallowest solution be s
- Search takes time $O(b^s)$

How much space does the fringe take?


- Has roughly the last tier, so $O(b^s)$

Is it complete?

- s must be finite if a solution exists

Is it optimal?

- Only if costs are all 1 (more on costs later)





BFS Memory requirements

- Assume that BFS expands **1,000,000 nodes / sec** for a problem with branching factor **b=10** and each node uses **1000 bytes node**

Depth	No. of nodes	Time	Memory
2	111	0.1 ms	11 KB
4	11,111	11ms	11 MB
6	$\sim 10^6$	1 sec	~ 1 GB
8	$\sim 10^8$	100 sec	~ 100 GB
10	$\sim 10^{10}$	3 hours	~ 10 TB
12	$\sim 10^{12}$	11 days	~ 1 PB
14	$\sim 10^{14}$	3 years	~ 100 PB

Time requirements are reasonable

Memory requirements are not

Depth-First Search

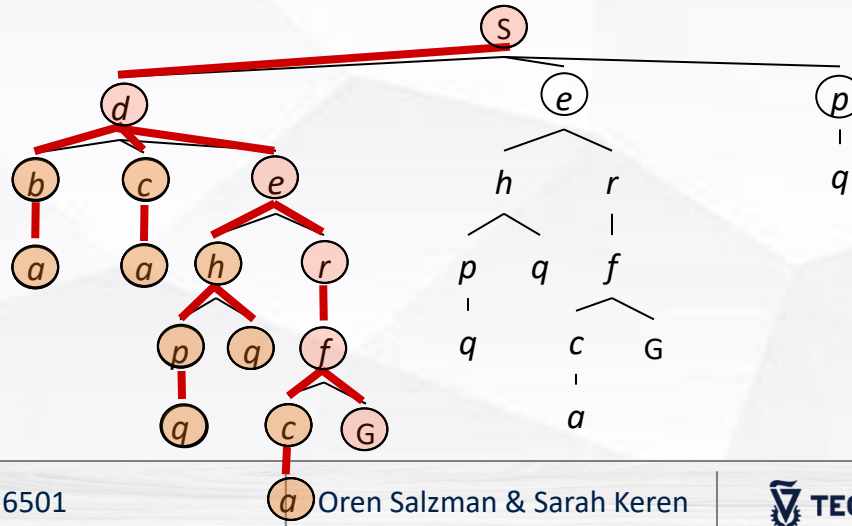





Depth-First Search

Strategy: expand a deepest node first


Implementation: Fringe is a LIFO stack





Depth-First Search (DFS) Properties

- ▶ What nodes DFS expand?
 - Some left prefix of the tree.
 - Could process the whole tree!
 - If m is finite, takes time $O(b^m)$
- ▶ How much space does the fringe take?
 - Only has siblings on path to root, so $O(bm)$
- ▶ Is it complete?
 - m could be infinite
 - To avoid this, we can call the algorithm with a bound L on the maximum depth – this is called **DFS-L**
- ▶ Is it optimal?
 - No, it finds the “leftmost” solution, regardless of depth or cost






Depth-First Search Variants

► DFS-L

- Same as **DFS** but with a maximum tree depth **L**

► DFS-L-BT

- Same as **DFS-L** but with the following modification:
- At every node, we run an operator **O** over the same node
- When we backtrack, we undo the previous operator
- We only need to store one node in memory (with some additional book keeping)




BFS vs. DFS






Quiz: DFS vs BFS






Video of Demo Maze Water DFS/BFS (part 1)





Video of Demo Maze Water DFS/BFS (part 2)





BFS vs. DFS (-L)

	BFS	DFS	DFS-L
Complete	YES	NO*	NO**
Optimal [#]	YES	NO	NO
Memory	Exponential in solution length	Infinite	Linear in L



BFS vs. DFS (-L)

	BFS	DFS	DFS-L
Complete	YES	NO*	NO**
Optimal [#]	YES	NO	NO
Memory	Exponential in solution length	Infinite	Linear in L

* If tree is unbounded, can run infinitely even if a solution exists

** If **L** is less than the solution depth, no solution is returned even though one exists


[#] Assuming uniform cost edges



Empirical demonstration

- ▶ Python implementation by Shaul Markovitz of BFS, DFS-L and DFS-L-BT
- ▶ Random **3X3** puzzles whose solution-lengths range between **6** and **22** (**20** problem instances per solution length)
- ▶ DFS-L was run with **L=22** (for all instances)

5	2	7
8	4	
1	3	6




1	2	3
4	5	6
7	8	



Empirical demonstration


- **Expanded nodes** as a function of **problem depth**






Empirical demonstration

► **Expanded nodes** as a function of **problem depth**




► **Time** as a function of **problem depth**



Empirical demonstration

- **Solution length** as a function of **problem depth**





Empirical demonstration

► Memory footprint as a function of problem depth

