



Slides adapted from ai.berkeley.edu by

Dan Klein, Pieter Abbeel and Anca Dragan

and from Shaul Markovitz @ Technion

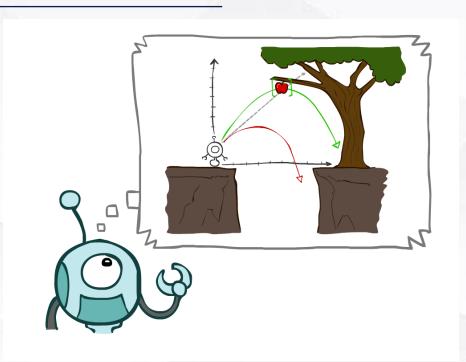


The Henry and Marilyn Taub

Faculty of Computer Science



- Uninformed Search Methods
  - Iterative deepening
  - **Uniform-Cost Search**
- ▶ Informed Search methods
  - **Greedy Best-First Search**
  - **A**\*
  - Heuristics
  - IDA\*





#### **General Tree Search (recap)**

function Tree-Search (problem, strategy) returns a solution, or failure initialize the search tree using the initial state of *problem* loop do if there are no candidates for expansion then return failure choose a leaf node for expansion according to strategy if the node contains a goal state then return the corresponding solution else expand the node and add the resulting nodes to the search tree end

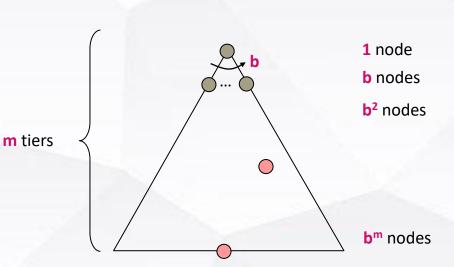
- Important ideas:
  - Fringe
  - Expansion
  - **Exploration strategy**
- Main question: which fringe nodes to explore?



## .ıll

#### **Search Algorithm Properties**

- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of search tree:
  - b is the branching factor
  - m is the maximum depth
  - solutions at various depths
- Number of nodes in entire tree?
  - $1 + b + b^2 + \ldots + b^m = O(b^m)$



## Iterative deepening



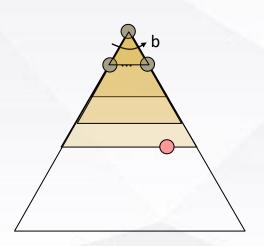




#### **Iterative Deepening**

- ▶ Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run a DFS with depth limit 1. If no solution...
  - Run a DFS with depth limit 2. If no solution...
  - Run a DFS with depth limit 3. .....

- Isn't that wastefully redundant?
  - Generally most work happens in the lowest level searched, so not so bad!





#### **Iterative Deepening - properties**

- Completeness if a solution of length L exists, ID will find it in the L'th iteration
- ▶ Optimality if ID found a solution in the L'th iteration, there cannot be a solution of length <L (assuming uniform cost edges)
- ▶ Memory complexity linear in solution depth (similar to DFS-L which is linear in the bound L)
- ▶ Time complexity
  - If solution has length L, the search tree will have L+1 layers
  - If the branching factor is b the i'th layer will have  $b^i$  nodes

$$t = \sum_{i=0}^{L}$$

$$(L+1-i)$$

X



no. of times i'th layer is expanded

no. of nodes in i'th layer



#### **Iterative Deepening vs. BFS - example**

- Assume that b=4 and L=20
- ▶ Time (BFS):  $\sum_{i=0}^{20} 4^i = \frac{4^{2i}-1}{4-1} = 1,466,015,503,701$
- ▶ Time (ID):  $\sum_{i=0}^{20} (20+1-i) \cdot 4^i = 1,954,687,338,261$
- ▶ **Memory (**BFS**):** equals to Time (BFS) i.e., 1, 466, 015, 503, 701
- Memory(ID):  $4 \cdot 20 = 80$

Ratio -

$$\frac{\text{Time(ID)}}{\text{Time(BFS)}} = 1.333$$

$$\frac{\text{Memory(BFS)}}{\text{Memory(ID)}} = 1.8 \cdot 10^{10}$$

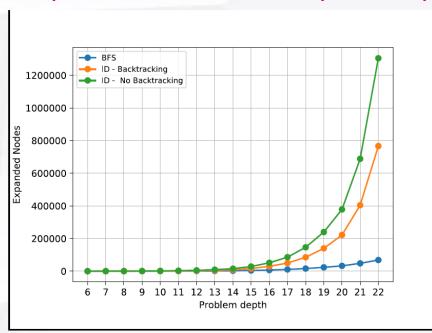
= > reduction by a factor of  $1.8 \cdot 10^{10}$  in memory for 33% more time





#### **Empirical demonstration**

**Expanded nodes** as a function of **problem depth** 

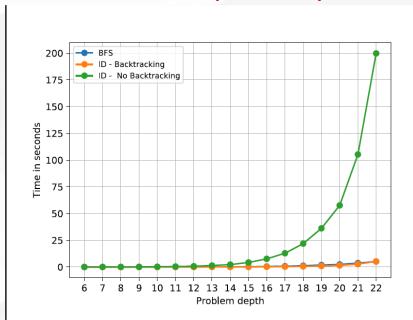


- ▶ ID expands more nodes because
  - It repeats expanding nodes in different depth levels
  - It repeats expanding nodes because it does not perform a graph search (i.e., eliminating expanding duplicate nodes)



#### **Empirical demonstration**

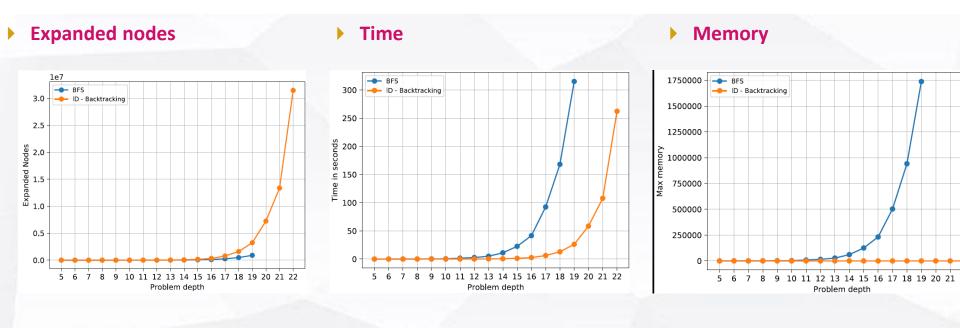
Time as a function of problem depth



- ▶ ID requires much more time due to the additional number of nodes expanded
- ID-BT is much more efficient and has running time comparable to BFS

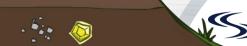


#### Empirical demonstration – 4X4 puzzle, BFS vs ID-BT



### Uniform Cost Search

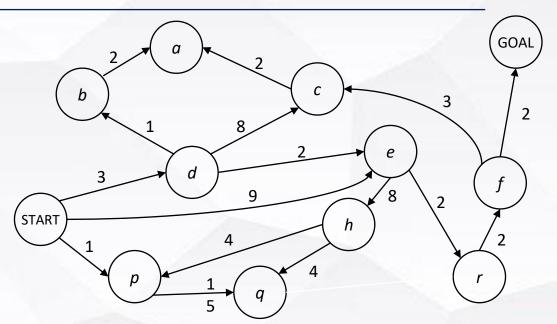








#### **Cost-Sensitive Search**



BFS finds the shortest path in terms of number of actions.

It does not find the least-cost path.

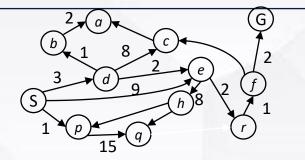
We will now cover a similar algorithm which does find the least-cost path.

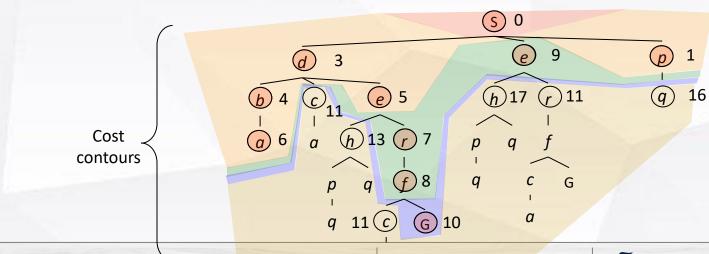


#### **Uniform Cost Search**

Strategy: expand a cheapest node first:

Fringe is a priority queue (priority: cumulative cost)





<sup>a</sup> Oren Salzman & Sarah Keren

TECHNION |

The Henry and Marilyn Taub

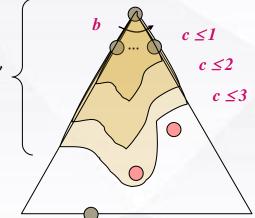
Faculty of Computer Science



#### **Uniform Cost Search (UCS) Properties**

- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\varepsilon$ , then the "effective depth" is roughly  $C^*/\varepsilon$
  - Takes time  $O(b^{C^*/\varepsilon})$  (exponential in effective depth)

C\*/€ "tiers"



- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C^*/\varepsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
  - Yes! (Proof next lecture via A\*)



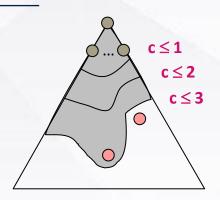
#### **Uniform Cost Issues**

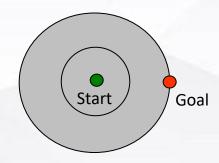
Remember: UCS explores increasing cost contours

▶ The good: UCS is complete and optimal!

- The bad:
  - Explores options in every "direction"
  - No information about goal location

We'll fix that soon!







#### **The One Queue**

- ▶ BFS FIFO Queue
- ▶ DFS LIFO Queue (Stack)

- ▶ UCTS Priority Queue => log(n) overhead
- ▶ All these search algorithms are the same except for fringe strategies
  - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
  - Practically, for DFS and BFS, you can avoid the log(n) overhead from an actual priority queue, by using stacks and queues
  - Can even code one implementation that takes a variable queuing object



#### **General Tree Search (recap)**

function Tree-Search (problem, strategy) returns a solution, or failure initialize the search tree using the initial state of *problem* loop do if there are no candidates for expansion then return failure choose a leaf node for expansion according to strategy if the node contains a goal state then return the corresponding solution else expand the node and add the resulting nodes to the search tree end

- Important ideas:
  - Fringe
  - Expansion
  - **Exploration strategy**
- Main question: which fringe nodes to explore?







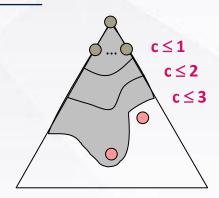
#### **Uniform Cost Issues**

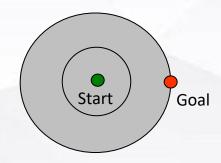
Remember: UCS explores increasing cost contours

▶ The good: UCS is complete and optimal!

- The bad:
  - Explores options in every "direction"
  - No information about goal location

We'll fix that soon!





## Informed Search

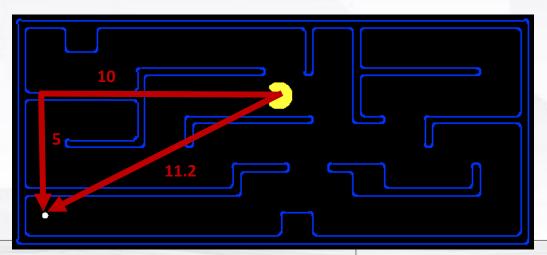


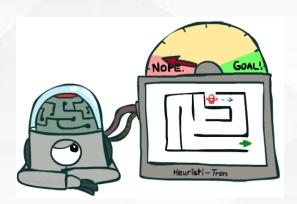


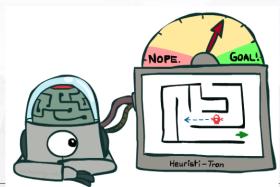


#### **Search Heuristics**

- A heuristic is:
  - A function that estimates how close a state is to a goal
  - Designed for a particular search problem
  - Examples (path finding): Manhattan distance,
     Euclidean distance





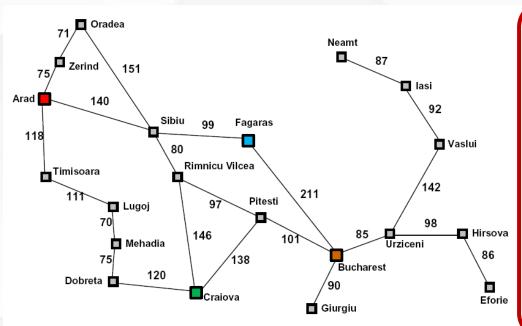








#### **Example: Heuristic Function**



| Straight-line distance to Bucharest |     |
|-------------------------------------|-----|
| Arad                                | 366 |
| Bucharest                           | 0   |
| Craiova                             | 160 |
| Dobreta                             | 242 |
| Eforie                              | 161 |
| Fagaras                             | 178 |
| Giurgiu                             | 77  |
| Hirsova                             | 151 |
| Iasi                                | 226 |
| Lugoj                               | 244 |
| Mehadia                             | 241 |
| Neamt                               | 234 |
| Oradea                              | 380 |
| Pitesti                             | 98  |
| Rimnicu Vilcea                      | 193 |
| Sibiu                               | 253 |
| Timisoara                           | 329 |
| Urziceni                            | 80  |
| Vaslui                              | 199 |
| Zerind                              | 374 |
| \                                   |     |

h(x)



## Greedy best-first Search





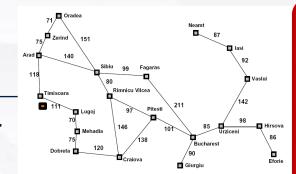


The Henry and Marilyn Taub Faculty of Computer Science



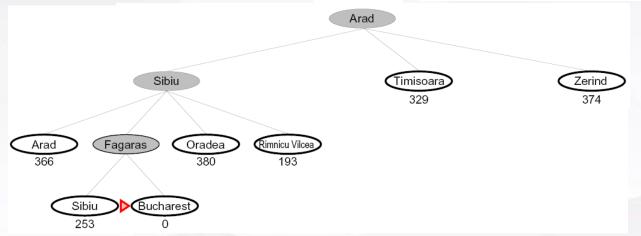
#### **Greedy Search**

Expand the node that seems closest...





h(x)



- ▶ Is it optimal?
  - No. Resulting path to Bucharest is not the shortest!



#### **Greedy Search – pseudo code**

Greedy best first search (search problem - P)

- ▶ OPEN <- make\_node(P.start, NIL, h(P.start)) //order according to h-value</p>
- ▶ CLOSE <- {}</p>
- While OPEN  $\neq \emptyset$ 
  - n <- OPEN.pop min()</pre>
  - CLOSE <- CLOSE | | {n}</li>
  - If P.goal\_test(n)
    - Return path(n)
  - For s in P.SUCC (n)
    - o If s ∉ OPEN ∪ CLOSED
      - n' <- make\_node(s, n, h(s))</pre>
      - OPEN.insert(n')
- Return failure

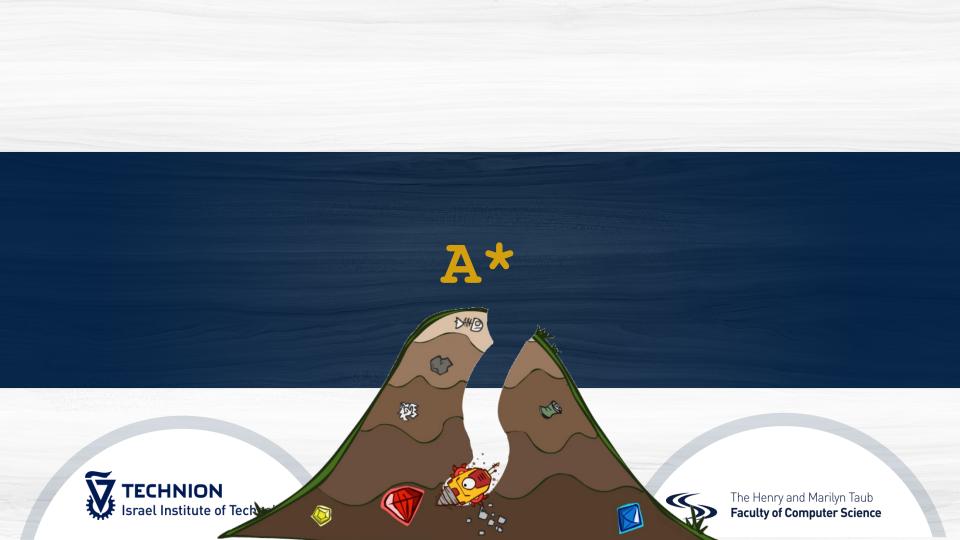
A node consists of:

- State
- Pointer to parent
- h-value



#### **Greedy best-first search - properties**

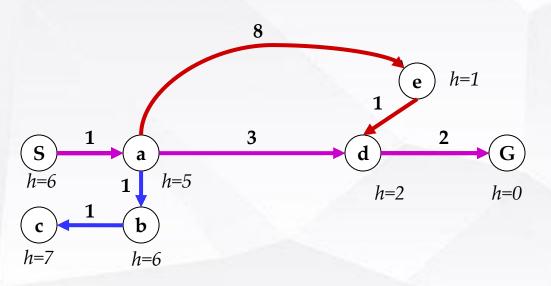
- In finite spaces the algorithm is complete
- Not complete in infinite state spaces
  - Can be arbitrarily bad
- Possibly non-optimal solutions
- ▶ Time and space complexity depends on heuristic
  - Proportional to the number of nodes in OPEN and CLOSED
  - Could be much larger or smaller than BFS

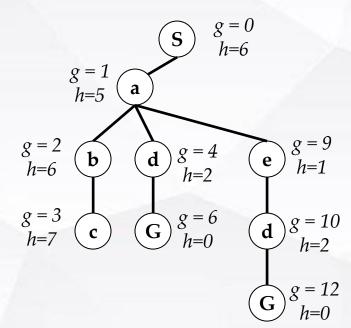




#### **Combining UCS and Greedy**

- Uniform-cost orders by path cost, or backward cost g(n) (aka cost-to-come)
- Greedy orders by goal proximity, or forward cost h(n) (aka cost-to-go)





A\* Search orders by the sum: f(n) = g(n) + h(n)



#### A\* Search – pseudo code (1/2)

#### A\* search (search problem - P)

- OPEN <- make node(P.start, NIL, 0, h(P.start))</p>
- CLOSE <- {}</p>
- While OPEN  $\neq \emptyset$ 
  - n <- OPEN.pop min()</pre>
  - CLOSE <- CLOSE  $\bigcup \{n\}$
  - If P.goal\_test(n)
    - Return path(n)
  - For s in P.SUCC (n)
    - new g <- g(n) + P.COST(n.state,s)</li>
    - o If s ∉ OPEN U CLOSED
      - n' <- make node(s, n, new g, h(s))</pre>
      - OPEN.insert(n')

//order according to **f**-value

A node consists of:

- State
- Pointer to parent
- o g-value
- f-value

//newly-computed cost to reach s



#### A\* Search – pseudo code (2/2)

CLOSED.remove(n\_curr)

```
n_curr <- node in OPEN with state s</p>
     ❖ If new g < g(n curr)</p>
                                                                    //found better path to s
            n_curr <- update node(s, n, new_g , h(s))</pre>
             OPEN.update key(n curr)
                                                    //don't forget to update place in OPEN...
                                                   //else do nothing – existing path is better
Else // s
            CLOSED
     n_@rr <- node in CLOSED with state s</p>
     ❖ If new g < g(n curr)</p>
                                                                   //found better path to s
            n_curr <- update_node(s, n, new_g , h(s))</pre>
             OPEN.insert(n curr)
```

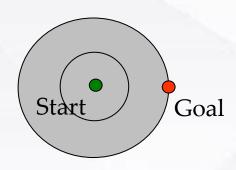
Return failure



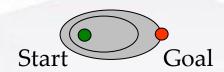


#### **UCS vs A\* Contours**

Uniform-cost expands equally in all "directions"

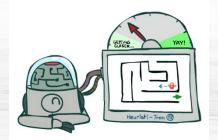


▶ A\* expands mainly toward the goal, but does hedge its bets to ensure optimality



# A\* - optimality (admissible heuristics)









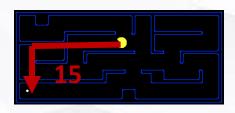
#### **Admissible Heuristics**

A heuristic *h* is *admissible* (optimistic) if:

$$\forall x \ 0 \le h(x) \le h^*(x)$$

where  $h^*(x)$  is the true cost to a nearest goal

**Examples:** 





0.0

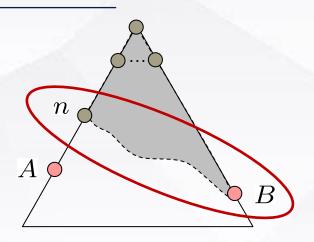
▶ Coming up with admissible heuristics is most of what's involved in using A\* in practice.



#### **Optimality of A\* Tree Search: Blocking**

#### **Proof**:

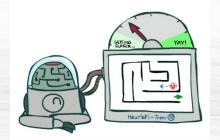
- ▶ Imagine **B** is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- ▶ Claim: n will be expanded before B
  - 1. f(n) is less or equal to f(A)
  - 2. f(A) is less than f(B)
  - 3. n expands before B
- All ancestors of A expand before B
- A expands before B
- A\* search is optimal (with an admissible heuristic)



$$f(n) \le f(A) < f(B)$$

# Efficiency of A\* (heuristics quality)







#### Efficiency of A\*

- ▶ The efficiency of A\* depends on the heuristic quality
- What happens when  $h \equiv 0$  ?

▶ What happens when  $h \equiv h^*$  (optimal heuristic)?



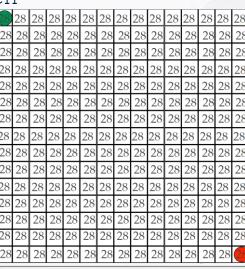
- ▶ The efficiency of A\* depends on the heuristic quality
- What happens when  $h \equiv 0$  ?
  - Uses only g-values => identical to Uniform-cost search
- ▶ What happens when  $h \equiv h^*$  (optimal heuristic)?



- ▶ The efficiency of A\* depends on the heuristic quality
- What happens when  $h \equiv 0$  ?
  - Uses only g-values => identical to Uniform-cost search
- Mhat happens when  $h \equiv h^*$  (optimal heuristic)?
  - Will only expand nodes along an optimal path (prove at home)
  - What is the difference between this setting and greedybest first with a perfect heuristic?
  - What if there is more than one optimal path?



- ▶ The efficiency of A\* depends on the heuristic quality
- What happens when  $h \equiv 0$ ?
  - Uses only g-values => identical to Uniform-cost search
- ▶ What happens when  $h \equiv h^*$  (optimal heuristic)?
  - Will only expand nodes along an optimal path (prove at home)
  - What is the difference between this setting and greedybest first with a perfect heuristic?
  - What if there is more than one optimal path?

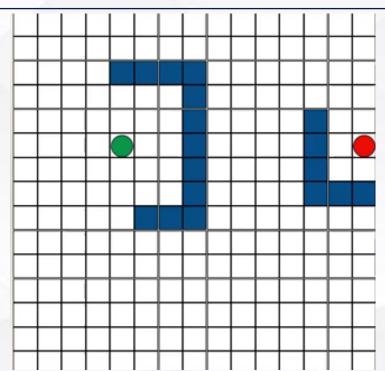




- What happens in the general case?
  - An informative heuristic allows ignoring paths that cannot be part of an optimal path
  - Typically, the more informative the heuristic is the more focused the search can be
  - Sometimes, computing a highly-informative heuristics takes more time than using a simple (but effective) one



## **Greedy Best-first Search Demo**

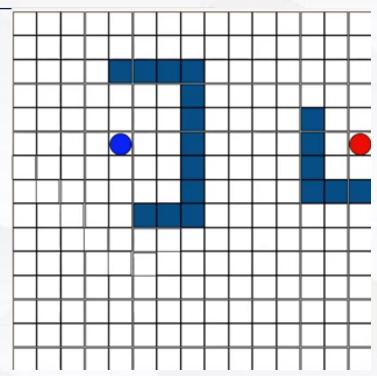


Visualization by Shaul Markovitz, video available at https://youtu.be/n1Cm1XiGd48



#### **Greedy best-first search demo**

▶ Heuristic: Manhattan distance

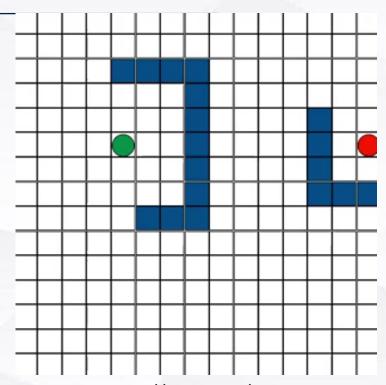


Visualization by Shaul Markovitz, video available at https://youtu.be/TdHbO3w68fY



#### A\* demo

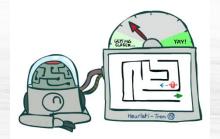
▶ Heuristic: Manhattan distance



Visualization by Shaul Markovitz, video available at https://youtu.be/F6HQ\_KzWuYQ

# A\* (consistent heuristics)









$$\forall s \in S, \ \forall s' \in \text{SUCC}(s) \ h(s) - h(s') \le \text{COST}(s, s')$$

- ▶ Roughly speaking a consistent heuristic is not only **globally optimistic** (like admissible heuristics) but also **locally optimistic** 
  - Let's re-write the definition:  $h(s') + \mathrm{COST}(s, s') \geq h(s)$



$$\forall s \in S, \ \forall s' \in \text{SUCC}(s) \ h(s) - h(s') \le \text{COST}(s, s')$$

- ▶ Roughly speaking a consistent heuristic is not only globally optimistic (like admissible heuristics) but also locally optimistic
  - Let's re-write the definition:  $h(s') + \text{COST}(s, s') \ge h(s)$
  - Now, let's show that the f-values monotonically increase during the search f(s') = a(s') + h(s')



$$\forall s \in S, \ \forall s' \in \text{SUCC}(s) \ h(s) - h(s') \le \text{COST}(s, s')$$

- ▶ Roughly speaking a consistent heuristic is not only **globally optimistic** (like admissible heuristics) but also locally optimistic
  - Let's re-write the definition:  $h(s') + \text{COST}(s, s') \ge h(s)$
  - Now, let's show that the f-values monotonically increase during the search

$$f(s') = g(s') + h(s') = g(s) + COST(s, s') + h(s')$$

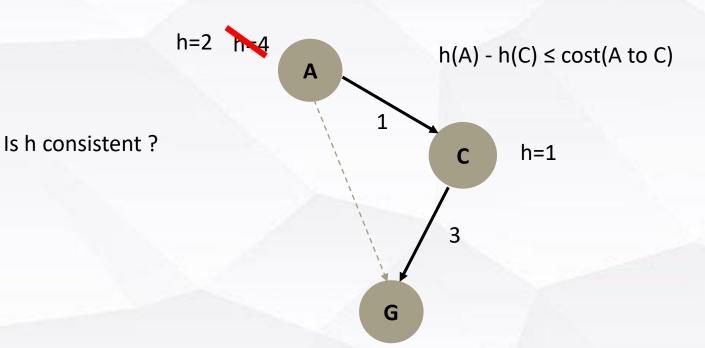


$$\forall s \in S, \ \forall s' \in \text{SUCC}(s) \ h(s) - h(s') \leq \text{COST}(s, s')$$

- ▶ Roughly speaking a consistent heuristic is not only **globally optimistic** (like admissible heuristics) but also **locally optimistic** 
  - Let's re-write the definition:  $h(s') + \mathrm{COST}(s, s') \geq h(s)$
  - Now, let's show that the f-values monotonically increase during the search

$$f(s') = g(s') + h(s') = g(s) + COST(s, s') + h(s') \ge g(s) + h(s) = f(s)$$







- ▶ When A\* uses a consistent heuristic the path to every node that was expanded is optimal (proof in Pearl's book)
- Namely,  $\forall n \in \text{CLOSED } g(n) = g^*(n)$

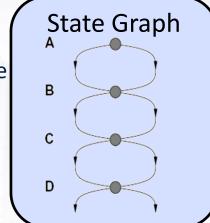
$$\forall s \in CLOSED: g(s) = g^*(s)$$

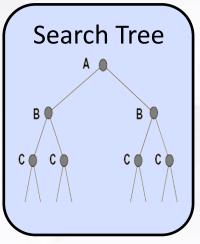
> => When using a consistent heuristic, no need to move nodes back from CLOSED to OPEN



## Closed list (actually not a list but a set...)

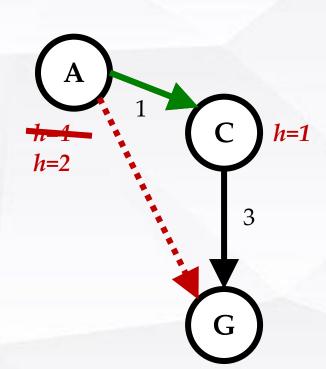
- Used to avoid expanding a state twice
  - Failure to detect repeated states can cause exponentially more work
- But why do we need to move nodes from the CLOSED list to the OPEN list?
  - Because the heuristic mislead us...







# **Admissibility and Consistency of Heuristics**



- Main idea: estimated heuristic costs ≤ actual costs
  - **Admissibility**: heuristic cost ≤ actual cost to goal  $h(A) \le actual cost from A to G$
  - **Consistency**: heuristic "arc" cost ≤ actual cost for each arc  $h(A) - h(C) \le cost(A to C)$
- Consequences of admissibility:
  - A\* search is optimal
- Consequences of consistency:
  - The **f**-value along a path never decreases

$$h(A) \le cost(A \text{ to } C) + h(C)$$

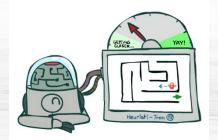


# **Optimality of A\*'s efficiency**

- When A\* uses a consistent heuristic it can be shown that any other (optimal) search algorithm that uses the same set heuristic expands at least the same number of nodes
- Is that what we always care about?

# Computing heuristics



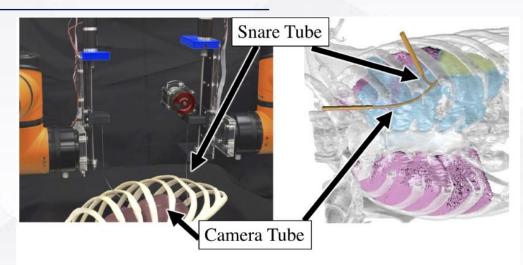


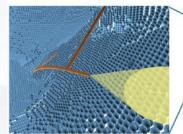


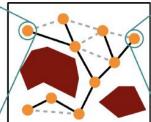


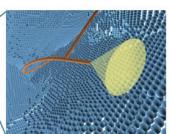
## **A\*** Applications

- Video games
- Path finding / routing problems
- Resource-planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition







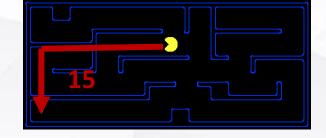




## **Creating Admissible Heuristics**

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to relaxed problems, where new actions are available

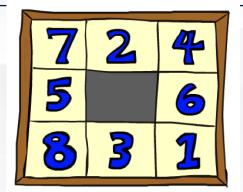




Inadmissible heuristics are often useful too

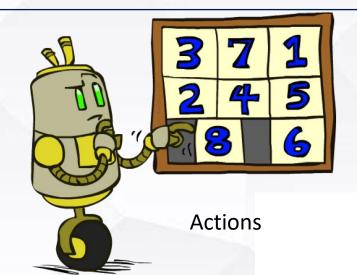


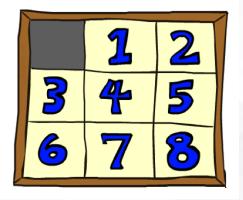
# **Example: 8 Puzzle**





- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?



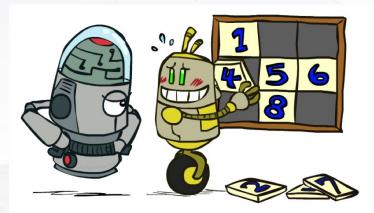


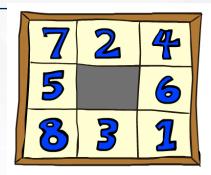
**Goal State** 

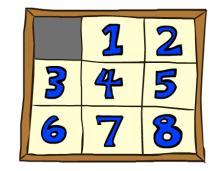
Admissible heuristics?



- Heuristic: Number of tiles misplaced
- Why is it admissible?
- ▶ h(start) = 8
- ▶ This is a *relaxed-problem* heuristic







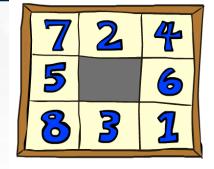
**Start State** 

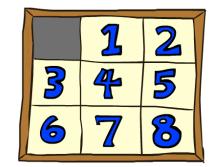
**Goal State** 

|            | Average nodes expanded when the optimal path has |         |                   |  |
|------------|--|---------|-------------------|--|
|            | 4 steps  | 8 steps | 12 steps          |  |
| UCS        | 112  | 6,300   | $3.6 \times 10^6$ |  |
| A* + TILES | 13   | 39      | 227               |  |

# 8 Puzzle II

What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?





▶ Total *Manhattan* distance

Start State

**Goal State** 

| Why | is it | admis | sible?  |
|-----|-------|-------|---------|
| ,   |       | 0. 0  | 0.10.0. |

$$h(start) = 3 + 1 + 2 + ... = 18$$

|                | optimal path has |         |          |  |
|----------------|------------------|---------|----------|--|
|                | 4 steps          | 8 steps | 12 steps |  |
| A* + TILES     | 13               | 39      | 227      |  |
| A* + MANHATTAN | 12               | 25      | 73       |  |

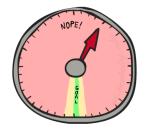
Average nodes expanded when the



- How about using the actual cost as a heuristic?
  - Would it be admissible?
  - Would we save on nodes expanded?
  - What's wrong with it?







- ▶ With A\*: a trade-off between quality of estimate and work per node
  - As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

# Bounded suboptimal search Trading quality for efficiency







#### Trading quality for efficiency in a structured manner

- ▶ The A\* algorithm computes an optimal solution
- Sometimes we are willing to compromise slightly on the quality of the solution if we can get it faster
- An algorithm is said to be bounded suboptimal if the quality of the solution is bounded by a multiplicative factor when compared to the optimal solution

# Weighted A\*

lacktriangle Exactly like A\* but given a user-provided  $w \in [0,1]$  , orders nodes in OPEN according to

$$f(n) = (1 - w) \cdot g(n) + w \cdot h(n)$$

- What happens when
  - $\mathbf{w} = \mathbf{0}$ ?
  - w=1?
  - w = 0.5?
  - w<0.5?

# Weighted A\*

lacktriangle Exactly like  $\mathbb{A}^{\star}$  but given a user-provided  $w \in [0,1]$  , orders nodes in OPEN according to

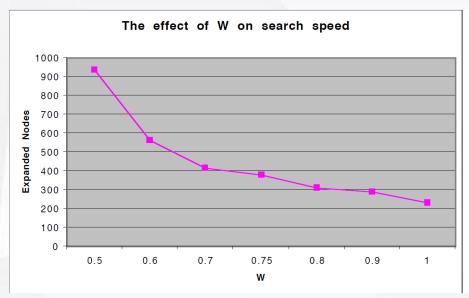
$$f(n) = (1 - w) \cdot g(n) + w \cdot h(n)$$

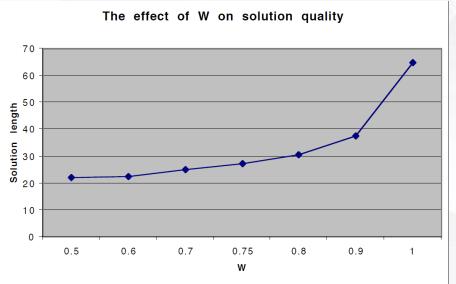
- What happens when
  - w=0? UCS
  - w=1? Greedy Best First Search
  - w=0.5? A\*
  - w<0.5?



#### **Empirical demonstration of wA\***

Every point is the average result of 100 random instances of the alg. On a 3X3 puzzle using the Manhattan heuristic







#### Weighted A\* (alternative formulation)

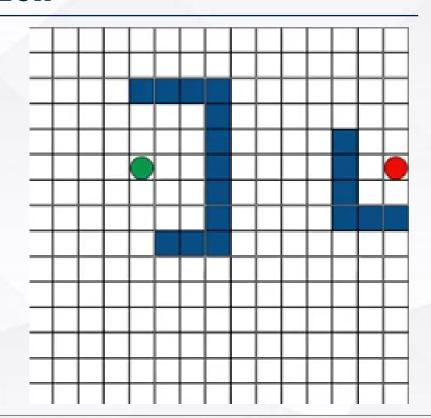
- Alternative formulation: given a user-provided  $w \ge 1$  , orders nodes in OPEN according to  $f(n) = g(n) + w \cdot h(n)$
- **Thm**: wA\* (under this formulation) returns a solution whose cost is at most  $w\cdot C^*$  where,  $C^*$  is the cost of the optimal solution



- $\blacktriangleright$  Similar to A\* in addition to a heuristic function h we are also given
  - a user-provided approximation factor  $\varepsilon \geq 0$
  - lacktriangle a (possibly-inadmissible) heuristic function  $h_{
    m focal}$
- ▶ In A\* we expand a node with the minimal f-value
- Here, we
  - Compute the set of "almost best nodes" according to h  $FOCAL = \{n \in OPEN \mid f(n) \leq (1+\varepsilon) \cdot \min_{n' \in OPEN} f(n')\}$
  - Choose a node according to  $h_{\text{focal}}$  with  $\min_{n \in \text{FOCAL}} h_{\text{FOCAL}}(n)$



• Using  $\varepsilon=0.1$ 





## A\*-epsilon - bounded suboptimality

- Thm: A\*-epsilon returns a solution whose cost is at most  $(1+\varepsilon)\cdot C^*$  where,  $C^*$  is the cost of the optimal solution under the condition that h is admissible.
- Proof: Let
  - $n_0$  be the node in OPEN with the minimal f-value (has to be in FOCAL)
  - ullet be the node returned by <code>A\*-epsilon</code> when terminated
  - $n_\ell$  be a node in OPEN that belongs to the optimal path with all ancestors in CLOSED
  - The cost of the solution returned is C(t) = f(t) = g(t)



## A\*-epsilon - bounded suboptimality (cont.)

#### Note that

- $f(n_{\ell}) \leq C^*$  since the heuristic is admissible
- $f(n_0) \leq f(n_\ell)$  by definition of  $n_0$
- $f(t) \leq (1+\varepsilon) \cdot f(n_0)$  since t was chosen from FOCAL

#### ▶ Thus,

$$C(t) = f(t) \le (1 + \varepsilon) \cdot f(n_0) \le (1 + \varepsilon) \cdot f(n_\ell) \le (1 + \varepsilon) \cdot C^*$$

# IDA\*







- ▶ One of the main drawbacks of best-first search algorithms (like A\*) is there large memory footprint (due to the OPEN list)
- ▶ Iterative deepening A\* (IDA\*) uses f-values to bound the depth of the search it performs
- It starts with a small bound on the maximal f-values and increases it between iterations
  - Start with a bound of  $f \leq h(s_{\text{start}})$
  - Bound in the next iteration is the minimal f-value found in the current iteration (that is larger than the existing bound)



#### **IDA\*** Search – pseudo code

```
function Iterative-Deepening-A* (problem):

new\_limit \leftarrow h(problem.init\_state)

While Not Interrupted:

f\_limit \leftarrow new\_limit

new\_limit \leftarrow \infty /* Global variable */

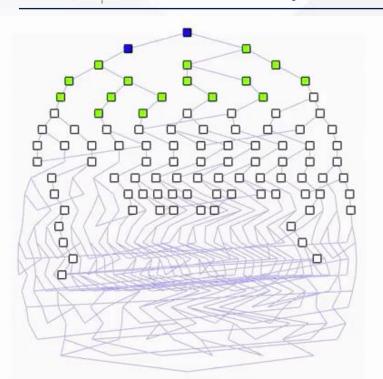
result \leftarrow DFS-f (problem.init_state, o, null, f_limit, problem)

if result \neq failure then return result

return failure
```



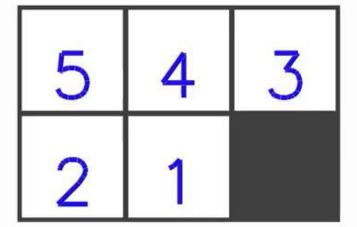
#### **Visualization (2X3 sliding puzzle)**





#### Video adapted from

https://movingai.com/





# **Summary**

| Algorithm |                          | Complete | Optimal | Time   | Space  |
|-----------|--------------------------|----------|---------|--------|--------|
| Blind     | Breadth First Search     | Υ        | γ*      | O(b^m) | O(b^m) |
|           | Depth First Search       | N**      | N       | O(b^m) | O(mb)  |
|           | Iterative Deepening      | Υ        | Υ       | O(b^m) | O(mb)  |
|           | Uniform Cost Search      | γ***     | γ***    | O(b^m) | O(b^m) |
| Informed  | Greedy Best First Search | N**      | N       | O(b^m) | O(b^m) |
|           | A*                       | Y***+^   | Y***+^  | O(b^m) | O(b^m) |
|           | IDA*                     | Y***+^   | Y***+^  | O(b^m) | O(mb)  |

<sup>\*</sup> When edge costs are one

<sup>\*\*</sup> Y when finite graph with no cycles

<sup>\*\*\*</sup> When costs are positive and lower bounded

<sup>^</sup> When j is admissible